



Linux kernel and driver development training

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Latest update: March 20, 2026.

Document updates and training details:
<https://bootlin.com/training/linux-kernel>


Corrections, suggestions, contributions and translations are welcome!
Send them to feedback@bootlin.com





Linux kernel and driver development training

- ▶ These slides are the training materials for Bootlin's *Linux kernel and driver development* training course.
- ▶ If you are interested in following this course with an experienced Bootlin trainer, we offer:
 - **Public online sessions**, opened to individual registration. Dates announced on our site, registration directly online.
 - **Dedicated online sessions**, organized for a team of engineers from the same company at a date/time chosen by our customer.
 - **Dedicated on-site sessions**, organized for a team of engineers from the same company, we send a Bootlin trainer on-site to deliver the training.
- ▶ Details and registrations:
<https://bootlin.com/training/linux-kernel>
- ▶ Contact: training@bootlin.com




Icon by Eucalyp, Flaticon



About Bootlin

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Bootlin introduction

- ▶ Engineering company
 - In business since 2004
 - Before 2018: *Free Electrons*
- ▶ Team based in France and Italy
- ▶ Serving **customers worldwide**
- ▶ **Highly focused and recognized expertise**
 - Embedded Linux
 - Linux kernel
 - Embedded Linux build systems
- ▶ **Strong open-source** contributor
- ▶ Activities
 - **Engineering** services
 - **Training** courses
- ▶ <https://bootlin.com>

bootlin



Bootloader / firmware development

U-Boot, Barebox,
OP-TEE, TF-A, .../

Linux kernel porting and driver development

Linux BSP development, maintenance and upgrade

Embedded Linux build systems

Yocto, OpenEmbedded,
Buildroot, ...

Embedded Linux integration

Boot time, real-time,
security, multimedia,
networking

Open-source upstreaming

Get code integrated
in upstream
Linux, U-Boot, Yocto,
Buildroot, ...



Bootlin training courses

Embedded Linux system development

On-site: 4 or 5 days
Online: 7 * 4 hours

Linux kernel driver development

On-site: 5 days
Online: 7 * 4 hours

Yocto Project system development

On-site: 3 days
Online: 4 * 4 hours

Buildroot system development

On-site: 3 days
Online: 5 * 4 hours

Embedded Linux networking

On-site: 3 days
Online: 4 * 4 hours

Understanding the Linux graphics stack

On-site: 2 days
Online: 4 * 4 hours

Embedded Linux audio

On-site: 2 days
Online: 4 * 4 hours

Real-Time Linux with PREEMPT_RT

On-site: 2 days
Online: 3 * 4 hours

Linux debugging, tracing, profiling and performance analysis

On-site: 3 days
Online: 4 * 4 hours

All our training materials are freely available
under a free documentation license (CC-BY-SA 3.0)
See <https://bootlin.com/training/>



Bootlin, an open-source contributor

- ▶ Strong contributor to the **Linux** kernel
 - In the top 30 of companies contributing to Linux worldwide
 - Contributions in most areas related to hardware support
 - Several engineers maintainers of subsystems/platforms
 - 9000 patches contributed
 - <https://bootlin.com/community/contributions/kernel-contributions/>
- ▶ Contributor to **Yocto Project**
 - Maintainer of the official documentation
 - Core participant to the QA effort
- ▶ Contributor to **Buildroot**
 - Co-maintainer
 - 6000 patches contributed
- ▶ Significant contributions to U-Boot, OP-TEE, Barebox, etc.
- ▶ Fully **open-source training materials**



Bootlin on-line resources

- ▶ Website with a technical blog:
<https://bootlin.com>
- ▶ Engineering services:
<https://bootlin.com/engineering>
- ▶ Training services:
<https://bootlin.com/training>
- ▶ LinkedIn:
<https://www.linkedin.com/company/bootlin>
- ▶ Elixir - browse Linux kernel sources on-line:
<https://elixir.bootlin.com>




Icon by Freepik, Flaticon



Generic course information

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Beaglebone Black / Beaglebone black wireless shopping list

- ▶ BeagleBone Black or BeagleBone Black Wireless, from BeagleBoard.org
 - Texas Instruments AM335x (ARM Cortex-A8 CPU)
 - 512 MB of RAM
 - 4 GB of on-board eMMC storage
 - Plenty of peripherals and features
 - 2 x 46 pins headers, with access to many expansion buses (I2C, SPI, UART and more)
- ▶ MicroUSB cable
- ▶ USB Serial Cable - 3.3 V - Female ends (for serial console) ¹
- ▶ Nintendo Nunchuk with UEXT connector ²
- ▶ Breadboard jumper wires - Male ends (to connect the Nunchuk) ³
- ▶ MicroSD card



¹ <https://www.olimex.com/Products/USB-Modules/Interfaces/USB-SERIAL-F>

² <https://www.olimex.com/Products/Modules/Sensors/MOD-WII/MOD-Wii-UEXT-NUNCHUK/>

³ <https://www.olimex.com/Products/Breadboarding/JUMPER-WIRES/JW-110x10/>



Beagleplay shopping list

- ▶ BeaglePlay, from [BeagleBoard.org](https://beagleboard.org)
 - Texas Instruments AM625x (4xARM Cortex-A53 CPU)
 - 2 GB of RAM
 - 16 GB of on-board eMMC storage
 - Plenty of peripherals: SPI, I2C, UART, USB...
- ▶ USB-C cable for the power supply
- ▶ A USB-FTDI cable
- ▶ RJ45 cable for networking
- ▶ A micro SD card with at least 2G of capacity
- ▶ Nintendo Nunchuk with UEXT connector ¹
- ▶ Breadboard jumper wires - Male ends (to connect the Nunchuk) ²



¹ <https://www.olimex.com/Products/Modules/Sensors/MOD-WII/MOD-Wii-UEXT-NUNCHUCK/>

² <https://www.olimex.com/Products/Breadboarding/JUMPER-WIRES/JW-110x10/>



IMX93 FRDM shopping list

- ▶ NXP i.MX93 11x11 FRDM board Available from Mouser (76 EUR + VAT)
 - NXP i.MX 93 (Dual ARM Cortex-A55 + Cortex-M33)
 - 2 GB LPDDR4
 - 32 GB of on-board eMMC storage
 - Plenty of peripherals: I2C, SPI, UART, USB...
- ▶ 2 USB-C cable for the power supply and the serial console
- ▶ RJ45 cable for networking
- ▶ USB Serial Cable - 3.3 V - Female ends (for serial labs, two if possible) ¹
imx93-frdm-nunchuk
- ▶ Nintendo Nunchuk with UEXT connector ²
- ▶ Breadboard jumper wires - Male/Female ends (to connect the Nunchuk) ³
- ▶ RJ45 cable for networking

¹ <https://www.olimex.com/Products/USB-Modules/Interfaces/USB-SERIAL-F>

² <https://www.olimex.com/Products/Modules/Sensors/MOD-WII/MOD-Wii-UEXT-NUNCHUK/>

³ <https://www.olimex.com/Products/Breadboarding/JUMPER-WIRES/JW-200x10-FM/>



imx93-frdm-nunchuk





Training quiz and certificate

- ▶ To get your training certificate you must
 1. Attend all sessions of this training course
 2. Achieve more than 50% of correct answers at our final quiz
 - The final quiz questions are identical to the pre-training quiz
 - The final quiz must be completed within two weeks of the session end's date
- ▶ The training certificate will be sent to you two weeks after the session end's date.



Participate!

During the lectures...

- ▶ Don't hesitate to ask questions. Other people in the audience may have similar questions too.
- ▶ Don't hesitate to share your experience too, for example to compare Linux with other operating systems you know.
- ▶ Your point of view is most valuable, because it can be similar to your colleagues' and different from the trainer's.
- ▶ In on-line sessions
 - Please always keep your camera on!
 - Also make sure your name is properly filled.
 - You can also use the "Raise your hand" button when you wish to ask a question but don't want to interrupt.
- ▶ All this helps the trainer to engage with participants, see when something needs clarifying and make the session more interactive, enjoyable and useful for everyone.



- ▶ Use the dedicated Matrix channel for this session to add questions.
- ▶ If your session offers practical labs, you can also report issues, share screenshots and command output there.
- ▶ Don't hesitate to share your own answers and to help others especially when the trainer is unavailable.
- ▶ The Matrix channel is also a good place to ask questions outside of training hours, and after the course is over.

 embedded-linux-nov2020 Channel for

 **Srinath**

 **michael.o**: What should be CROSS_COMPILE variable set to in case of the XPlained board? I ran into some issues with my USB hub so doing the u-boot again

 **michael.o**

 **Srinath**: you should look at the name of the cross-compiler in the toolchain's bin/ directory. CROSS_COMPILE should be set to what's before "gcc" in the name, including the trailing "-". Like if the compiler is arm-buildroot-linux-gcc, CROSS_COMPILE should be arm-buildroot-linux-

 2 messages deleted.


 **Srinath**

We'll ask them here since I am going to do labs after the session is over! Thanks!

 **Srinath** changed their display name to srinath.f.

 **Anand A**

I tried to finalize Kernel - Cross-compiling was the new kernel. Does anyone know what can be the root cause?



 Decrypt Image.png (109.2 KB)

 **arnaud.a**

I had the same because I accidentally executed the `rm -rf ..` from the journal

 Send an encrypted message...



Prepare your lab environment

- ▶ Download and extract the lab archive




Linux Kernel Introduction

© Copyright 2004-2026, Bootlin.

Creative Commons BY-SA 3.0 license.

Corrections, suggestions, contributions and translations are welcome!





Origin


- ▶ The Linux kernel was created as a hobby in 1991 by a Finnish student, Linus Torvalds.
 - Linux quickly started to be used as the kernel for free software operating systems
- ▶ Linus Torvalds has been able to create a large and dynamic developer and user community around Linux.
- ▶ As of today, about 2,000+ people contribute to each kernel release, individuals or companies big and small.



Linus Torvalds in 2014
Image credits (Wikipedia):
<https://bit.ly/2UIa1TD>



Linux kernel in the system





Linux kernel main roles

- ▶ **Manage all the hardware resources:** CPU, memory, I/O.
- ▶ Provide a **set of portable, architecture and hardware independent APIs** to allow user space applications and libraries to use the hardware resources.
- ▶ **Handle concurrent accesses and usage** of hardware resources from different applications.
 - Example: a single network interface is used by multiple user space applications through various network connections. The kernel is responsible for “multiplexing” the hardware resource.



System calls

- ▶ The main interface between the kernel and user space is the set of system calls
- ▶ About 400 system calls that provide the main kernel services
 - File and device operations, networking operations, inter-process communication, process management, memory mapping, timers, threads, synchronization primitives, etc.
- ▶ This interface is stable over time: only new system calls can be added by the kernel developers This system call interface is wrapped by the C library, and user space applications usually never make a system call directly but rather use the corresponding C library function



Image credits (Wikipedia):

<https://bit.ly/2U2rdGB>



Pseudo filesystems

- ▶ Linux makes system and kernel information available in user space through **pseudo filesystems**, sometimes also called **virtual filesystems**
- ▶ Pseudo filesystems allow applications to see directories and files that do not exist on any real storage: they are created and updated on the fly by the kernel
- ▶ The two most important pseudo filesystems are
 - `proc`, usually mounted on `/proc`:
Operating system related information (processes, memory management parameters...)
 - `sysfs`, usually mounted on `/sys`:
Representation of the system as a tree of devices connected by buses. Information gathered by the kernel frameworks managing these devices.



Linux kernel sources



Location of the official kernel sources

- ▶ The mainline versions of the Linux kernel, as released by Torvalds
 - These versions follow the development model of the kernel (`master` branch)
 - They may not contain the latest developments from a specific area yet
 - A good pick for products development phase
 - <https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git>




Linux versioning scheme

- ▶ Until 2003, there was a new “stabilized” release branch of Linux every 2 or 3 years (2.0, 2.2, 2.4). Development branches took 2-3 years to be merged (too slow!).
- ▶ Since 2003, there is a new official release of Linux about every 10 weeks:
 - Versions 2.6 (Dec. 2003) to 2.6.39 (May 2011)
 - Versions 3.0 (Jul. 2011) to 3.19 (Feb. 2015)
 - Versions 4.0 (Apr. 2015) to 4.20 (Dec. 2018)
 - Versions 5.0 (Mar. 2019) to 5.19 (July 2022)
 - Versions 6.0 (Oct. 2022) to 6.19 (Feb. 2026)
 - Version 7.0 will be released in March/April 2026.
- ▶ Features are added to the kernel in a progressive way. Since 2003, kernel developers have managed to do so without having to introduce a massively incompatible development branch.
- ▶ Major updates **have no specific meaning!**



Linux development model


- ▶ Each new release starts with a two-week merge window for new features
- ▶ Follow about 8 release candidates (one week each)
- ▶ Until adoption of a new official release.





Need to further stabilize the official kernels

- ▶ Issue: bug and security fixes only merged into the master branch, need to update to the latest kernel to benefit from them.
- ▶ Solution: a stable maintainers team goes through all the patches merged into Torvald's tree and backports the relevant ones into their stable branches for at least a few months.





Location of the stable kernel sources

- ▶ The stable versions of the Linux kernel, as maintained by a maintainers group
 - These versions do not bring new features compared to Linus' tree
 - Only bug fixes and security fixes are pulled there
 - Each version is stabilized during the development period of the next mainline kernel
 - A good pick for products commercialization phase
 - <https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git>
 - Certain versions will be maintained much longer



Need for long term support

- ▶ Issue: bug and security fixes only released for most recent kernel versions.
- ▶ Solution: the last release of each year is made an LTS (*Long Term Support*) release, and is supposed to be supported (and receive bug and security fixes) for at least 2 years. These projected EOL may be extended up to 6 years based on the industry interest at large.

Version	Maintainer	Released	Projected EOL
6.18	Greg Kroah-Hartman & Sasha Levin	2025-11-30	Dec, 2027
6.12	Greg Kroah-Hartman & Sasha Levin	2024-11-17	Dec, 2026
6.6	Greg Kroah-Hartman & Sasha Levin	2023-10-29	Dec, 2026
6.1	Greg Kroah-Hartman & Sasha Levin	2022-12-11	Dec, 2027
5.15	Greg Kroah-Hartman & Sasha Levin	2021-10-31	Dec, 2026
5.10	Greg Kroah-Hartman & Sasha Levin	2020-12-13	Dec, 2026

Captured on <https://kernel.org> in Feb. 2026, following the [Releases](#) link.

- ▶ Example at Google: starting from *Android O (2017)*, all new Android devices have to run such an LTS kernel.



Need for even longer term support

- ▶ You could also get long term support from a commercial embedded Linux provider.
 - Wind River Linux can be supported for up to 15 years.
 - Ubuntu Core can be supported for up to 10 years.
- ▶ *"If you are not using a supported distribution kernel, or a stable / longterm kernel, you have an insecure kernel"* - Greg KH, 2019
Some vulnerabilities are fixed in stable without ever getting a CVE.
- ▶ The *Civil Infrastructure Platform* project is an industry / Linux Foundation effort to support much longer (at least 10 years) selected LTS versions (currently 4.4, 4.19, 5.10, 6.1 and 6.12) on selected architectures. See <https://wiki.linuxfoundation.org/civilinfrastructureplatform/start>.



Location of non-official kernel sources

- ▶ Many chip vendors supply their own kernel sources
 - Focusing on hardware support first
 - Can have a very important delta with mainline Linux
 - Sometimes they break support for other platforms/devices without caring
 - Useful in early phases only when mainline hasn't caught up yet (many vendors invest in the mainline kernel at the same time)
 - Suitable for PoC, not suitable for products on the long term as usually no updates are provided to these kernels
 - Getting stuck with a deprecated system with broken software that cannot be updated has a real cost in the end
- ▶ Many kernel sub-communities maintain their own kernel, with usually newer but fewer stable features, only for cutting-edge development
 - Architecture communities (ARM, MIPS, PowerPC, etc)
 - Device drivers communities (I2C, SPI, USB, PCI, network, etc)
 - Other communities (filesystems, memory-management, scheduling, etc)
 - Not suitable to be used in products



Linux kernel size and structure

- ▶ Linux v6.19 sources: close to 92k files, 42M lines, 1.4GiB
- ▶ But a compressed Linux kernel just sizes a few megabytes.
- ▶ So, why are these sources so big?
Because they include numerous device drivers, network protocols, architectures, filesystems... The core is pretty small!
- ▶ As of kernel version v6.19 (in percentage of total number of lines):

- | | | |
|--------------------------------------|--------------------------------|--|
| ▶ <code>drivers/</code> : 61% | ▶ <code>include/</code> : 3.4% | ▶ <code>security/</code> , <code>rust/</code> , <code>crypto/</code> , <code>block/</code> , |
| ▶ <code>arch/</code> : 12% | ▶ <code>net/</code> : 3.1% | <code>samples/</code> , <code>io_uring/</code> , <code>virt/</code> , <code>usr/</code> , |
| ▶ <code>tools/</code> : 5.4% | ▶ <code>kernel/</code> : 1.3% | <code>LICENSES/</code> , <code>ipc/</code> , <code>init/</code> , <code>certs/</code> : $\leq 0.3\%$ |
| ▶ <code>sound/</code> : 3.9% | ▶ <code>lib/</code> : 0.8% | ▶ Build system files: <code>Kbuild</code> , <code>Kconfig</code> , |
| ▶ <code>fs/</code> : 3.8% | ▶ <code>mm/</code> : 0.5% | <code>Makefile</code> |
| ▶ <code>Documentation/</code> : 3.8% | ▶ <code>scripts/</code> : 0.3% | ▶ Other files: <code>COPYING</code> , <code>CREDITS</code> , |
| | | <code>MAINTAINERS</code> , <code>README</code> |



- ▶ Clone the mainline Linux source tree with git



Linux kernel source code



Programming language

- ▶ Implemented in C like all UNIX systems
- ▶ A little Assembly is used too:
 - CPU and machine initialization, exceptions
 - Critical library routines.
- ▶ No C++ used, see <https://lkml.org/lkml/2004/1/20/20>
- ▶ Rust support is currently being introduced: [drivers/net/phy/ax88796b_rust.rs](#) is a first driver written in Rust.
- ▶ All the code compiled with gcc
 - Many gcc specific extensions used in the kernel code, any ANSI C compiler will not compile the kernel
 - See <https://gcc.gnu.org/onlinedocs/gcc-10.2.0/gcc/C-Extensions.html>
- ▶ A subset of the supported architectures can be built with the LLVM C compiler (Clang) too: <https://clangbuiltlinux.github.io/>



No C library

- ▶ The kernel has to be standalone and can't use user space code.
- ▶ Architectural reason: user space is implemented on top of kernel services, not the opposite.
- ▶ Technical reason: the kernel is on its own during the boot up phase, before it has accessed a root filesystem.
- ▶ Hence, kernel code has to supply its own library implementations (string utilities, cryptography, uncompression...)
- ▶ So, you can't use standard C library functions in kernel code (`printf()`, `memset()`, `malloc()`,...).
- ▶ Fortunately, the kernel provides similar C functions for your convenience, like `printk()`, `memset()`, `kmalloc()`, ...



- ▶ The Linux kernel code is designed to be portable
- ▶ All code outside `arch/` should be portable
- ▶ To this aim, the kernel provides macros and functions to abstract the architecture specific details
 - Endianness
 - I/O memory access
 - Memory barriers to provide ordering guarantees if needed
 - DMA API to flush and invalidate caches if needed
- ▶ Never use floating point numbers in kernel code. Your code may need to run on a low-end processor without a floating point unit.



Linux kernel to user API/ABI stability

Linux kernel to userspace API is stable

- ▶ Source code for userspace applications will not have to be updated when compiling for a more recent kernel
 - System calls, `/proc` and `/sys` content cannot be removed or changed. Only new entries can be added.

Linux kernel to userspace ABI is stable

- ▶ Binaries are portable and can be executed on a more recent kernel
 - The way memory is accessed, the size of the variables in memory, how structures are organized, the calling convention, etc, are all stable over time.


Linux kernel to user **API**

- ☑ API stability **is** guaranteed, source code is portable!



Linux kernel to user **ABI**

- ☑ compatible ABI **can be** guaranteed, binaries are portable



Modified Image from Wikipedia:

<https://bit.ly/2U2rdGB>



Linux internal API/ABI instability

Linux internal API is not stable


- ▶ The source code of a driver is not portable across versions
 - In-tree drivers are updated by the developer proposing the API change: works great for mainline code
 - An out-of-tree driver compiled for a given version may no longer compile or work on a more recent one
 - See [process/stable-api-nonsense](#) for reasons why

Linux internal ABI is not stable

- ▶ A binary module compiled for a given kernel version cannot be used with another version
 - The module loading utilities will perform this check prior to the insertion


Linux internal API

- ☒ API stability is **not** guaranteed, source code portability is not given



Linux internal ABI

- ☒ **no** stable ABI over Linux kernel releases, binaries are not portable



Modified Image from Wikipedia:
<https://bit.ly/2U2rdGB>



Kernel memory constraints

- ▶ No memory protection
- ▶ The kernel doesn't try to recover from attempts to access illegal memory locations. It just dumps *oops* messages on the system console.
- ▶ Fixed size stack (8 or 4 KB). Unlike in user space, no mechanism was implemented to make it grow. Don't use recursion!
- ▶ Swapping is not implemented for kernel memory either (except *tmpfs* which lives completely in the page cache and on swap)



Linux kernel licensing constraints

- ▶ The Linux kernel is licensed under the GNU General Public License version 2
 - This license gives you the right to use, study, modify and share the software freely
- ▶ However, when the software is redistributed, either modified or unmodified, the GPL requires that you redistribute the software under the same license, with the source code
 - If modifications are made to the Linux kernel (for example to adapt it to your hardware), it is a derivative work of the kernel, and therefore must be released under GPLv2.
- ▶ The GPL license has been successfully enforced in courts:
https://en.wikipedia.org/wiki/Gpl-violations.org#Notable_victories
- ▶ However, you're only required to do so
 - At the time the device starts to be distributed
 - To your customers, not to the entire world




Proprietary code and the kernel

- ▶ It is illegal to distribute a binary kernel that includes statically compiled proprietary drivers
- ▶ The kernel modules are a gray area: unclear if they are legal or not
 - The general opinion of the kernel community is that proprietary modules are bad: [process/kernel-driver-statement](#)
 - From a legal point of view, each driver is probably a different case:
 - Are they derived works of the kernel?
 - Are they designed to be used with another operating system?
- ▶ Is it really useful to keep drivers secret anyway?




Abusing the kernel licensing constraints

- ▶ There are some examples of proprietary drivers
 - Nvidia uses a wrapper between their drivers and the kernel
 - They claim the drivers could be used with a different OS with another wrapper
 - Unclear whether it makes it legal or not



- ▶ The current trend is to hide the logic in the firmware or in userspace. The GPL kernel driver is almost empty and either:
 - Blindly writes an incoming flow of bytes in the hardware
 - Exposes a huge MMIO region to userspace through `mmap`





Advantages of GPL drivers

- ▶ You don't have to write your driver from scratch. You can reuse code from similar free software drivers.
- ▶ Your drivers can be freely and easily shipped by others (for example by Linux distributions or embedded Linux build systems).
- ▶ Legal certainty, you are sure that a GPL driver is fine from a legal point of view.



Advantages of mainlining your kernel drivers

- ▶ The community, reviewers and maintainers will review your code before accepting it, offering you the opportunity to enhance it and understand better the internal APIs.
- ▶ Once accepted, you will get cost-free bug and security fixes, support for new features, and general improvements.
- ▶ Your work will automatically follow the API changes.
- ▶ Accessing your code will be much easier for users.
- ▶ Your code will remain valid no matter the kernel version.

This will for sure reduce your maintenance and support work



User space device drivers 1/2

- ▶ The kernel provides some mechanisms to access hardware from userspace:
 - USB devices with *libusb*, <https://libusb.info/>
 - SPI devices with *spidev*, [spi/spidev](https://github.com/spidev/spidev)
 - I2C devices with *i2cdev*, [i2c/dev-interface](https://github.com/linux-i2c/i2c-dev)
 - GPIOs with *libgpiod*, <https://libgpiod.readthedocs.io>
 - Memory-mapped devices with *UIO*, including interrupt handling, [driver-api/uio-howto](https://www.kernel.org/doc/Documentation/driver-api/uio-howto)
- ▶ These solutions can only be used if:
 - There is no need to leverage an existing kernel subsystem such as the networking stack or filesystems.
 - There is no need for the kernel to act as a “multiplexer” for the device: only one application accesses the device.
- ▶ Certain classes of devices like printers and scanners do not have any kernel support, they have always been handled in user space for historical reasons.
- ▶ Otherwise this is **not** how the system should be architected. Kernel drivers should always be preferred!



User space device drivers 2/2

► Advantages

- No need for kernel coding skills.
- Drivers can be written in any language, even Perl!
- Drivers can be kept proprietary.
- Driver code can be killed and debugged. Cannot crash the kernel.
- Can use floating-point computation.
- Potentially higher performance, especially for memory-mapped devices, thanks to the avoidance of system calls.

► Drawbacks

- The kernel has no longer access to the device.
- None of the standard applications will be able to use it.
- Cannot use any hardware abstraction or software helpers from the kernel
- Need to adapt applications when changing the hardware.
- Less straightforward to handle interrupts: increased latency.




- ▶ Explore kernel sources manually
- ▶ Use automated tools to explore the source code



Linux Kernel Usage

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Kernel configuration



Kernel configuration

- ▶ The kernel contains thousands of device drivers, filesystem drivers, network protocols and other configurable items
- ▶ Thousands of options are available, that are used to selectively compile parts of the kernel source code
- ▶ The kernel configuration is the process of defining the set of options with which you want your kernel to be compiled
- ▶ The set of options depends
 - On the target architecture and on your hardware (for device drivers, etc.)
 - On the capabilities you would like to give to your kernel (network capabilities, filesystems, real-time, etc.). Such generic options are available in all architectures.



Kernel configuration and build system

- ▶ The kernel configuration and build system is based on multiple Makefiles
- ▶ One only interacts with the main `Makefile`, present at the **top directory** of the kernel source tree
- ▶ Interaction takes place
 - using the `make` tool, which parses the Makefile
 - through various **targets**, defining which action should be done (configuration, compilation, installation, etc.).
 - Run `make help` to see all available targets.
- ▶ Example
 - `cd linux/`
 - `make <target>`



Specifying the target architecture

First, specify the architecture for the kernel to build

- ▶ Set ARCH to the name of a directory under [arch/](#):
ARCH=arm or ARCH=arm64 or ARCH=riscv, etc
- ▶ By default, the kernel build system assumes that the kernel is configured and built for the host architecture (x86 in our case, native kernel compiling)
- ▶ The kernel build system will use this setting to:
 - Use the configuration options for the target architecture.
 - Compile the kernel with source code and headers for the target architecture.



Choosing a compiler

The compiler invoked by the kernel Makefile is `$(CROSS_COMPILE)gcc`

- ▶ Specifying the compiler is already needed at configuration time, as some kernel configuration options depend on the capabilities of the compiler.
- ▶ When compiling natively
 - Leave `CROSS_COMPILE` undefined and the kernel will be natively compiled for the host architecture using `gcc`.
- ▶ When using a cross-compiler
 - Specify the prefix of your cross-compiler executable, for example for `arm-linux-gnueabi-gcc`:
`CROSS_COMPILE=arm-linux-gnueabi-`

Set `LLVM` to `1` to compile your kernel with Clang.

See our [LLVM tools for the Linux kernel](#) presentation.



Specifying ARCH and CROSS_COMPILE

There are actually two ways of defining ARCH and CROSS_COMPILE:

- ▶ Pass ARCH and CROSS_COMPILE on the make command line:

```
make ARCH=arm CROSS_COMPILE=arm-linux- ...
```

Drawback: it is easy to forget to pass these variables when you run any make command, causing your build and configuration to be screwed up.

- ▶ Define ARCH and CROSS_COMPILE as environment variables:

```
export ARCH=arm
```

```
export CROSS_COMPILE=arm-linux-
```

Drawback: it only works inside the current shell or terminal. You could put these settings in a file that you source every time you start working on the project, see also the <https://direnv.net/> project.



Initial configuration

Difficult to find which kernel configuration will work with your hardware and root filesystem. Start with one that works!

▶ Desktop or server case:

- Advisable to start with the configuration of your running kernel:

```
cp /boot/config-`uname -r` .config
```

▶ Embedded platform case:

- Default configurations stored in-tree as minimal configuration files (only listing settings that are different with the defaults) in `arch/<arch>/configs/`
- `make help` will list the available configurations for your platform
- To load a default configuration file, just run `make foo_defconfig` (will erase your current `.config`!)
 - On ARM 32-bit, there is usually one default configuration per CPU family
 - On ARM 64-bit, there is only one big default configuration to customize



Create your own default configuration

- ▶ Use a tool such as `make menuconfig` to make changes to the configuration
- ▶ Saving your changes will overwrite your `.config` (not tracked by Git)
- ▶ When happy with it, create your own default configuration file:
 - Create a minimal configuration (non-default settings) file:
`make savedefconfig`
 - Save this default configuration in the right directory:
`mv defconfig arch/<arch>/configs/myown_defconfig`
 - Add this file to Git.
- ▶ This way, you can share a reference configuration inside the kernel sources and other developers can now get the same `.config` as you by running
`make myown_defconfig`
- ▶ When you use an embedded build system (Buildroot, OpenEmbedded) use its specific commands. E.g. `make linux-menuconfig` and `make linux-update-defconfig` in Buildroot.



Built-in or module?

- ▶ The **kernel image** is a **single file**, resulting from the linking of all object files that correspond to features enabled in the configuration
 - This is the file that gets loaded in memory by the bootloader
 - All built-in features are therefore available as soon as the kernel starts, at a time where no filesystem exists
- ▶ Some features (device drivers, filesystems, etc.) can however be compiled as **modules**
 - These are *plugins* that can be loaded/unloaded dynamically to add/remove features to the kernel
 - Each **module is stored as a separate file in the filesystem**, and therefore access to a filesystem is mandatory to use modules
 - This is not possible in the early boot procedure of the kernel, because no filesystem is available



Kernel option types

There are different types of options, defined in Kconfig files:

- ▶ `bool` options, they are either
 - *true* (to include the feature in the kernel) or
 - *false* (to exclude the feature from the kernel)
- ▶ `tristate` options, they are either
 - *true* (to include the feature in the kernel image) or
 - *module* (to include the feature as a kernel module) or
 - *false* (to exclude the feature)
- ▶ `int` options, to specify integer values
- ▶ `hex` options, to specify hexadecimal values
Example: `CONFIG_PAGE_OFFSET=0xC0000000`
- ▶ `string` options, to specify string values
Example: `CONFIG_LOCALVERSION=-no-network`

Useful to distinguish between two kernels built from different options



Kernel option dependencies

Enabling a network driver requires the network stack to be enabled, therefore configuration symbols have two ways to express dependencies:

▶ **depends on dependency:**

```
config B
    depends on A
```

- B is not visible until A is enabled
- Works well for dependency chains

▶ **select dependency:**

```
config A
    select B
```

- When A is enabled, B is enabled too (and cannot be disabled manually)
- Should preferably not select symbols with `depends on dependencies`
- Used to declare hardware features or select libraries

```
config SPI_ATH79
    tristate "Atheros AR71XX/AR724X/AR913X SPI controller driver"
    depends on ATH79 || COMPILE_TEST
    select SPI_BITBANG
    help
        This enables support for the SPI controller present on the
        Atheros AR71XX/AR724X/AR913X SoCs.
```



Kernel configuration details

- ▶ The configuration is stored in the `.config` file at the root of kernel sources
 - Simple text file, `CONFIG_PARAM=value`
 - Options are grouped by sections and are prefixed with `CONFIG_`
 - "No" value is encoded as
`# CONFIG_FOO is not set`
 - Included by the top-level kernel Makefile
 - Typically not edited by hand because of the dependencies


```
#
# CD-ROM/DVD Filesystems
#
CONFIG_ISO9660_FS=m
CONFIG_JOLIET=y
CONFIG_ZISOFS=y
CONFIG_UDF_FS=y
# end of CD-ROM/DVD Filesystems

#
# DOS/FAT/EXFAT/NT Filesystems
#
CONFIG_FAT_FS=y
CONFIG_MSDOS_FS=y
# CONFIG_VFAT_FS is not set
CONFIG_FAT_DEFAULT_CODEPAGE=437
# CONFIG_EXFAT_FS is not set
```



make xconfig

- ▶ A graphical interface to configure the kernel.
- ▶ File browser: easy to load configuration files
- ▶ Search interface to look for parameters (`[Ctrl] + [f]`)
- ▶ Required Debian/Ubuntu packages:
`qtbse5-dev` on Ubuntu 22.04





menuconfig

make menuconfig

- ▶ Useful when no graphics are available. Very efficient interface.
- ▶ Same interface found in other tools: BusyBox, Buildroot...
- ▶ Convenient number shortcuts to jump directly to search results.
- ▶ Required Debian/Ubuntu packages: libncurses-dev
- ▶ Alternative: make nconfig (now also has the number shortcuts)

```
Linux/arm 5.11.0 Kernel Configuration

Arrow keys navigate the menu. <Enter> selects submenus ---> (or empty submenus ----).
Highlighted letters are hotkeys. Pressing <Y> includes, <N> excludes, <M> modularizes
features. Press <Esc><Esc> to exit, <?> for Help, </> for Search. Legend: [*]
built-in [ ] excluded <M> module <?> module capable

General setup --->
(8) Maximum PAGE_SIZE order of alignment for DMA IOMMU buffers
System Type --->
Bus support --->
Kernel Features --->
Boot options --->
CPU Power Management --->
Floating point emulation --->
Power management options --->
Firmware Drivers --->
[*] ARM Accelerated Cryptographic Algorithms --->
General architecture-dependent options --->
[*] Enable loadable module support --->
[*] Enable the block layer --->
IO Schedulers --->
Executable file formats --->
** Memory Management options --->
[*] Networking support --->
Device Drivers --->
File systems --->
Security options --->
-* Cryptographic API --->
Library routines --->
Kernel hacking --->

<select> < Exit > < Help > < Save > < Load >
```



Kernel configuration options

You can switch from one tool to another, they all load/save the same `.config` file, and show the same set of options

Compiled as a module:

`CONFIG_ISO9660_FS=m`

Additional driver options:

`CONFIG_JOLIET=y`

`CONFIG_ZISOFS=y`

Statically built:

`CONFIG_UDF_FS=y`

☒ ISO 9660 CDROM file system support
☒ Microsoft Joliet CDROM extensions
☒ Transparent decompression extension
☒ UDF file system support

```
<M> ISO 9660 CDROM file system support
[*]  Microsoft Joliet CDROM extensions
[*]  Transparent decompression extension
<*> UDF file system support
```

Values in resulting `.config` file

Parameter values as displayed by `xconfig`

Parameter values as displayed by `menuconfig`



make oldconfig

`make oldconfig`

- ▶ Useful to upgrade a `.config` file from an earlier kernel release
- ▶ Asks for values for new parameters.
- ▶ ... unlike `make menuconfig` and `make xconfig` which silently set default values for new parameters.

If you edit a `.config` file by hand, it's useful to run `make oldconfig` afterwards, to set values to new parameters that could have appeared because of dependency changes.



Undoing configuration changes

A frequent problem:

- ▶ After changing several kernel configuration settings, your kernel no longer works.
- ▶ If you don't remember all the changes you made, you can get back to your previous configuration:

```
$ cp .config.old .config
```
- ▶ All the configuration tools keep this `.config.old` backup copy.



Compiling and installing the kernel




Kernel compilation

make


- ▶ Only works from the top kernel source directory
- ▶ Should not be performed as a privileged user
- ▶ Run several jobs in parallel. Our advice: `$(nproc)` to fully load the CPU and I/Os at all times.
Example: `make -j20`
- ▶ To **recompile** faster (7x according to some benchmarks), use the `ccache` compiler cache:
`export CROSS_COMPILE="ccache arm-linux-"`

Benefits of parallel compile jobs (make -j<n>)

Tests on Linux 5.11 on arm
`make allnoconfig configuration`
gnome-system-monitor showing the load on 4 threads / 2 CPUs



Command: `make`
Total time: 129 s



Command: `make -j8`
Total time: 67 s



Kernel compilation results

- ▶ `arch/<arch>/boot/Image`, uncompressed kernel image that can be booted
- ▶ `arch/<arch>/boot/*Image*`, compressed kernel images that can also be booted
 - `bzImage` for x86, `zImage` for ARM, `Image.gz` for RISC-V, `vmlinux.bin.gz` for ARC, etc.
- ▶ `arch/<arch>/boot/dts/<vendor>/*.dtb`, compiled Device Tree Blobs
- ▶ All kernel modules, spread over the kernel source tree, as `.ko` (*Kernel Object*) files.
- ▶ `vmlinux`, a raw uncompressed kernel image in the ELF format, useful for debugging purposes but generally not used for booting purposes



Kernel installation: native case

- ▶ `sudo make install`
 - Does the installation for the host system by default
- ▶ Installs
 - `/boot/vmlinuz-<version>`
Compressed kernel image. Same as the one in `arch/<arch>/boot`
 - `/boot/System.map-<version>`
Stores kernel symbol addresses for debugging purposes (obsolete: such information is usually stored in the kernel itself)
 - `/boot/config-<version>`
Kernel configuration for this version
- ▶ In GNU/Linux distributions, typically re-runs the bootloader configuration utility to make the new kernel available at the next boot.



Kernel installation: embedded case

- ▶ `make install` is rarely used in embedded development, as the kernel image is a single file, easy to handle.
- ▶ Another reason is that there is no standard way to deploy and use the kernel image.
- ▶ Therefore making the kernel image available to the target is usually manual or done through scripts in build systems.
- ▶ It is however possible to customize the `make install` behavior in `arch/<arch>/boot/install.sh`



Module installation: native case

- ▶ `sudo make modules_install`
 - Does the installation for the host system by default, so needs to be run as root
- ▶ Installs all modules in `/lib/modules/<version>/`
 - `kernel/`
Module `.ko` (Kernel Object) files, in the same directory structure as in the sources.
 - `modules.alias`, `modules.alias.bin`
Aliases for module loading utilities , see next slide
 - `modules.dep`, `modules.dep.bin`
Module dependencies. Kernel modules can depend on other modules, based on the symbols (functions and data structures) they use.
 - `modules.symbols`, `modules.symbols.bin`
Tells which module a given symbol belongs to (related to module dependencies).
 - `modules.builtin`
List of built-in modules of the kernel.



Module alias: *modules.alias*

Kernel compiling

```
static const struct usb_device_id  products[] = {
{
    // Linksys USB200M
    USB_DEVICE(0x077b, 0x2226),
    .driver_info = (unsigned long) &ax8817x_info,
}, {
    // Netgear FA120
    USB_DEVICE(0x0846, 0x1040),
    .driver_info = (unsigned long) &netgear_fa120_info,
},
...
{ },    // END
};
MODULE_DEVICE_TABLE(usb, products);

drivers/net/usb/asix_devices.c
```

The device driver source code lists which devices it supports

make modules


Module file
asix.ko

Containing the list of supported devices
(module metadata)

make modules_install
(depmod)

```
alias usb:v077bp2226d*dc*dsc*dp*ic*isc*ip*in* asix
alias usb:v0846p1040d*dc*dsc*dp*ic*isc*ip*in* asix
....
modules.alias
```

System operation





Module installation: embedded case

- ▶ In embedded development, you can't directly use `make modules_install` as it would install target modules in `/lib/modules` on the host!
- ▶ The `INSTALL_MOD_PATH` variable is needed to generate the module related files and install the modules in the target root filesystem instead of your host root filesystem (no need to be root):

```
make INSTALL_MOD_PATH=<dir>/ modules_install
```



Kernel cleanup targets

► From make help:

Cleaning targets:


- | | |
|------------------------|--|
| <code>clean</code> | - Remove most generated files but keep the config and enough build support to build external modules |
| <code>mrproper</code> | - Remove all generated files + config + various backup files |
| <code>distclean</code> | - <code>mrproper</code> + remove editor backup and patch files |

- If you are in a git tree, remove all files not tracked (and ignored) by git:
`git clean -fdx`





Kernel building overview





Booting the kernel



Hardware description

- ▶ Many embedded architectures have a lot of non-discoverable hardware (serial, Ethernet, I2C, Nand flash, USB controllers...)
- ▶ This hardware needs to be described and passed to the Linux kernel.
- ▶ The bootloader/firmware is expected to provide this description when starting the kernel:
 - On x86: using ACPI tables
 - On most embedded devices: using an OpenFirmware Device Tree (DT)
- ▶ This way, a kernel supporting different SoCs knows which SoC and device initialization hooks to run on the current board.



Booting with U-Boot

- ▶ On ARM32, U-Boot can boot zImage (bootz command)
- ▶ On ARM64 or RISC-V, it boots the Image file (booti command)
- ▶ In addition to the kernel image, U-Boot should also pass a DTB to the kernel.
- ▶ The typical boot process is therefore:
 1. Load zImage at address X in memory
 2. Load <board>.dtb at address Y in memory
 3. Start the kernel with `boot[z|i] X - Y`
The - in the middle indicates no *initramfs*



Kernel command line


- ▶ In addition to the compile time configuration, the kernel behavior can be adjusted with no recompilation using the **kernel command line**
- ▶ The kernel command line is a string that defines various arguments to the kernel
 - It is very important for system configuration
 - `root=` for the root filesystem (covered later)
 - `console=` for the destination of kernel messages
 - Example: `console=ttyS0 root=/dev/mmcblk0p2 rootwait`
 - Many more exist. The most important ones are documented in [admin-guide/kernel-parameters](#) in kernel documentation.



Passing the kernel command line

- ▶ U-Boot carries the Linux kernel command line string in its `bootargs` environment variable
- ▶ Right before starting the kernel, it will store the contents of `bootargs` in the chosen section of the Device Tree
- ▶ The kernel will behave differently depending on its configuration:
 - If `CONFIG_CMDLINE_FROM_BOOTLOADER` is set:
The kernel will use only the string from the bootloader
 - If `CONFIG_CMDLINE_FORCE` is set:
The kernel will only use the string received at configuration time in `CONFIG_CMDLINE`
 - If `CONFIG_CMDLINE_EXTEND` is set:
The kernel will concatenate both strings

See the "Understanding U-Boot Falcon Mode" presentation from Michael Opdenacker, for details about how U-Boot boots Linux.



Slides: <https://bootlin.com/pub/conferences/2021/lee/>
Video: <https://www.youtube.com/watch?v=LFe3x2QMhSo>



Kernel log

- ▶ The kernel keeps its messages in a circular buffer in memory
 - The size is configurable using `CONFIG_LOG_BUF_SHIFT`
- ▶ When a module is loaded, related information is available in the kernel log.
- ▶ Kernel log messages are available through the `dmesg` command (**d**iagnostics **m**essage)
- ▶ Kernel log messages are also displayed on the console pointed by the `console=` kernel command line argument
 - Console messages can be filtered by level using the `loglevel` parameter
 - Example: `console=ttyS0 loglevel=5`
- ▶ It is possible to write to the kernel log from user space:
`echo "<n>Debug info" > /dev/kmsg`



Practical lab - Kernel compiling and booting



1st lab: board and bootloader setup:

- ▶ Prepare the board and access its serial port
- ▶ Configure its bootloader to use TFTP

2nd lab: kernel compiling and booting:

- ▶ Set up a cross-compiling environment
- ▶ Cross-compile a kernel for an ARM target platform
- ▶ Boot this kernel from a directory on your workstation, accessed by the board through NFS




Using kernel modules



Advantages of modules

- ▶ Modules make it easy to develop drivers without rebooting: load, test, unload, rebuild, load...
- ▶ Useful to keep the kernel image size to the minimum (essential in GNU/Linux distributions for PCs).
- ▶ Also useful to reduce boot time: you don't spend time initializing devices and kernel features that you only need later.
- ▶ Caution: once loaded, have full control and privileges in the system. No particular protection. That's why only the `root` user can load and unload modules.
- ▶ To increase security, possibility to allow only signed modules, or to disable module support entirely.

Using kernel modules to support many different devices and setups



The modules in the initramfs are updated every time a kernel upgrade is available.



Module utilities: extracting information

`<module_name>`: name of the module file without the trailing `.ko`

▶ `modinfo <module_name>` (for modules in `/lib/modules`)

`modinfo <module_path>.ko`

Gets information about a module without loading it: parameters, license, description and dependencies.



Module utilities: loading

- ▶ `sudo insmod <module_path>.ko`
Tries to load the given module. The full path to the module object file must be given.
- ▶ `sudo modprobe <top_module_name>`
Most common usage of `modprobe`: tries to load all the dependencies of the given top module, and then this module. Lots of other options are available. `modprobe` automatically looks in `/lib/modules/<version>/` for the object file corresponding to the given module name.
- ▶ `lsmod`
Displays the list of loaded modules
Compare its output with the contents of `/proc/modules!`



Understanding module loading issues

- ▶ When loading a module fails, `insmod` often doesn't give you enough details!
- ▶ Details are often available in the kernel log.
- ▶ Example:

```
$ sudo insmod ./intr_monitor.ko
insmod: error inserting './intr_monitor.ko': -1 Device or resource busy
$ dmesg
[17549774.552000] Failed to register handler for irq channel 2
```




Module utilities: removals

- ▶ `sudo rmmod <module_name>`

Tries to remove the given module.

Will only be allowed if the module is no longer in use (for example, no more processes opening a device file)

- ▶ `sudo modprobe -r <top_module_name>`

Tries to remove the given top module and all its no longer needed dependencies



Passing parameters to modules

- ▶ Find available parameters:

```
modinfo usb-storage
```

- ▶ Through `insmod`:

```
sudo insmod ./usb-storage.ko delay_use=0
```

- ▶ Through `modprobe`:

Set parameters in `/etc/modprobe.conf` or in any file in `/etc/modprobe.d/`:

```
options usb-storage delay_use=0
```

- ▶ Through the kernel command line, when the module is built statically into the kernel:

```
usb-storage.delay_use=0
```

- `usb-storage` is the *module name*
- `delay_use` is the *module parameter name*. It specifies a delay before accessing a USB storage device (useful for rotating devices).
- `0` is the *module parameter value*



Check module parameter values

How to find/edit the current values for the parameters of a loaded module?


- ▶ Check `/sys/module/<name>/parameters`.
- ▶ There is one file per parameter, containing the parameter value.
- ▶ Also possible to change parameter values if these files have write permissions (depends on the module code).
- ▶ Example:

```
echo 0 > /sys/module/usb_storage/parameters/delay_use
```



Developing kernel modules

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Hello module 1/2

```
// SPDX-License-Identifier: GPL-2.0
/* hello.c */
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>

static int __init hello_init(void)
{
    pr_alert("Good morrow to this fair assembly.\n");
    return 0;
}

static void __exit hello_exit(void)
{
    pr_alert("Alas, poor world, what treasure hast thou lost!\n");
}

module_init(hello_init);
module_exit(hello_exit);
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Greeting module");
MODULE_AUTHOR("William Shakespeare");
```



Hello module 2/2

▶ Code marked as `__init`:

- Removed after initialization (static kernel or module.)
- See how init memory is reclaimed when the kernel finishes booting:

```
[ 2.689854] VFS: Mounted root (nfs filesystem) on device 0:15.  
[ 2.698796] devtmpfs: mounted  
[ 2.704277] Freeing unused kernel memory: 1024K  
[ 2.710136] Run /sbin/init as init process
```

▶ Code marked as `__exit`:

- Discarded when module compiled statically into the kernel, or when module unloading support is not enabled.

▶ Code of this example module available on

<https://raw.githubusercontent.com/bootlin/training-materials/master/code/hello/hello.c>



Hello module explanations

- ▶ Headers specific to the Linux kernel: `linux/xxx.h`
 - No access to the usual C library, we're doing kernel programming
- ▶ An initialization function
 - Called when the module is loaded, returns an error code (0 on success, negative value on failure)
 - Declared by the `module_init()` macro: the name of the function doesn't matter, even though `<modulename>_init()` is a convention.
- ▶ A cleanup function
 - Called when the module is unloaded
 - Declared by the `module_exit()` macro.
- ▶ Metadata information declared using `MODULE_LICENSE()`, `MODULE_DESCRIPTION()` and `MODULE_AUTHOR()`







Symbols exported to modules 1/2

- ▶ From a kernel module, only a limited number of kernel functions can be called
- ▶ Functions and variables have to be explicitly exported by the kernel to be visible to a kernel module
- ▶ Two macros are used in the kernel to export functions and variables:
 - `EXPORT_SYMBOL(symbolname)`, which exports a function or variable to all modules
 - `EXPORT_SYMBOL_GPL(symbolname)`, which exports a function or variable only to GPL modules
 - Linux 5.3: contains the same number of symbols with `EXPORT_SYMBOL()` and symbols with `EXPORT_SYMBOL_GPL()`
- ▶ A normal driver should not need any non-exported function.



Symbols exported to modules 2/2





► Several usages

- Used to restrict the kernel functions that the module can use if it isn't a GPL licensed module.
 - Difference between `EXPORT_SYMBOL()` and `EXPORT_SYMBOL_GPL()`.
- One reason a kernel can become “tainted” is proprietary modules, among others.
 - See [admin-guide/tainted-kernels](#) for other taint flags.
 - This attribute is visible in kernel crashes and oopses for bug reports.
- Useful for users to check that their system is 100% free (for the kernel, check `/proc/sys/kernel/tainted`; run `vrms` to check installed packages).

► Values

- GPL compatible (see [include/linux/license.h](#): GPL, GPL v2, GPL and additional rights, Dual MIT/GPL, Dual BSD/GPL, Dual MPL/GPL)
- Proprietary



Two solutions

- ▶ *Out of tree*, when the code is outside of the kernel source tree, in a different directory
 - Not integrated into the kernel configuration/compilation process
 - Needs to be built separately
 - The driver cannot be built statically, only as a module
- ▶ *Inside the kernel tree*
 - Well integrated into the kernel configuration/compilation process
 - The driver can be built statically or as a module



Compiling an out-of-tree module 1/2

- ▶ The below Makefile should be reusable for any single-file out-of-tree Linux module
- ▶ The source file is `hello.c`
- ▶ Just run `make` to build the `hello.ko` file


```
ifneq ($(KERNELRELEASE),)
obj-m := hello.o
else
KDIR := /path/to/kernel/sources

all:
<tab>$(MAKE) -C $(KDIR) M=$$PWD
endif
```

- ▶ KDIR: kernel source or headers directory (see next slides)



Compiling an out-of-tree module 2/2



- ▶ The module Makefile is interpreted with `KERNELRELEASE` undefined, so it calls the kernel Makefile, passing the module directory in the `M` variable
- ▶ The kernel Makefile knows how to compile a module, and thanks to the `M` variable, knows where the Makefile for our module is. This module Makefile is then interpreted with `KERNELRELEASE` defined, so the kernel sees the `obj-m` definition.



Modules and kernel version

- ▶ To be compiled, a kernel module needs access to *kernel headers*, containing the definitions of functions, types and constants.
- ▶ Two solutions
 - Full kernel sources (configured + `make modules_prepare`)
 - Only kernel headers (`linux-headers-*` packages in Debian/Ubuntu distributions, or directory created by `make headers_install`).
- ▶ The sources or headers must be configured (`.config` file)
 - Many macros or functions depend on the configuration
- ▶ You also need the kernel [Makefile](#), the [scripts/](#) directory, and a few others.
- ▶ A kernel module compiled against version X of kernel headers will not load in kernel version Y
 - `modprobe` / `insmod` will say `Invalid module format`



New driver in kernel sources 1/2

- ▶ To add a new driver to the kernel sources:
 - Add your new source file to the appropriate source directory. Example: [drivers/usb/serial/navman.c](#)
 - Single file drivers in the common case, even if the file is several thousand lines of code big. Only really big drivers are split in several files or have their own directory.
 - Describe the configuration interface for your new driver by adding the following lines to the `Kconfig` file in this directory:

```
config USB_SERIAL_NAVMAN
    tristate "USB Navman GPS device"
    depends on USB_SERIAL
    help
        To compile this driver as a module, choose M
        here: the module will be called navman.
```



New driver in kernel sources 2/2

- ▶ Add a line in the Makefile file based on the Kconfig setting:
`obj-$(CONFIG_USB_SERIAL_NAVMAN) += navman.o`
- ▶ It tells the kernel build system to build `navman.c` when the `USB_SERIAL_NAVMAN` option is enabled. It works both if compiled statically or as a module.
 - Run `make xconfig` and see your new options!
 - Run `make` and your new files are compiled!
 - See [kbuild/](#) for details and more elaborate examples like drivers with several source files, or drivers in their own subdirectory, etc.



Hello module with parameters 1/2

```
// SPDX-License-Identifier: GPL-2.0
/* hello_param.c */
#include <linux/init.h>
#include <linux/module.h>

MODULE_LICENSE("GPL");

static char *whom = "world";
module_param(whom, charp, 0644);
MODULE_PARM_DESC(whom, "Recipient of the hello message");

static int howmany = 1;
module_param(howmany, int, 0644);
MODULE_PARM_DESC(howmany, "Number of greetings");
```



Hello module with parameters 2/2

```
static int __init hello_init(void)
{
    int i;

    for (i = 0; i < howmany; i++)
        pr_alert("(%d) Hello, %s\n", i, whom);
    return 0;
}

static void __exit hello_exit(void)
{
    pr_alert("Goodbye, cruel %s\n", whom);
}

module_init(hello_init);
module_exit(hello_exit);
```

Thanks to Jonathan Corbet for the examples

Source code available on: https://github.com/bootlin/training-materials/blob/master/code/hello-param/hello_param.c



Declaring a module parameter

```
module_param(  
    name, /* name of an already defined variable */  
    type, /* standard types (different from C types) are:  
        * byte, short, ushort, int, uint, long, ulong  
        * charp: a character pointer  
        * bool: a bool, values 0/1, y/n, Y/N.  
        * invbool: the above, only sense-reversed (N = true). */  
    perm /* for /sys/module/<module_name>/parameters/<param>,  
        * 0: no such module parameter value file */  
);
```

```
/* Example: drivers/block/loop.c */  
static int max_loop;  
module_param(max_loop, int, 0444);  
MODULE_PARM_DESC(max_loop, "Maximum number of loop devices");
```

Modules parameter arrays are also possible with `module_param_array()`.




- ▶ Create, compile and load your first module
- ▶ Add module parameters
- ▶ Access kernel internals from your module



Describing hardware devices

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Discoverable hardware: USB and PCI



Discoverable hardware

- ▶ Some busses have built-in hardware discoverability mechanisms
- ▶ Most common busses: USB and PCI
- ▶ Hardware devices can be enumerated, and their characteristics retrieved with just a driver or the bus controller
- ▶ Useful Linux commands
 - `lsusb`, lists all USB devices detected
 - `lspci`, lists all PCI devices detected
 - A detected device does not mean it has a kernel driver associated to it!
- ▶ Association with kernel drivers done based on product ID/vendor ID, or some other characteristics of the device: device class, device sub-class, etc.



Describing non-discoverable hardware



Describing non-discoverable hardware

1. Directly in the **OS/bootloader code**

- ▶ Using compiled data structures, typically in C
- ▶ How it was done on most embedded platforms in Linux, U-Boot.
- ▶ Considered not maintainable/sustainable on ARM32, which motivated the move to another solution.



Describing non-discoverable hardware

2. Using **ACPI** tables

- ▶ On *x86* systems, but also on a subset of ARM64 platforms
- ▶ Tables provided by the firmware




3. Using a **Device Tree**

- ▶ Originates from **OpenFirmware**, defined by Sun, used on SPARC and PowerPC
 - That's why many Linux/U-Boot functions related to DT have a `of_` prefix
- ▶ Now used by most embedded-oriented CPU architectures that run Linux: ARC, ARM64, RISC-V, ARM32, PowerPC, Xtensa, MIPS, etc.
- ▶ Writing/tweaking a DT is necessary when porting Linux to a new board, or when connecting additional peripherals



Device Tree: from source to blob

- ▶ A tree data structure describing the hardware is written by a developer in a **Device Tree Source** file, `.dts`
- ▶ Processed by the **Device Tree Compiler**, `dtc`
- ▶ Produces a more efficient representation: **Device Tree Blob**, `.dtb`
- ▶ Additional C preprocessor pass
- ▶ `.dtb` → accurately describes the hardware platform in an **OS-agnostic** way.
- ▶ `.dtb` $\approx$ few dozens of kilobytes
- ▶ DTB also called **FDT**, *Flattened Device Tree*, once loaded into memory.
 - `fdt` command in U-Boot
 - `fdt_` APIs





dts example

```
$ cat foo.dts
/dts-v1/;

/ {
    welcome = <0xBADCAFE>;
    bootlin {
        webinar = "great";
        demo = <1>, <2>, <3>;
    };
};
```



dts example

```
$ cat foo.dts
/dts-v1/;

/ {
    welcome = <0xBADCAFE>;
    bootlin {
        webinar = "great";
        demo = <1>, <2>, <3>;
    };
};
```

```
$ dtc -I dts -O dtb -o foo.dtb foo.dts
$ ls -l foo.dt*
-rw-r--r-- 1 thomas thomas 169 ... foo.dtb
-rw-r--r-- 1 thomas thomas 102 ... foo.dts
```



dtc example

```
$ cat foo.dts
/dts-v1/;

/ {
    welcome = <0xBADCAFE>;
    bootlin {
        webinar = "great";
        demo = <1>, <2>, <3>;
    };
};
```

```
$ dtc -I dts -O dtb -o foo.dtb foo.dts
$ ls -l foo.dt*
-rw-r--r-- 1 thomas thomas 169 ... foo.dtb
-rw-r--r-- 1 thomas thomas 102 ... foo.dts
```

```
$ dtc -I dtb -O dts foo.dtb
/dts-v1/;

/ {
    welcome = <0xbadcafe>;

    bootlin {
        webinar = "great";
        demo = <0x01 0x02 0x03>;
    };
};
```



Where are Device Tree Sources located?

- ▶ Even though they are OS-agnostic, **no central and OS-neutral** place to host Device Tree sources and share them between projects
 - Often discussed, never done
- ▶ In practice, the Linux kernel sources can be considered as the **canonical location** for Device Tree Source files
 - `arch/<ARCH>/boot/dts/<vendor>/`
 - `arch/arm/boot/dts` (on ARM 32 architecture before Linux 6.5)
 - $\approx$ 4500 Device Tree Source files (`.dts` and `.dtsi`) in Linux as of 6.0.
- ▶ Duplicated/synced in various projects
 - U-Boot, Barebox, TF-A



Device Tree base syntax

- ▶ Tree of **nodes**
- ▶ Nodes with **properties**
- ▶ Node $\approx$ a device or IP block
- ▶ Properties $\approx$ device characteristics
- ▶ Notion of **cells** in property values
- ▶ Notion of **phandle** to point to other nodes
- ▶ dtc only does syntax checking, no semantic validation

```
/ {  
    node@0 {  
        a-string-property = "A string";  
        a-string-list-property = "first string", "second string";  
        a-byte-data-property = [0x01 0x23 0x34 0x56];  
  
        child-node@0 {  
            first-child-property;  
            second-child-property = <1>;  
            a-reference-to-something = <&node1>;  
        };  
  
        child-node@1 {  
        };  
    };  
    node1: node@1 {  
        an-empty-property;  
        a-cell-property = <1 2 3 4>;  
  
        child-node@0 {  
        };  
    };  
};
```


Diagram illustrating the Device Tree base syntax with annotations:

- Node name**: Points to `node@0`.
- Unit address**: Points to `@0`.
- Property name**: Points to `a-string-property`.
- Property value**: Points to the value of `a-string-property`.
- Properties of node@0**: Points to the list of properties under `node@0`.
- Bytestring**: Points to the value of `a-byte-data-property`.
- A handle (reference to another node)**: Points to the value of `a-reference-to-something`.
- Label**: Points to the label `node1`.
- Four cells (32 bits values)**: Points to the value of `a-cell-property`.



DT overall structure: simplified example

```
/ {  
  #address-cells = <1>;  
  #size-cells = <1>;  
  model = "TI AM335x BeagleBone Black";  
  compatible = "ti,am335x-bone-black", "ti,am335x-bone", "ti,am33xx";  
  
  cpus { ... };  
  memory@80000000 { ... };  
  chosen { ... };  
  ocp {  
    intc: interrupt-controller@48200000 { ... };  
    usb0: usb@47401300 { ... };  
    l4_per: interconnect@44c00000 {  
      i2c0: i2c@40012000 { ... };  
    };  
  };  
};
```






DT overall structure: simplified example

```
/ {
  cpus {
    #address-cells = <1>;
    #size-cells = <0>;
    cpu0: cpu@0 {
      compatible = "arm,cortex-a8";
      enable-method = "ti,am3352";
      device_type = "cpu";
      reg = <0>;
    };
  };

  memory@0x80000000 {
    device_type = "memory";
    reg = <0x80000000 0x10000000>; /* 256 MB */
  };

  chosen {
    bootargs = "";
    stdout-path = &uart0;
  };

  ocp { ... };
};
```






DT overall structure: simplified example

```
/ {
  cpus { ... };
  memory@0x80000000 { ... };
  chosen { ... };
  ocp {

    intc: interrupt-controller@48200000 {
      compatible = "ti,am33xx-intc";
      interrupt-controller;
      #interrupt-cells = <1>;
      reg = <0x48200000 0x1000>;
    };

    usb0: usb@47401300 {
      compatible = "ti,musb-am33xx";
      reg = <0x1400 0x400>, <0x1000 0x200>;
      reg-names = "mc", "control";
      interrupts = <18>;
      dr_mode = "otg";
      dmas = <&cppi41dma 0 0 &cppi41dma 1 0 ...>;
      status = "okay";
    };

    l4_per: interconnect@44c00000 {
      i2c0: i2c@40012000 { ... };
    };
  };
};
```





DT overall structure: simplified example


```
/ {
  cpus { ... };
  memory@0x80000000 { ... };
  chosen { ... };

  ocp {
    compatible = "simple-pm-bus";
    clocks = <&l3_clkctrl AM3_L3_L3_MAIN_CLKCTRL 0>;
    clock-names = "fck";
    #address-cells = <1>;
    #size-cells = <1>;

    intc: interrupt-controller@48200000 { ... };
    usb0: usb@47401300 { ... };

    l4_per: interconnect@44c00000 {
      compatible = "ti,am33xx-l4-wkup", "simple-pm-bus";
      reg = <0x44c00000 0x800>, <0x44c00800 0x800>,
            <0x44c01000 0x400>, <0x44c01400 0x400>;
      reg-names = "ap", "la", "ia0", "ia1";
      #address-cells = <1>;
      #size-cells = <1>;


      i2c0: i2c@40012000 { ... };
    };
  };
};
```





DT overall structure: simplified example

```
/ {  
  cpus { ... };  
  memory@0x80000000 { ... };  
  chosen { ... };  
  ocp {  
    intc: interrupt-controller@48200000 { ... };  
    usb0: usb@47401300 { ... };  
    l4_per: interconnect@44c00000 {  
  
      i2c0: i2c@40012000 {  
        compatible = "ti,omap4-i2c";  
        #address-cells = <1>;  
        #size-cells = <0>;  
        reg = <0x0 0x1000>;  
        interrupts = <70>;  
        status = "okay";  
        pinctrl-names = "default";  
        pinctrl-0 = <&i2c0_pins>;  
        clock-frequency = <400000>;  
  
        baseboard_eeprom: eeprom@50 {  
          compatible = "atmel,24c256";  
          reg = <0x50>;  
        };  
      };  
    };  
  };  
};
```





Device Tree inheritance

- ▶ Device Tree files are not monolithic, they can be split in several files, including each other.
- ▶ `.dtsi` files are included files, while `.dts` files are *final* Device Trees
 - Only `.dts` files are accepted as input to `dtc`
- ▶ Typically, `.dtsi` will contain
 - definitions of SoC-level information
 - definitions common to several boards
- ▶ The `.dts` file contains the board-level information
- ▶ The inclusion works by **overlaying** the tree of the including file over the tree of the included file, according to the order of the `#include` directives.
- ▶ Allows an including file to **override** values specified by an included file.
- ▶ Uses the C pre-processor `#include` directive



Device Tree inheritance example

Definition of the AM33xx SoC family

```
&l4_wkup {
    target-module@b000 {
        i2c0: i2c@0 {
            compatible = "ti,omap4-i2c";
            reg = <0x0 0x1000>;
            interrupts = <70>;
            status = "disabled";
        };
    };
};
```

am33xx-l4.dtsi



Definition of the Bone Black board

```
#include "am33xx-l4.dtsi"

&i2c0 {
    pinctrl-names = "default";
    pinctrl-0 = <&i2c0_pins>;
    status = "okay";

    baseboard_eeprom: eeprom@50 {
        compatible = "atmel,24c256";
        reg = <0x50>;
    };
};
```

am33xx-boneblack.dts



Compiled DTB

```
&l4_wkup {
    target-module@b000 {
        i2c0: i2c@0 {
            compatible = "ti,omap4-i2c";
            reg = <0x0 0x1000>;
            interrupts = <70>;
            pinctrl-names = "default";
            pinctrl-0 = <&i2c0_pins>;
            status = "okay";

            baseboard_eeprom: eeprom@50 {
                compatible = "atmel,24c256";
                reg = <0x50>;
            };
        };
    };
};
```

am33xx-boneblack.dtb

Note 1

The actual Device Trees for this platform are more complicated. This example is highly simplified.

Note 2

The real DTB is in binary format. Here we show the text equivalent of the DTB contents.



Inheritance and labels

Doing:

soc.dtsi

```
/ {
    ocp {
        uart0: serial@0 {
            compatible = "ti,am3352-uart", "ti,omap3-uart";
            reg = <0x0 0x1000>;
            status = "disabled";
        };
    };
};
```

board.dts

```
#include "soc.dtsi"

/ {
    ocp {
        serial@0 {
            status = "okay";
        };
    };
};
```



Inheritance and labels

Doing:

soc.dtsi

```
/ {
    ocp {
        uart0: serial@0 {
            compatible = "ti,am3352-uart", "ti,omap3-uart";
            reg = <0x0 0x1000>;
            status = "disabled";
        };
    };
};
```

board.dts

```
#include "soc.dtsi"

/ {
    ocp {
        serial@0 {
            status = "okay";
        };
    };
};
```

Is exactly equivalent to:

soc.dtsi

```
/ {
    ocp {
        uart0: serial@0 {
            compatible = "ti,am3352-uart", "ti,omap3-uart";
            reg = <0x0 0x1000>;
            status = "disabled";
        };
    };
};
```

board.dts


```
#include "soc.dtsi"

&uart0 {
    status = "okay";
};
```

→ this solution is now often preferred



DT inheritance in Bone Black support





Device Tree design principles

- ▶ **Describe hardware** (how the hardware is), not configuration (how I choose to use the hardware)
- ▶ **OS-agnostic**
 - For a given piece of HW, Device Tree should be the same for U-Boot, FreeBSD or Linux
 - There should be no need to change the Device Tree when updating the OS
- ▶ Describe **integration of hardware components**, not the internals of hardware components
 - The details of how a specific device/IP block is working is handled by code in device drivers
 - The Device Tree describes how the device/IP block is connected/integrated with the rest of the system: IRQ lines, DMA channels, clocks, reset lines, etc.
- ▶ Like all beautiful design principles, these principles are sometimes violated.



The properties

Device tree properties can:

- ▶ Be generic and apply to most nodes
 - Their meaning is usually described in one place: the core DT schema available at <https://github.com/devicetree-org/dt-schema>.
 - `compatible`, `reg`, `#address-cells`, etc
- ▶ Cover common consumer-provider relationships
 - Their meaning is either described in the [dt-schema](#) GitHub repository or under [Documentation/devicetree/bindings](#).
 - `clocks`, `interrupts`, `regulators`, etc
- ▶ Subsystem specific
 - All devices of a certain class may use them, often starting with the class name
 - `spi-cpha`, `i2c-scl-internal-delay-ns`, `nand-ecc-engine`, `mac-address`, etc
- ▶ Vendor/device specific
 - To describe uncommon or very specific properties
 - Always described in the device's binding file and prefixed with `<vendor>`,
 - `ti`, `hwmodes`, `xlnx`, `num-channels`, `nxp`, `tx-output-mode`, etc
- ▶ Some of them are deprecated, watch out the bindings!



The compatible property

- ▶ Is a list of strings
 - From the most specific to the least specific
- ▶ Describes the specific **binding** to which the node complies.
- ▶ It uniquely identifies the **programming model** of the device.
- ▶ Practically speaking, it is used by the operating system to find the **appropriate driver** for this device.
- ▶ When describing real hardware, the typical form is `vendor,model`
- ▶ Examples:
 - `compatible = "arm,armv7-timer";`
 - `compatible = "st,stm32mp1-dwmac", "snps,dwmac-4.20a";`
 - `compatible = "regulator-fixed";`
 - `compatible = "gpio-keys";`
- ▶ Special value: `simple-bus` → bus where all sub-nodes are memory-mapped devices



compatible property and Linux kernel drivers

- ▶ Linux identifies as **platform devices**:
 - Top-level DT nodes with a `compatible` string
 - Sub-nodes of `simple-bus`
 - Instantiated automatically at boot time
- ▶ Sub-nodes of I2C controllers → *I2C devices*
- ▶ Sub-nodes of SPI controllers → *SPI devices*
- ▶ Each Linux driver has a table of compatible strings it supports
 - `struct of_device_id[]`
- ▶ When a DT node compatible string matches a given driver, the device is *bound* to that driver.

```
/ {  
    timer {  
        compatible = "...";  
    };  
    soc {  
        compatible = "simple-bus";  
        uart@1000 {  
            compatible = "...";  
        };  
        i2c@2000 {  
            compatible = "...";  
            eeprom@65 {  
                compatible = "...";  
            };  
        };  
    };  
};
```

Platform device

Platform device

Platform device

I2C device



Matching with drivers in Linux: platform driver

drivers/i2c/busses/i2c-omap.c

```
static const struct of_device_id omap_i2c_of_match[] = {
    {
        .compatible = "ti,omap4-i2c",
        .data = &omap4_pdata,
    },
    {
        .compatible = "ti,omap3-i2c",
        .data = &omap3_pdata,
    },
    [...],
    { },
};
MODULE_DEVICE_TABLE(of, omap_i2c_of_match);

[...]
```

```
static struct platform_driver omap_i2c_driver = {
    .probe      = omap_i2c_probe,
    .remove     = omap_i2c_remove,
    .driver     = {
        .name    = "omap_i2c",
        .pm      = &omap_i2c_pm_ops,
        .of_match_table = of_match_ptr(omap_i2c_of_match),
    },
};
```




Matching with drivers in Linux: I2C driver

sound/soc/codecs/cs42l51.c

```
const struct of_device_id cs42l51_of_match[] = {
    { .compatible = "cirrus,cs42l51", },
    { }
};
MODULE_DEVICE_TABLE(of, cs42l51_of_match);
```

sound/soc/codecs/cs42l51-i2c.c

```
static struct i2c_driver cs42l51_i2c_driver = {
    .driver = {
        .name = "cs42l51",
        .of_match_table = cs42l51_of_match,
        .pm = &cs42l51_pm_ops,
    },
    .probe = cs42l51_i2c_probe,
    .remove = cs42l51_i2c_remove,
    .id_table = cs42l51_i2c_id,
};
```



reg property

- ▶ Most important property after `compatible`
- ▶ **Memory-mapped** devices: base physical address and size of the memory-mapped registers. Can have several entries for multiple register areas.

```
sai4: sai@50027000 {  
    reg = <0x50027000 0x4>, <0x500273f0 0x10>;  
};
```



reg property

- ▶ Most important property after `compatible`
- ▶ **Memory-mapped** devices: base physical address and size of the memory-mapped registers. Can have several entries for multiple register areas.
- ▶ **I2C** devices: address of the device on the I2C bus.

```
&i2c1 {  
    hdmi-transmitter@39 {  
        reg = <0x39>;  
    };  
    cs42l51: cs42l51@4a {  
        reg = <0x4a>;  
    };  
};
```



reg property

- ▶ Most important property after `compatible`
- ▶ **Memory-mapped** devices: base physical address and size of the memory-mapped registers. Can have several entries for multiple register areas.
- ▶ **I2C** devices: address of the device on the I2C bus.
- ▶ **SPI** devices: chip select number

```
&qspi {  
    flash0: mx66l51235l@0 {  
        reg = <0>;  
    };  
    flash1: mx66l51235l@1 {  
        reg = <1>;  
    };  
};
```



reg property

- ▶ Most important property after `compatible`
- ▶ **Memory-mapped** devices: base physical address and size of the memory-mapped registers. Can have several entries for multiple register areas.
- ▶ **I2C** devices: address of the device on the I2C bus.
- ▶ **SPI** devices: chip select number
- ▶ The unit address must be the address of the first `reg` entry.

```
sai4: sai@50027000 {  
    reg = <0x50027000 0x4>, <0x500273f0 0x10>;  
};
```



cells property

- ▶ Property numbers shall fit into 32-bit containers called `cells`
- ▶ The compiler does not maintain information about the number of entries, the OS just receives 4 independent `cells`
 - Example with a `reg` property using 2 entries of 2 `cells`:

```
reg = <0x50027000 0x4>, <0x500273f0 0x10>;
```

- The OS cannot make the difference with:

```
reg = <0x50027000>, <0x4>, <0x500273f0>, <0x10>;  
reg = <0x50027000 0x4 0x500273f0>, <0x10>;  
reg = <0x50027000>, <0x4 0x500273f0 0x10>;  
reg = <0x50027000 0x4 0x500273f0 0x10>;
```



cells property

- ▶ Property numbers shall fit into 32-bit containers called `cells`
- ▶ The compiler does not maintain information about the number of entries, the OS just receives 4 independent `cells`
- ▶ Need for other properties to declare the right formatting:
 - `#address-cells`: Indicates the number of cells used to carry the address
 - `#size-cells`: Indicates the number of cells used to carry the size of the range
- ▶ The parent-node declares the children `reg` property formatting
 - Platform devices need memory ranges

```
module@a0000 {  
    #address-cells = <1>;  
    #size-cells = <1>;  
  
    serial@1000 {  
        reg = <0x1000 0x10>, <0x2000 0x10>;  
    };  
};
```



cells property

- ▶ Property numbers shall fit into 32-bit containers called `cells`
- ▶ The compiler does not maintain information about the number of entries, the OS just receives 4 independent `cells`
- ▶ Need for other properties to declare the right formatting:
 - `#address-cells`: Indicates the number of cells used to carry the address
 - `#size-cells`: Indicates the number of cells used to carry the size of the range
- ▶ The parent-node declares the children `reg` property formatting
 - Platform devices need memory ranges
 - SPI devices need chip-selects

```
spi@300000 {  
    #address-cells = <1>;  
    #size-cells = <0>;  
  
    flash@1 {  
        reg = <1>;  
    };  
};
```




Status property

- ▶ The `status` property indicates if the device is really in use or not
 - `okay` or `ok` → the device is really in use
 - any other value, by convention `disabled` → the device is not in use
- ▶ In Linux, controls if a device is instantiated
- ▶ In `.dtsi` files describing SoCs: all devices that interface to the outside world have `status = "disabled";`
- ▶ Enabled on a per-device basis in the board `.dts`



Resources: interrupts, clocks, DMA, reset lines, ...

- ▶ Common pattern for resources shared by multiple hardware blocks
 - Interrupt lines
 - Clock controllers
 - DMA controllers
 - Reset controllers
 - ...
- ▶ A Device Tree node describing the *controller* as a device
- ▶ References from other nodes that use resources provided by this *controller*

```
intc: interrupt-controller@a0021000 {
    compatible = "arm,cortex-a7-gic";
    #interrupt-cells = <3>;
    interrupt-controller;
    reg = <0xa0021000 0x1000>, <0xa0022000 0x2000>;
};

rcc: rcc@50000000 {
    compatible = "st,stm32mp1-rcc", "syscon";
    reg = <0x50000000 0x1000>;
    #clock-cells = <1>;
    #reset-cells = <1>;
};

dmamux1: dma-router@48002000 {
    compatible = "st,stm32h7-dmamux";
    reg = <0x48002000 0x1c>;
    #dma-cells = <3>;
    clocks = <&rcc DMAMUX>;
    resets = <&rcc DMAMUX_R>;
};

spi3: spi@4000c000 {
    interrupts = <GIC_SPI 51 IRQ_TYPE_LEVEL_HIGH>;
    clocks = <&rcc SPI3_K>;
    resets = <&rcc SPI3_R>;
    dmas = <&dmamux1 61 0x400 0x05>, <&dmamux1 62 0x400 0x05>;
};
```



Generic suffixes

- ▶ xxx-gpios
 - When drivers need access to GPIOs
 - May be subsystem-specific or vendor-specific
 - Examples: enable-gpios, cts-gpios, rts-gpios
- ▶ xxx-names
 - Sometimes naming items is relevant
 - Allows drivers to perform lookups by name rather than ID
 - The order of definition of each item still matters
 - Examples: gpio-names, clock-names, reset-names

```
uart0@4000c000 {  
    dmas = <&edma 26 0>, <&edma 27 0>;  
    dma-names = "tx", "rx";  
    ...  
};
```



How to validate Device Tree content? 1/2

- ▶ `compatible` properties enforce a specific programming model
- ▶ OS expect a specific set of properties in each node
 - The syntax is fixed
 - The content is defined (number of items, their size, their meaning)
 - Some properties are mandatory
- ▶ How do I check the validity of a DT snippet?
 - How do I avoid losing half a day on a typo?
 - Looking at drivers to understand the DT structure tends to make it OS-specific



How to validate Device Tree content? 2/2

- ▶ **Device Tree Specifications** → base Device Tree syntax + number of standard properties.
 - <https://www.devicetree.org/specifications/>
 - Not sufficient to describe the wide variety of hardware.
- ▶ **Device Tree Bindings** → describes how a piece of HW should be described
 - Common bindings are defined in an external repository <https://github.com/devicetree-org/dt-schema/tree/main/dtschema/schemas>
 - Generic properties: `reg` or `#address-cells`
 - Consumer bindings: `interrupts`, `clocks`, `dmab`, etc
 - Device-specific descriptions are in the Linux kernel sources <Documentation/devicetree/bindings/>



Devicetree Specification
Release v0.3

[devicetree.org](https://www.devicetree.org)

13 February 2020



Device Tree bindings

- ▶ Bindings are improved as part of the Linux kernel contribution process
- ▶ They are carefully reviewed by DT binding maintainers and can only be merged once approved by them
- ▶ Need for automated verifications:
 - Legacy: human readable .txt documents, hardly parsable by tools
 - Current norm: YAML-written specifications, easy to parse by humans and tools at the same time!



Device Tree binding: legacy style

[Documentation/devicetree/bindings/i2c/i2c-omap.txt](#)

I2C for OMAP platforms

-Required properties :

- compatible : Must be

"ti,omap2420-i2c" for OMAP2420 SoCs

"ti,omap2430-i2c" for OMAP2430 SoCs

"ti,omap3-i2c" for OMAP3 SoCs

"ti,omap4-i2c" for OMAP4+ SoCs

"ti,am654-i2c", "ti,omap4-i2c" for AM654 SoCs

"ti,j721e-i2c", "ti,omap4-i2c" for J721E SoCs

"ti,am64-i2c", "ti,omap4-i2c" for AM64 SoCs

- ti,hwmods : Must be "i2c<n>", n being the instance number (1-based)

- #address-cells = <1>;

- #size-cells = <0>;

Recommended properties :

- clock-frequency : Desired I2C bus clock frequency in Hz. Otherwise the default 100 kHz frequency will be used.

Optional properties:

- Child nodes conforming to i2c bus binding

Note: Current implementation will fetch base address, irq and dma from omap hwmod data base during device registration. Future plan is to migrate hwmod data base contents into device tree blob so that, all the required data will be used from device tree dts file.

Examples :

```
i2c1: i2c@0 {
    compatible = "ti,omap3-i2c";
    #address-cells = <1>;
    #size-cells = <0>;
    ti,hwmods = "i2c1";
    clock-frequency = <400000>;
};
```



Device Tree binding: YAML style

Documentation/devicetree/bindings/i2c/ti,omap4-i2c.yaml

```
# SPDX-License-Identifier: (GPL-2.0-only OR BSD-2-Clause)
%YAML 1.2
---
$id: http://devicetree.org/schemas/i2c/ti,omap4-i2c.yaml#
$schema: http://devicetree.org/meta-schemas/core.yaml#

title: I2C controllers on TI's OMAP and K3 SoCs

maintainers:
- Vignesh Raghavendra <vigneshr@ti.com>

properties:
  compatible:
    oneOf:
      - enum:
          - ti,omap2420-i2c
          - ti,omap2430-i2c
          - ti,omap3-i2c
          - ti,omap4-i2c
      - items:
          - enum:
              - ti,am4372-i2c
              - ti,am64-i2c
              - ti,am654-i2c
              - ti,j721e-i2c
          - const: ti,omap4-i2c

  reg:
    maxItems: 1
```

```
  interrupts:
    maxItems: 1

  clocks:
    maxItems: 1

  clock-names:
    const: fck

  clock-frequency: true

  power-domains: true

  "#address-cells":
    const: 1

  "#size-cells":
    const: 0

  ti,hwmods:
    description:
      Must be "i2c<n>", n being [...]
    $ref: /schemas/types.yaml#/definitions/string
    deprecated: true

  required:
    - compatible
    - reg
    - interrupts
```

```
  additionalProperties: false

  if:
    properties:
      compatible:
        enum:
          - ti,omap2420-i2c
          - ti,omap2430-i2c
          - ti,omap3-i2c
          - ti,omap4-i2c
    then:
      properties:
        ti,hwmods:
          items:
            - pattern: "^i2c([1-9])$"
    else:
      properties:
        ti,hwmods: false

  examples:
    - |
      #include <dt-bindings/interrupt-controller/irq.h>
      #include <dt-bindings/interrupt-controller/arm-gic.h>

      main_i2c0: i2c@20000000 {
        compatible = "ti,j721e-i2c", "ti,omap4-i2c";
        reg = <0x20000000 0x100>;
        interrupts = <GIC_SPI 200 IRQ_TYPE_LEVEL_HIGH>;
      };
```




Validating Device Trees

- ▶ `dtc` only does syntactic validation
- ▶ YAML bindings allow to do semantic validation
- ▶ Linux kernel make rules:
 - `make dt_binding_check`
verify that YAML bindings are valid, particularly useful if you write examples!
 - `make dtbs_check`
validate DTs currently enabled against YAML bindings
- ▶ The combination of DTS and bindings growing, it may sometimes be relevant to only check against a subset of matching schema by adding the `DT_SCHEMA_FILES` specifier on the `make` command line:
 - eg. `make DT_SCHEMA_FILES=Documentation/devicetree/bindings/trivial-devices.yaml dtbs_check`
 - Can be used with both `dt_binding_check` and `dtbs_check`



Bindings syntax: base structure

```
# SPDX-License-Identifier: (GPL-2.0-only OR BSD-2-Clause)
%YAML 1.2
---
$id: http://devicetree.org/schemas/<path>/<file-name.yaml>#
$schema: http://devicetree.org/meta-schemas/core.yaml#

title: <Type and name of the device>

maintainers:
- John Doe <john@doe.com>

description: |
    Some multiline text.

    At an additional indentation level.

# This line is a comment
properties:
  prop-a:
    ...

  prop-b:
    ...
```

Each YAML file defines one DT hierarchical level (up to two when there are children nodes expected)

- ▶ %YAML defines the expected language version
- ▶ \$id maybe not a real URL, but a unique identifier
- ▶ \$schema refers to the base meta-schema this file should be validated against (in the Github repository mentioned previously)
- ▶ properties: where the definitions start
- ▶ All possible properties should be listed
 - dash-separated lowercase names
 - names followed by a colon ':' and a new line
- ▶ Every indentation level is 2 spaces
- ▶ An empty line between property definitions



Bindings syntax: types

```
properties:
  # A boolean property, basically a yes or no
  pin-x-not-wired:      # pin-x-not-wired;
    type: boolean

  # Expects a single 32-bit numerical value
  start-offset:         # start-offset: <0x1000>;
    $ref: /schemas/types.yaml#/definitions/uint32

  # The suffix already enforces a numerical value!
  # In this case if there is no additional constraint
  # we set the property to 'true'
  my-freq-hz: true      # my-freq-hz = <100000>;

  # Expects an array of 32-bit numerical values
  supported-rates:      # supported-rates = <25>, <50>;
    $ref: /schemas/types.yaml#/definitions/uint32-array

  # A string value is expected
  instruction-set:      # instruction-set = "extended";
    $ref: /schemas/types.yaml#/definitions/string

  # Phandles will be expected
  sampling-lines:       # sampling-lines = <&pioA 1>, <&pioA 5>;
    $ref: /schemas/types.yaml#/definitions/phandle-array

  # Here as well, but no need to repeat the constraint
  # because '-gpios' is a generic suffix
  reset-gpios: true     # reset-gpios = <&gpio SOC_SPEC_IDX>;
```

- ▶ Properties must be typed, either with the `type:` or the `ref:` keyword.
 - Boolean properties require no value
 - Numerical values can be signed or unsigned but should always be 32-bit wide
 - Strings should always be fully defined (see next slides)
 - Arrays and matrices are possible as well
- ▶ Generic bindings already set the type for many properties:
 - Their values/items numbers can be constrained further
 - The types don't need to be repeated however
- ▶ `dt-schema` will enforce a type based on the property name suffix, eg: `-hz`, `-ohms`, `-us`



Bindings syntax: child nodes

```
properties:
  # The sub-node can only be named: child-node
  child-node:
    type: object

patternProperties:
  # The sub-node name is flexible, eg: child@1000, child@2a, etc
  "^child@[a-f0-9]+$":
    type: object
```

- ▶ From a yaml-schema perspective, children nodes are just another property
- ▶ A specific type shall however be enforced:
 - type: object
- ▶ Under the main `properties` keyword, property/sub-node names are fixed
 - If the sub-node name is dynamic, we shall define it under another top-level keyword, `patternProperties` and use pattern-matching regexes for the naming



Bindings syntax: expressing constraints

Besides defining precisely the different properties and their type, the content of the property values must also be constrained.

- ▶ All properties can get an additional `description` parameter, which is only readable by humans
- ▶ We try to maximize the constraints to minimize human errors
- ▶ One new line per constraint



Bindings syntax: numerical constraints

```
properties:
# The numerical value is bounded
# This is valid:
# frequency-hz = <100000>;
# frequency-hz = <0x40000>; /* 262144 Hz */
# This is not:
# frequency-hz = <0>;
# frequency-hz = <&gpio 10>;
frequency-hz:
    minimum: 10000
    maximum: 400000
    default: 100000

# This is an array with either 1 or 2 members
# This is valid:
# cs-gpios = <&gpioA 1>;
# cs-gpios = <&gpioA 1>, <gpioA 5>;
# This is not:
# cs-gpios = <&gpioA 1>, <gpioA 5>, <gpioA 6>;
# cs-gpios = <50>;
cs-gpios:
    minItems: 1
    maxItems: 2
```

▶ Example of constraints:

- minimum:/maximum: min/max values for a single value
- default: for a default value
- minItems:/maxItems: min/max number of items in an array



Bindings syntax: lists and dictionaries

```
properties:
# This is a very common compatible definition
# The only allowed combinations are (order matters):
# compatible = "vendor1,compat", "generic,compat";
# compatible = "vendor2,compat", "generic,compat";
# compatible = "legacy-compat";
compatible:
  oneOf:
    - items:
      - enum:
        - vendor1,compat
        - vendor2,compat
      - const: generic,compat
    - items:
      - const: legacy-compat

# Property name is known by dt-schema, type will be inferred
# No need for minItems/maxItems, 2 will be implied from
# the main items list!
clocks:
  items:
    - description: Interconnect
    - description: External bus

# This is valid: strength = <0>, <5>;
# This is invalid: strength = <0>;
# strength = <0>, <8>;
strength:
  $ref: /schemas/types.yaml#/definitions/uint32-array
  minItems: 2
  maxItems: 2
  items:
    maximum: 5
```

- ▶ Expressing several possible property values (works with numbers and strings):
 - Force a single expected value: `const`
 - Allow taking one value from a list: `enum`
 - watch out the indentation: 2 spaces from the previous keyword and a dash
- ▶ `const/enum` can be grouped within an `items` list, where each `items` sub-entry must be observed
- ▶ We can build abstract conditional lists (eg. on top of `items` rather than proper values like with `const/enum`):
 - XOR using `oneOf`
 - OR using `anyOf`
 - AND using `allOf`



Bindings syntax: referencing other bindings

```
# All properties/constraints defined in generic-controller.yaml
# will apply (but they can be tuned/overwritten below)
allOf:
  - $ref: generic-controller.yaml

properties:
  # Tune a property defined in generic-controller.yaml
  prop-a:
    maximum: 1

  # Allow a new, more specific property
  vendor-specific-prop: true

  # common-child-constraints.yaml will enforce a base set of
  # properties and rules
  child-node:
    type: object
    $ref: common-child-constraints.yaml
```

- ▶ It is possible to write "common" constraints in a YAML file and refer to it
 - Very usual when describing a certain type of controller
 - Refer to the generic constraints with a top-level `allOf`
 - Add constraints which are specific to the hardware implementation
 - Possible to constrain children nodes by referencing another YAML file



Bindings syntax: altering on presence of properties

```
properties:
  compatible:
    enum:
      - compat1
      - compat2

prop-a: true

prop-b: true

prop-c: true

dependencies:
  prop-a: [ 'prop-b' ]
  prop-b: [ 'prop-a' ]

allOf:
  - if:
      properties:
        compatible:
          contains:
            const: compat1
      then:
        properties:
          prop-c: false
```

- ▶ Sometimes more dynamic descriptions are needed
 - Dependencies between properties
 - A property may be needed if there is another property
 - If both or none shall be present, the dependency should be expressed twice (in both directions)
 - Changing constraints based on a property
 - Can be expressed using `if/else` statements under the top-level `allOf`
 - Typical case: a `compatible` implies tweaking a constraint



Bindings syntax: enforcing correct properties only

```
allOf:
- $ref: generic-file.yaml

properties:
  prop-a: true

  prop-b: true

child-node:
  type: object
  properties:
    prop-c: true
    prop-d: true

  required:
    - prop-c

  # No additional property than the ones above
  # will be allowed inside child-node
  additionalProperties: false

required:
- prop-a

# Only properties defined below or coming from
# generic-file.yaml will be allowed
unevaluatedProperties: false
```

- ▶ YAML files list properties and add constraints to them
 - It is still possible to add undefined properties
 - It is still possible to forget defining a mandatory property
- ▶ We need further constraints to spot typos and unexpected properties
 - `required` forces the presence
 - `additionalProperties` prevents any property not defined in **this** file to be used
 - `unevaluatedProperties` prevents any property not defined in this file nor referenced (through `allOf` or `$ref`) to be used



Bindings syntax: validating your own bindings

```
properties:
  prop-a: true
  prop-b: true

child-node:
  type: object
  additionalProperties: false

required:
  - prop-a

unevaluatedProperties: false


example:
  - |
    node@1000 {
      prop-a;
    };
```

- ▶ It is very recommended to test your bindings before testing your DTS
 - Add `examples` at the end of your file!
 - Examples are indented with 4 spaces



References

- ▶ Device Tree 101 webinar, Thomas Petazzoni (2021):
Slides: <https://bootlin.com/blog/device-tree-101-webinar-slides-and-videos/>
Video: <https://youtu.be/a9CZ1Uk30YQ>
- ▶ Kernel documentation
 - [driver-api/driver-model/](#)
 - [devicetree/](#)
 - [filesystems/sysfs](#)
- ▶ <https://devicetree.org>
- ▶ The kernel source code
 - Full of examples of other drivers!
 - Reference DT binding implementation:
[Documentation/devicetree/bindings/example-schema.yaml](#)






- ▶ Browse and update Device Trees.
- ▶ Use GPIO LEDs.
- ▶ Modify the Device Tree to enable an I2C controller and describe an I2C device.
- ▶ Write a yaml binding to validate a device description.



Introduction to pin muxing

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!






What is pin muxing?

- ▶ Modern SoCs (System on Chip) include more and more hardware blocks, many of which need to interface with the outside world using *pins*.
- ▶ However, the physical size of the chips remains small, and therefore the number of available pins is limited.
- ▶ For this reason, not all of the internal hardware block features can be exposed on the pins simultaneously.
- ▶ The pins are **multiplexed**: they expose either the functionality of hardware block A **or** the functionality of hardware block B.
- ▶ This *multiplexing* is usually software configurable.



Pin muxing diagram






Pin muxing in the Linux kernel

- ▶ Since Linux 3.2, a `pinctrl` subsystem has been added.
- ▶ This subsystem, located in [drivers/pinctrl/](#) provides a generic subsystem to handle pin muxing. It offers:
 - A pin muxing driver interface, to implement the system-on-chip specific drivers that configure the muxing.
 - A pin muxing consumer interface, for device drivers.
- ▶ Most *pinctrl* drivers provide a Device Tree binding, and the pin muxing must be described in the Device Tree.
 - The exact Device Tree binding depends on each driver. Each binding is defined in [Documentation/devicetree/bindings/pinctrl](#).



pinctrl subsystem diagram





Device Tree properties for consumer devices

The devices that require certain pins to be muxed will use the `pinctrl-<x>` and `pinctrl-names` Device Tree properties.

- ▶ The `pinctrl-0`, `pinctrl-1`, `pinctrl-<x>` properties link to a pin configuration for a given state of the device.
- ▶ The `pinctrl-names` property associates a name to each state. The name `default` is special, and is automatically selected by a device driver, without having to make an explicit *pinctrl* function call.
- ▶ See [Documentation/devicetree/bindings/pinctrl/pinctrl-bindings.txt](https://bootlin.com/documentation/devicetree/bindings/pinctrl/pinctrl-bindings.txt) for details.



Device Tree properties for consumer devices - Examples

```
i2c0: i2c@11000 {  
    ...  
    pinctrl-0 = <&pmx_twsio>;  
    pinctrl-names = "default";  
    ...  
};
```

Most common case ([arch/arm/boot/dts/marvell/kirkwood.dtsi](#))

```
i2c0: i2c@f8014000 {  
    ...  
    pinctrl-names = "default", "gpio";  
    pinctrl-0 = <&pinctrl_i2c0>;  
    pinctrl-1 = <&pinctrl_i2c0_gpio>;  
    ...  
};
```

Case with multiple pin states ([arch/arm/boot/dts/microchip/sama5d4.dtsi](#))



Defining pinctrl configurations

- ▶ The different *pinctrl configurations* must be defined as child nodes of the main *pinctrl device* (which controls the muxing of pins).
- ▶ The configurations may be defined at:
 - the SoC level (*.dtsi* file), for pin configurations that are often shared between multiple boards
 - at the board level (*.dts* file) for configurations that are board specific.
- ▶ The `pinctrl-<x>` property of the consumer device points to the pin configuration it needs through a DT *phandle*.
- ▶ The description of the configurations is specific to each *pinctrl driver*. See [Documentation/devicetree/bindings/pinctrl](https://bootlin.com/documentation/devicetree/bindings/pinctrl) for the pinctrl bindings.



Example on OMAP/AM33xx

- ▶ On OMAP/AM33xx, the `pinctrl-single` driver is used. It is common between multiple SoCs and simply allows to configure pins by writing a value to a register.
 - In each pin configuration, a `pinctrl-single,pins` value gives a list of (*register, value*) pairs needed to configure the pins.
- ▶ To know the correct values, one must use the SoC and board datasheets.

```
/* Excerpt from am335x-bone-common.dts */
```

```
&am33xx_pinmux {  
    ...  
    i2c2_pins: pinmux_i2c2_pins {  
        pinctrl-single,pins = <  
            AM33XX_PADCONF(AM335X_PIN_UART1_CTSN, PIN_INPUT_PULLUP, M  
            /* uart1_ctsn.i2c2_sda */  
            AM33XX_PADCONF(AM335X_PIN_UART1_RTSN, PIN_INPUT_PULLUP, M  
            /* uart1_rtsn.i2c2_scl */  
        >;  
    };  
};  
  
&i2c2 {  
    pinctrl-names = "default";  
    pinctrl-0 = <&i2c2_pins>;  
  
    status = "okay";  
    clock-frequency = <400000>;  
    ...  
  
    pressure@76 {  
        compatible = "bosch,bmp280";  
        reg = <0x76>;  
    };  
};
```



Example on the Allwinner A20 SoC

SoC level

```
/ {
    ...
    soc {
        ...
        pio: pinctrl@01c20800 {
            compatible = "allwinner,sun7i-a20-pinctrl";
            reg = <0x01c20800 0x400>;
            ...
        }

        UART0
        pin mux
        config | uart@pb_pins: uart@pb_pins {
            pins = "PB22", "PB23";
            function = "uart0";
        };
        ...
    };
};
};
```

arch/arm/boot/dts/sun7i-a20.dtsi

Board level

```
/ {
    ...
    leds {
        compatible = "gpio-leds";
        pinctrl-names = "default";
        pinctrl-0 = <&led_pins_olinuxino>;

        led {
            label = "a20-olinuxino-micro:green:usr";
            gpios = <&pio 7 2 GPIO_ACTIVE_HIGH>;
            default-state = "on";
        };
    };
    ...
};

&pio {
    ...

    led_pins_olinuxino: led_pins@0 {
        pins = "PH2";
        function = "gpio_out";
        drive-strength = <20>;
    };
    ...
};

&uart0 {
    pinctrl-names = "default";
    pinctrl-0 = <&uart0_pb_pins>;
    status = "okay";
};
```

arch/arm/boot/dts/sun7i-a20-olinuxino-micro.dts



Practical lab - Setup pinmuxing to enable I2C communication




- ▶ Configure the pinmuxing for the I2C bus used to communicate with the Nunchuk
- ▶ Validate that the I2C communication works with user space tools.



Linux device and driver model

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Introduction



The need for a device model?

- ▶ The Linux kernel runs on a wide range of architectures and hardware platforms, and therefore needs to **maximize the reusability** of code between platforms.
- ▶ For example, we want the same *USB device driver* to be usable on a x86 PC, or an ARM platform, even though the USB controllers used on these platforms are different.
- ▶ This requires a clean organization of the code, with the *device drivers* separated from the *controller drivers*, the hardware description separated from the drivers themselves, etc.
- ▶ This is what the Linux kernel **Device Model** allows, in addition to other advantages covered in this section.




Kernel and device drivers

In Linux, a driver is always interfacing with:

- ▶ a **framework** that allows the driver to expose the hardware features in a generic way.
- ▶ a **bus infrastructure**, part of the device model, to detect/communicate with the hardware.

This section focuses on the *bus infrastructure*, while *kernel frameworks* are covered later in this training.





Device model data structures

- ▶ The *device model* is organized around three main data structures:
 - The `struct bus_type` structure, which represents one type of bus (USB, PCI, I2C, etc.)
 - The `struct device_driver` structure, which represents one driver capable of handling certain devices on a certain bus.
 - The `struct device` structure, which represents one device connected to a bus
- ▶ The kernel uses inheritance to create more specialized versions of `struct device_driver` and `struct device` for each bus subsystem.



Bus drivers

- ▶ The first component of the device model is the bus driver
 - One bus driver for each type of bus: USB, PCI, SPI, MMC, I2C, etc.
- ▶ It is responsible for
 - Registering the bus type (`struct bus_type`)
 - Allowing the registration of adapter drivers (USB controllers, I2C adapters, etc.), able to detect the connected devices (if possible), and providing a communication mechanism with the devices
 - Allowing the registration of device drivers (USB devices, I2C devices, PCI devices, etc.), managing the devices
 - Matching the device drivers against the devices detected by the adapter drivers.
 - Provides an API to implement both adapter drivers and device drivers
 - Defining driver and device specific structures, eg. `struct usb_driver` and `struct usb_interface`




- ▶ The bus, device, drivers, etc. structures are internal to the kernel
- ▶ The `sysfs` virtual filesystem offers a mechanism to export such information to user space
- ▶ Used for example by `udev` to provide automatic module loading, firmware loading, mounting of external media, etc.
- ▶ `sysfs` is usually mounted in `/sys`
 - `/sys/bus/` contains the list of buses
 - `/sys/devices/` contains the list of devices
 - `/sys/class` enumerates devices by the framework they are registered to (net, input, block...), whatever bus they are connected to. Very useful!



Example of the USB bus




Example: USB bus 1/3



Hardware view of the bus



Example: USB bus 2/3



Device model view of the bus




Example: USB bus 3/3

- ▶ Core infrastructure (bus driver)
 - [drivers/usb/core/](#)
 - `struct bus_type` is defined in [drivers/usb/core/driver.c](#) and registered in [drivers/usb/core/usb.c](#)
- ▶ Adapter drivers
 - [drivers/usb/host/](#)
 - For EHCI, UHCI, OHCI, XHCI, and their implementations on various systems (Microchip, IXP, Xilinx, OMAP, Samsung, PXA, etc.)
- ▶ Device drivers
 - Everywhere in the kernel tree, classified by their type (Example: [drivers/net/usb/](#))



Example of device driver

- ▶ To illustrate how drivers are implemented to work with the device model, we will study the source code of a driver for a USB network card
 - It is USB device, so it has to be a USB device driver
 - It exposes a network device, so it has to be a network driver
 - Most drivers rely on a bus infrastructure (here, USB) and register themselves in a framework (here, network)
- ▶ We will only look at the device driver side, and not the adapter driver side
- ▶ The driver we will look at is [drivers/net/usb/rtl8150.c](#)





Device identifiers

- ▶ Defines the set of devices that this driver can manage, so that the USB core knows for which devices this driver should be used
- ▶ The `MODULE_DEVICE_TABLE()` macro allows `depmod` (run by `make modules_install`) to extract the relationship between device identifiers and drivers, so that drivers can be loaded automatically by `udev`. See `/lib/modules/$(uname -r)/modules.{alias,usbmap}`

```
static struct usb_device_id rtl8150_table[] = {  
    { USB_DEVICE(VENDOR_ID_REALTEK, PRODUCT_ID_RTL8150) },  
    { USB_DEVICE(VENDOR_ID_MELCO, PRODUCT_ID_LUAKTX) },  
    { USB_DEVICE(VENDOR_ID_MICRONET, PRODUCT_ID_SP128AR) },  
    { USB_DEVICE(VENDOR_ID_LONGSHINE, PRODUCT_ID_LCS8138TX) },  
    [...]  
};  
MODULE_DEVICE_TABLE(usb, rtl8150_table);
```



Instantiation of usb_driver

- ▶ `struct usb_driver` is a structure defined by the USB core. Each USB device driver must instantiate it, and register itself to the USB core using this structure
- ▶ This structure inherits from `struct device_driver`, which is defined by the device model.

```
static struct usb_driver rtl8150_driver = {  
    .name = "rtl8150",  
    .probe = rtl8150_probe,  
    .disconnect = rtl8150_disconnect,  
    .id_table = rtl8150_table,  
    .suspend = rtl8150_suspend,  
    .resume = rtl8150_resume  
};
```



Driver registration and unregistration

- ▶ When the driver is loaded / unloaded, it must register / unregister itself to / from the USB core
- ▶ Done using `usb_register()` and `usb_deregister()`, provided by the USB core.

```
static int __init usb_rtl8150_init(void)
{
    return usb_register(&rtl8150_driver);
}

static void __exit usb_rtl8150_exit(void)
{
    usb_deregister(&rtl8150_driver);
}

module_init(usb_rtl8150_init);
module_exit(usb_rtl8150_exit);
```


- ▶ All this code is actually replaced by a call to the `module_usb_driver()` macro:

```
module_usb_driver(rtl8150_driver);
```



At Initialization

- ▶ The USB adapter driver that corresponds to the USB controller of the system registers itself to the USB core
- ▶ The `rtl8150` USB device driver registers itself to the USB core




- ▶ The USB core now knows the association between the vendor/product IDs of `rtl8150` and the `struct usb_driver` structure of this driver



When a device is detected

Step 2: USB core looks up the registered IDs, and finds the matching driver





Probe method

- ▶ Invoked **for each device** bound to a driver
- ▶ The `probe()` method receives as argument a structure describing the device, usually specialized by the bus infrastructure (`struct pci_dev`, `struct usb_interface`, etc.)
- ▶ This function is responsible for
 - Initializing the device, mapping I/O memory, registering the interrupt handlers. The bus infrastructure provides methods to get the addresses, interrupt numbers and other device-specific information.
 - Registering the device to the proper kernel framework, for example the network infrastructure.



Example: probe() and disconnect() methods

```
static int rtl8150_probe(struct usb_interface *intf,
                        const struct usb_device_id *id)
{
    rtl8150_t *dev;
    struct net_device *netdev;

    netdev = alloc_etherdev(sizeof(rtl8150_t));
    [...]
    dev = netdev_priv(netdev);
    tasklet_init(&dev->tl, rx_fixup, (unsigned long)dev);
    spin_lock_init(&dev->rx_pool_lock);
    [...]
    netdev->netdev_ops = &rtl8150_netdev_ops;
    alloc_all_urbs(dev);
    [...]
    usb_set_intfdata(intf, dev);
    SET_NETDEV_DEV(netdev, &intf->dev);
    register_netdev(netdev);

    return 0;
}
```



```
static void rtl8150_disconnect(struct usb_interface *intf)
{
    rtl8150_t *dev = usb_get_intfdata(intf);

    usb_set_intfdata(intf, NULL);
    if (dev) {
        set_bit(RTL8150_UNPLUG, &dev->flags);
        tasklet_kill(&dev->tl);
        unregister_netdev(dev->netdev);
        unlink_all_urbs(dev);
        free_all_urbs(dev);
        free_skb_pool(dev);
        if (dev->rx_skb)
            dev_kfree_skb(dev->rx_skb);
        kfree(dev->intr_buff);
        free_netdev(dev->netdev);
    }
}
```

Source: [drivers/net/usb/rtl8150.c](#)



The model is recursive





Platform drivers



Platform devices

- ▶ Amongst the non-discoverable devices, a huge family are the devices that are directly part of a system-on-chip: UART controllers, Ethernet controllers, SPI or I2C controllers, graphic or audio devices, etc.
- ▶ In the Linux kernel, a special bus, called the **platform bus** has been created to handle such devices. Those get controlled through **memory-mapped registers**.
- ▶ It supports **platform drivers** that handle **platform devices**.
- ▶ It works like any other bus (USB, PCI), except that devices are enumerated statically instead of being discovered dynamically.



Implementation of a platform driver (1)

The driver implements a `struct platform_driver` structure (example taken from `drivers/tty/serial/imx.c`, simplified)

```
static struct platform_driver serial_imx_driver = {  
    .probe      = serial_imx_probe,  
    .remove     = serial_imx_remove,  
    .id_table   = imx_uart_devtype,  
    .driver     = {  
        .name    = "imx-uart",  
        .of_match_table = imx_uart_dt_ids,  
        .pm      = &imx_serial_port_pm_ops,  
    },  
};
```



Implementation of a platform driver (2)

... and registers its driver to the platform driver infrastructure

```
static int __init imx_serial_init(void) {  
    return platform_driver_register(&serial_imx_driver);  
}  
  
static void __exit imx_serial_cleanup(void) {  
    platform_driver_unregister(&serial_imx_driver);  
}  
  
module_init(imx_serial_init);  
module_exit(imx_serial_cleanup);
```

Most drivers actually use the `module_platform_driver()` macro when they do nothing special in `init()` and `exit()` functions:

```
module_platform_driver(serial_imx_driver);
```




- ▶ As platform devices cannot be detected dynamically, they are defined statically
 - Legacy way: by direct instantiation of `struct platform_device` structures, as done on a few old ARM platforms. The device was part of a list, and the list of devices was added to the system during board initialization.
 - Current way: by parsing an "external" description, like a *device tree* on most embedded platforms today, from which `struct platform_device` instances are created.



Using additional hardware resources

- ▶ Regular DT descriptions contain many information. It includes phandles (pointers) towards additional hardware blocks which cannot be discovered.
 - I/O register addresses and IRQ lines are available through a `struct resource` array associated to each `struct platform_device`.
 - Information relevant to a given subsystem is parsed by that specific subsystem. Examples are clocks, GPIOs or DMA. A subsystem is responsible for:
 - instantiating its components,
 - offering an API to use those objects from device drivers.
 - Specific properties are directly retrieved by device drivers, through (expensive) DT lookups.
- ▶ All these methods allow the same driver to be used with multiple devices functioning similarly, but with different addresses, IRQs, etc.



Using resources

- ▶ The platform driver has access to the resources provided by the platform bus:

```
res = platform_get_resource(pdev, IORESOURCE_MEM, 0);  
base = ioremap(res->start, PAGE_SIZE);  
sport->rxirq = platform_get_irq(pdev, 0);
```

- ▶ As well as the various subsystem-provided dependencies through individual APIs:
 - `clk_get()`
 - `gpio_request()`
 - `dma_request_channel()`



Driver data

- ▶ In addition to the per-device resources and information, drivers may require driver-specific information to behave slightly differently when different flavors of an IP block are driven by the same driver.
- ▶ A `const void *data` pointer can be used to store per-compatible specificities:

```
static const struct of_device_id marvell_nfc_of_ids[] = {  
    {  
        .compatible = "marvell,armada-8k-nand-controller",  
        .data = &marvell_armada_8k_nfc_caps,  
    },  
};
```


- ▶ Which can be retrieved in the probe with:

```
/* Get NAND controller capabilities */  
if (pdev->id_entry) /* legacy way */  
    nfc->caps = (void *)pdev->id_entry->driver_data;  
else /* current way */  
    nfc->caps = of_device_get_match_data(&pdev->dev);
```



Introduction to the I2C subsystem

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!






What is I2C?

- ▶ A very commonly used low-speed bus to connect on-board and external devices to the processor.
- ▶ Uses only two wires: SDA for the data, SCL for the clock.
- ▶ It is a master/slave bus: only the master can initiate transactions, and slaves can only reply to transactions initiated by masters.
- ▶ In a Linux system, the I2C controller embedded in the processor is typically the master, controlling the bus.
- ▶ Each slave device is identified by an I2C address (you can't have 2 devices with the same address on the same bus). Each transaction initiated by the master contains this address, which allows the relevant slave to recognize that it should reply to this particular transaction.



An I2C bus example





The I2C bus driver

- ▶ Like all bus subsystems, the I2C bus driver is responsible for:
 - Providing an API to implement I2C controller drivers
 - Providing an API to implement I2C device drivers, in kernel space
 - Providing an API to implement I2C device drivers, in user space
- ▶ The core of the I2C bus driver is located in [drivers/i2c/](#).
- ▶ The I2C controller drivers are located in [drivers/i2c/busses/](#).
- ▶ The I2C device drivers are located throughout [drivers/](#), depending on the framework used to expose the devices (e.g. [drivers/input/](#) for input devices).



Registering an I2C device driver

- ▶ Like all bus subsystems, the I2C subsystem defines a `struct i2c_driver` that inherits from `struct device_driver`, and which must be instantiated and registered by each I2C device driver.
 - As usual, this structure points to the `->probe()` and `->remove()` functions.
 - It also contains a legacy `id_table`, used for non-DT based probing of I2C devices.
- ▶ The `i2c_add_driver()` and `i2c_del_driver()` functions are used to register/unregister the driver.
- ▶ If the driver doesn't do anything else in its `init()/exit()` functions, it is advised to use the `module_i2c_driver()` macro instead.



Registering an I2C device driver: example

```
static const struct i2c_device_id adxl345_i2c_id[] = {
    { "adxl345", ADXL345 },
    { "adxl375", ADXL375 },
    { }
};

MODULE_DEVICE_TABLE(i2c, adxl345_i2c_id);

static const struct of_device_id adxl345_of_match[] = {
    { .compatible = "adi,adxl345" },
    { .compatible = "adi,adxl375" },
    { },
};

MODULE_DEVICE_TABLE(of, adxl345_of_match);

static struct i2c_driver adxl345_i2c_driver = {
    .driver = {
        .name = "adxl345_i2c",
        .of_match_table = adxl345_of_match,
    },
    .probe = adxl345_i2c_probe,
    .remove = adxl345_i2c_remove,
    .id_table = adxl345_i2c_id,
};

module_i2c_driver(adxl345_i2c_driver);
```

From `drivers/iio/accel/adxl345_i2c.c`



Registering an I2C device: non-DT

- ▶ On non-DT platforms, the `struct i2c_board_info` structure allows to describe how an I2C device is connected to a board.
- ▶ Such structures are normally defined with the `I2C_BOARD_INFO()` helper macro.
 - Takes as argument the device name and the slave address of the device on the bus.
- ▶ An array of such structures is registered on a per-bus basis using `i2c_register_board_info()`, when the platform is initialized.



Registering an I2C device, non-DT example

arch/arm/mach-iop32x/em7210.c

```
static struct i2c_board_info __initdata em7210_i2c_devices[] = {
    { I2C_BOARD_INFO("rs5c372a", 0x32) },
};

static void __init em7210_init_machine(void)
{
    register_iop32x_gpio();
    platform_device_register(&em7210_serial_device);
    platform_device_register(&iop3xx_i2c0_device);
    platform_device_register(&iop3xx_i2c1_device);
    platform_device_register(&em7210_flash_device);
    platform_device_register(&iop3xx_dma_0_channel);
    platform_device_register(&iop3xx_dma_1_channel);

    i2c_register_board_info(0, em7210_i2c_devices,
        ARRAY_SIZE(em7210_i2c_devices));
}
```



Registering an I2C device, in the DT

- ▶ In the Device Tree, the I2C controller device is typically defined in the `.dtsi` file that describes the processor.
 - Normally defined with `status = "disabled"`.
- ▶ At the board/platform level:
 - the I2C controller device is enabled (`status = "okay"`)
 - the I2C bus frequency is defined, using the `clock-frequency` property.
 - the I2C devices on the bus are described as children of the I2C controller node, where the `reg` property gives the I2C slave address on the bus.
- ▶ See the binding for the corresponding driver for a specification of the expected DT properties. Example:

[Documentation/devicetree/bindings/i2c/ti,omap4-i2c.yaml](#)



Registering an I2C device, DT example (1/2)

Definition of the I2C controller

```
i2c0: i2c@01c2ac00 {  
    compatible = "allwinner,sun7i-a20-i2c",  
                "allwinner,sun4i-a10-i2c";  
    reg = <0x01c2ac00 0x400>;  
    interrupts = <GIC_SPI 7 IRQ_TYPE_LEVEL_HIGH>;  
    clocks = <&apb1_gates 0>;  
    status = "disabled";  
    #address-cells = <1>;  
    #size-cells = <0>;  
};
```

From `arch/arm/boot/dts/allwinner/sun7i-a20.dtsi`

`#address-cells`: number of 32-bit values needed to encode the address fields

`#size-cells`: number of 32-bit values needed to encode the size fields

See details in https://elinux.org/Device_Tree_Usage



Registering an I2C device, DT example (2/2)

Definition of the I2C device

```
&i2c0 {  
    pinctrl-names = "default";  
    pinctrl-0 = <&i2c0_pins_a>;  
    status = "okay";  
  
    axp209: pmic@34 {  
        compatible = "x-powers,axp209";  
        reg = <0x34>;  
        interrupt-parent = <&nmi_intc>;  
        interrupts = <0 IRQ_TYPE_LEVEL_LOW>;  
  
        interrupt-controller;  
        #interrupt-cells = <1>;  
    };  
};
```

From [arch/arm/boot/dts/allwinner/sun7i-a20-olinuxino-micro.dts](#)



probe() and remove()

- ▶ The `->probe()` function is responsible for initializing the device and registering it in the appropriate kernel framework. It receives as argument:
 - An `struct i2c_client` pointer, which represents the I2C device itself. This structure inherits from `struct device`.
 - On older kernels (`< v6.4`), `->probe()` was taking a second (unused) argument, the removal of this other argument implied the use of another probe function for some kernel releases, called `->probe_new()`.
- ▶ The `->remove()` function is responsible for unregistering the device from the kernel framework and shut it down. It receives as argument:



Probe example

```
static int da311_probe(struct i2c_client *client)
{
    struct iio_dev *indio_dev;          // framework structure
    da311_data *data;                   // per device structure
    ...
    // Allocate framework structure with per device struct inside
    indio_dev = devm_iio_device_alloc(&client->dev, sizeof(*data));
    data = iio_priv(indio_dev);
    data->client = client;
    i2c_set_clientdata(client, indio_dev);
    // Prepare device and initialize indio_dev
    ...
    // Register device to framework
    ret = iio_device_register(indio_dev);
    ...
    return ret;
}
```

From `drivers/iio/accel/da311.c`



Remove example

```
static int da311_remove(struct i2c_client *client)
{
    struct iio_dev *indio_dev = i2c_get_clientdata(client);
    // Unregister device from framework
    iio_device_unregister(indio_dev);
    return da311_enable(client, false);
}
```

From `drivers/iio/accel/da311.c`



Communicating with the I2C device: raw API

The most **basic API** to communicate with the I2C device provides functions to either send or receive data:

- ▶ Send a `buf` to the I2C device with:

```
int i2c_master_send(const struct i2c_client *client, const char *buf, int count);
```

- ▶ Receive a `count` bytes from the I2C device and save them in `buf` with:

```
int i2c_master_recv(const struct i2c_client *client, char *buf, int count);
```

Both functions return a negative error number in case of failure, otherwise the number of transmitted bytes.



Communicating with the I2C device: message transfer

The message transfer API allows to describe **transfers** that consists of several **messages**, with each message being a transaction in one direction:

```
int i2c_transfer(struct i2c_adapter *adap, struct i2c_msg *msgs, int num);
```

- ▶ The `struct i2c_adapter` pointer can be found by using `client->adapter`
- ▶ The `struct i2c_msg` structure defines the length, location, and direction of the message.



I2C: message transfer example

```
static int st1232_ts_read_data(struct st1232_ts_data *ts)
{
    ...
    struct i2c_client *client = ts->client;
    struct i2c_msg msg[2];
    int error;
    ...
    u8 start_reg = ts->chip_info->start_reg;
    u8 *buf = ts->read_buf;

    /* read touchscreen data */
    msg[0].addr = client->addr;
    msg[0].flags = 0;
    msg[0].len = 1;
    msg[0].buf = &start_reg;

    msg[1].addr = ts->client->addr;
    msg[1].flags = I2C_M_RD;
    msg[1].len = ts->read_buf_len;
    msg[1].buf = buf;

    error = i2c_transfer(client->adapter, msg, 2);
    ...
}
```

From [drivers/input/touchscreen/st1232.c](#)



I2C functionality

- ▶ Not all I2C controllers support all functionalities.
- ▶ The I2C controller drivers therefore tell the I2C core which functionalities they support.
- ▶ An I2C device driver must check that the functionalities they need are provided by the I2C controller in use on the system.
- ▶ The `i2c_check_functionality()` function allows to make such a check.
- ▶ Examples of functionalities: `I2C_FUNC_I2C` to be able to use the raw I2C functions, `I2C_FUNC_SMBUS_BYTE_DATA` to be able to use SMBus commands to write a command and read/write one byte of data.
- ▶ See `include/uapi/linux/i2c.h` for the full list of existing functionalities.



- ▶ <https://en.wikipedia.org/wiki/I2C>, general presentation of the I2C protocol
- ▶ [i2c/](#), details about Linux support for I2C
 - [i2c/writing-clients](#)
How to write I2C kernel device drivers
 - [i2c/dev-interface](#)
How to write I2C user-space device drivers
 - [i2c/instantiating-devices](#)
How to instantiate devices
 - [i2c/smbus-protocol](#)
Details on the SMBus functions
 - [i2c/functionality](#)
How the functionality mechanism works
- ▶ See also Luca Ceresoli's introduction to I2C ([slides](#), [video](#)).



Practical lab - Communicate with the Nunchuk




- ▶ Explore the content of `/dev` and `/sys` and the devices available on the embedded hardware platform.
- ▶ Implement a driver that registers as an I2C driver.
- ▶ Communicate with the Nunchuk and extract data from it.



Kernel frameworks for device drivers

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!






Kernel and Device Drivers

In Linux, a driver is always interfacing with:

- ▶ a **framework** that allows the driver to expose the hardware features to user space applications.
- ▶ a **bus infrastructure**, part of the device model, to detect/communicate with the hardware.

This section focuses on the *kernel frameworks*, while the *bus infrastructure* was covered earlier in this training.





User space vision of devices



Types of devices

Under Linux, there are essentially four types of devices:

- ▶ **Network devices.** They are represented as network interfaces, visible in user space using `ip a`
- ▶ **Block devices.** They are used to provide user space applications access to raw storage devices (hard disks, USB keys). They are visible to the applications as *device files* in `/dev`.
- ▶ **Character devices.** They are used to provide user space applications access to all other types of devices (input, sound, graphics, serial, etc.). They are also visible to the applications as *device files* in `/dev`.
- ▶ **Sysfs devices.** They don't have any of the above user space interfaces, only a representation in sysfs. "Internal" device drivers fall under this (e.g. `pinctrl`), but also some user-space accessible devices. E.g. `gpio` (deprecated), IIO (Industrial I/O).

→ Most devices are *character devices*, so we will study these in more details.



Devices: everything is a file

- ▶ A very important UNIX design decision was to represent most *system objects* as files
- ▶ It allows applications to manipulate all *system objects* with the normal file API (open, read, write, close, etc.)
- ▶ So, devices had to be represented as files to the applications
- ▶ This is done through a special artifact called a **device file**
- ▶ It is a special type of file, associating a file name visible to user space applications to a triplet the kernel understands:
 - *type*: char or block
 - *major*: class of the device
 - *minor*: unique identifier in a class
- ▶ All *device files* are by convention stored in the `/dev` directory



Device files examples

Example of device files in a Linux system

```
$ ls -l /dev/ttyS0 /dev/tty1 /dev/sda /dev/sda1 /dev/sda2 /dev/sdc1 /dev/zero
brw-rw---- 1 root disk      8,  0 2011-05-27 08:56 /dev/sda
brw-rw---- 1 root disk      8,  1 2011-05-27 08:56 /dev/sda1
brw-rw---- 1 root disk      8,  2 2011-05-27 08:56 /dev/sda2
brw-rw---- 1 root disk      8, 32 2011-05-27 08:56 /dev/sdc
crw----- 1 root root       4,  1 2011-05-27 08:57 /dev/tty1
crw-rw---- 1 root dialout  4, 64 2011-05-27 08:56 /dev/ttyS0
crw-rw-rw- 1 root root       1,  5 2011-05-27 08:56 /dev/zero
```

Example C code that uses the usual file API to write data to a serial port

```
int fd;
fd = open("/dev/ttyS0", O_RDWR);
write(fd, "Hello", 5);
close(fd);
```



Creating device files

- ▶ Before Linux 2.6.32, on basic Linux systems, the device files had to be created manually using the `mknod` command
 - `mknod /dev/<device> [c|b] major minor`
 - Needed root privileges
 - Coherency between device files and devices handled by the kernel was left to the system developer
- ▶ The `devtmpfs` virtual filesystem can be mounted on `/dev` and contains all the devices registered to kernel frameworks. The `CONFIG_DEVTMPFS_MOUNT` kernel configuration option makes the kernel mount it automatically at boot time, except when booting on an `initramfs`.
- ▶ `devtmpfs` can be supplemented by userspace tools like `udev` or `mdev` to adjust permission/ownership, load kernel modules automatically and create symbolic links to devices.



Character drivers




A character driver in the kernel

- ▶ From the point of view of an application, a *character device* is essentially a **file**.
- ▶ Character device drivers therefore implement **operations** that let applications think the device is a file.
- ▶ In order to achieve this, a character driver implements the operations it wants from the `struct file_operations` structure: `read`, `write`, `ioctl`, etc.
- ▶ The Linux filesystem layer will ensure that the driver's operations are called when a user space application makes the corresponding system call.



From user space to the kernel: character devices





File operations

Here are the most important operations for a character driver, from the definition of `struct file_operations`:

```
struct file_operations {
    struct module *owner;
    ssize_t (*read) (struct file *, char __user *,
        size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *,
        size_t, loff_t *);
    long (*unlocked_ioctl) (struct file *, unsigned int,
        unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*release) (struct inode *, struct file *);
    ...
};
```

Many operations exist, they are all optional.



open() and release()

▶ `int foo_open(struct inode *i, struct file *f)`

- Called when user space opens the device file.
- **Only implement this function when you do something special with the device at open() time.**
- `struct inode` is a structure that uniquely represents a file in the filesystem (be it a regular file, a directory, a symbolic link, a character or block device)
- `struct file` is a structure created every time a file is opened. Several file structures can point to the same `inode` structure.
 - Contains information like the current position, the opening mode, etc.
 - Has a `void *private_data` pointer that one can freely use.
 - A pointer to the `file` structure is passed to all other operations

▶ `int foo_release(struct inode *i, struct file *f)`

- Called when user space closes the file.
- **Only implement this function when you do something special with the device at close() time.**



read() and write()

- ▶ `ssize_t foo_read(struct file *f, char __user *buf, size_t sz, loff_t *off)`
 - Called when user space uses the `read()` system call on the device.
 - Must read data from the device, write at most `sz` bytes to the user space buffer `buf`, and update the current position in the file `off`. `f` is a pointer to the same file structure that was passed in the `open()` operation
 - Must return the number of bytes read.
 - 0 is usually interpreted by userspace as the end of the file.
 - On UNIX, `read()` operations typically block when there isn't enough data to read from the device
- ▶ `ssize_t foo_write(struct file *f, const char __user *buf, size_t sz, loff_t *off)`
 - Called when user space uses the `write()` system call on the device
 - The opposite of `read`, must read at most `sz` bytes from `buf`, write it to the device, update `off` and return the number of bytes written.



Exchanging data with user space 1/3

- ▶ Kernel code isn't allowed to directly access user space memory, using `memcpy()` or direct pointer dereferencing
 - User pointer dereferencing is disabled by default to make it harder to exploit vulnerabilities.
 - If the address passed by the application was invalid, the kernel could segfault.
 - **Never** trust user space. A malicious application could pass a kernel address which you could overwrite with device data (`read` case), or which you could dump to the device (`write` case).
 - Doing so does not work on some architectures anyway.
- ▶ To keep the kernel code portable, secure, and have proper error handling, your driver must use special kernel functions to exchange data with user space.



Exchanging data with user space 2/3

▶ A single value

- `get_user(v, p);`
 - The kernel variable `v` gets the value pointed by the user space pointer `p`
- `put_user(v, p);`
 - The value pointed by the user space pointer `p` is set to the contents of the kernel variable `v`.


▶ A buffer

- `unsigned long copy_to_user(void __user *to, const void *from, unsigned long n);`
- `unsigned long copy_from_user(void *to, const void __user *from, unsigned long n);`

- ▶ The return value must be checked. Zero on success, non-zero on failure. If non-zero, the convention is to return `-EFAULT`.



Exchanging data with user space 3/3





Zero copy access to user memory

- ▶ Having to copy data to or from an intermediate kernel buffer can become expensive when the amount of data to transfer is large (video).
- ▶ *Zero copy* options are possible:
 - `mmap()` system call to allow user space to directly access memory mapped I/O space. See our `mmap()` chapter.
 - `get_user_pages()` and related functions to get a mapping to user pages without having to copy them.



unlocked_ioctl()

- ▶ `long unlocked_ioctl(struct file *f, unsigned int cmd, unsigned long arg)`
- Associated to the `ioctl()` system call.
 - Called `unlocked` because it didn't hold the Big Kernel Lock (gone now).
 - Allows to extend the driver capabilities beyond the limited read/write API.
 - For example: changing the speed of a serial port, setting video output format, querying a device serial number... Used extensively in the V4L2 (video) and ALSA (sound) driver frameworks.
 - `cmd` is a number identifying the operation to perform.
See [driver-api/ioctl](#) for the recommended way of choosing `cmd` numbers.
 - `arg` is the optional argument passed as third argument of the `ioctl()` system call.
Can be an integer, an address, etc.
 - The semantic of `cmd` and `arg` is driver-specific.



ioctl() example: kernel side

```
#include <linux/phantom.h>

static long phantom_ioctl(struct file *file, unsigned int cmd,
                          unsigned long arg)
{
    struct phm_reg r;
    void __user *argp = (void __user *)arg;

    switch (cmd) {
    case PHN_SET_REG:
        if (copy_from_user(&r, argp, sizeof(r)))
            return -EFAULT;
        /* Do something */
        break;
    ...
    case PHN_GET_REG:
        if (copy_to_user(argp, &r, sizeof(r)))
            return -EFAULT;
        /* Do something */
        break;
    ...
    default:
        return -ENOTTY;
    }

    return 0;
}
```

Selected excerpt from [drivers/misc/phantom.c](#)



ioctl() Example: Application Side

```
#include <linux/phantom.h>

int main(void)
{
    int fd, ret;
    struct phm_reg reg;

    fd = open("/dev/phantom");
    assert(fd > 0);

    reg.field1 = 42;
    reg.field2 = 67;

    ret = ioctl(fd, PHN_SET_REG, &reg);
    assert(ret == 0);

    return 0;
}
```



The concept of kernel frameworks




Beyond character drivers: kernel frameworks

- ▶ Many device drivers are not implemented directly as character drivers
- ▶ They are implemented under a *framework*, specific to a given device type (framebuffer, V4L, serial, etc.)
 - The framework allows to factorize the common parts of drivers for the same type of devices
 - From user space, they are still seen as character devices by the applications
 - The framework allows to provide a coherent user space interface (`ioctl`, etc.) for every type of device, regardless of the driver



Example: Some Kernel Frameworks





Example: the input subsystem




What is the input subsystem?

- ▶ The input subsystem takes care of all the input events coming from the human user.
- ▶ Initially written to support the USB *HID* (Human Interface Device) devices, it quickly grew up to handle all kinds of inputs (using USB or not): keyboards, mice, joysticks, touchscreens, etc.
- ▶ The input subsystem is split in two parts:
 - **Device drivers**: they talk to the hardware (for example via USB), and provide events (keystrokes, mouse movements, touchscreen coordinates) to the input core
 - **Event handlers**: they get events from drivers and pass them where needed via various interfaces (most of the time through `evdev`)
- ▶ In user space it is usually used by the graphic stack such as *X.Org*, *Wayland* or *Android's InputManager*.



Input subsystem diagram





Input subsystem overview

- ▶ Kernel option `CONFIG_INPUT`
 - menuconfig INPUT
 - tristate "Generic input layer (needed for keyboard, mouse, ...)"
- ▶ Implemented in `drivers/input/`
 - `input.c`, `input-poller.c`, `evdev.c`...
- ▶ Defines the user/kernel API
 - `include/uapi/linux/input.h`
- ▶ Defines the set of operations an input driver must implement and helper functions for the drivers
 - `struct input_dev` for the device driver part
 - `struct input_handler` for the event handler part
 - `include/linux/input.h`



Input subsystem API 1/3

An *input device* is described by a very long `struct input_dev` structure, an excerpt is:

```
struct input_dev {
    const char *name;
    [...]
    struct input_id id;
    [...]
    unsigned long evbit[BITS_TO_LONGS(EV_CNT)];
    unsigned long keybit[BITS_TO_LONGS(KEY_CNT)];
    [...]
    int (*getkeycode)(struct input_dev *dev,
                     struct input_keymap_entry *ke);
    [...]
    int (*open)(struct input_dev *dev);
    [...]
    int (*event)(struct input_dev *dev, unsigned int type,
                 unsigned int code, int value);
    [...]
};
```

Before being used, this structure must be allocated and initialized, typically with:

```
struct input_dev *devm_input_allocate_device(struct device *dev);
```



Input subsystem API 2/3

- ▶ Depending on the type of events that will be generated, the input bit fields `evbit` and `keybit` must be configured: For example, for a button we only generate `EV_KEY` type events, and from these only `BTN_0` events code:

```
set_bit(EV_KEY, myinput_dev.evbit);  
set_bit(BTN_0, myinput_dev.keybit);
```

- ▶ Once the *input device* is allocated and filled, the function to register it is:
`int input_register_device(struct input_dev *);`



Input subsystem API 3/3

The events are sent by the driver to the event handler using

```
void input_event(struct input_dev *dev, unsigned int type, unsigned int code, int value)
```

- ▶ The event types are documented in [input/event-codes](#)
- ▶ An event is composed by one or several input data changes (packet of input data changes) such as the button state, the relative or absolute position along an axis, etc..
- ▶ The input subsystem provides other wrappers such as:
 - [input_report_key\(\)](#)
 - [input_report_abs\(\)](#)

After submitting potentially multiple events, the *input* core must be notified by calling:

```
void input_sync(struct input_dev *dev)
```



Example from drivers/hid/usbhid/usbmouse.c

```
static void usb_mouse_irq(struct urb *urb)
{
    struct usb_mouse *mouse = urb->context;
    signed char *data = mouse->data;
    struct input_dev *dev = mouse->dev;
    ...

    input_report_key(dev, BTN_LEFT,  data[0] & 0x01);
    input_report_key(dev, BTN_RIGHT, data[0] & 0x02);
    input_report_key(dev, BTN_MIDDLE, data[0] & 0x04);
    input_report_key(dev, BTN_SIDE,  data[0] & 0x08);
    input_report_key(dev, BTN_EXTRA, data[0] & 0x10);

    input_report_rel(dev, REL_X,      data[1]);
    input_report_rel(dev, REL_Y,      data[2]);
    input_report_rel(dev, REL_WHEEL, data[3]);

    input_sync(dev);
    ...
}
```



Polling input devices

- ▶ The input subsystem provides an API to support simple input devices that *do not raise interrupts* but have to be *periodically scanned or polled* to detect changes in their state.
- ▶ Setting up polling is done using `input_setup_polling()`:

```
int input_setup_polling(struct input_dev *dev, void (*poll_fn)(struct input_dev *dev));
```
- ▶ `poll_fn` is the function that will be called periodically.
- ▶ The polling interval can be set using `input_set_poll_interval()` or `input_set_min_poll_interval()` and `input_set_max_poll_interval()`



- ▶ The main user space interface to *input devices* is the **event interface**
- ▶ Each *input device* is represented as a `/dev/input/event<X>` character device
- ▶ A user space application can use blocking and non-blocking reads, but also `select()` (to get notified of events) after opening this device.
- ▶ Each read will return `struct input_event` structures of the following format:

```
struct input_event {  
    struct timeval time;  
    unsigned short type;  
    unsigned short code;  
    unsigned int value;  
};
```

- ▶ A very useful application for *input device* testing is `evtest`, from <https://cgkit.freedesktop.org/evtest/>



Device-managed allocations



Device managed allocations

- ▶ The `probe()` function is typically responsible for allocating a significant number of resources: memory, mapping I/O registers, registering interrupt handlers, etc.
- ▶ These resource allocations have to be properly freed:
 - In the `probe()` function, in case of failure
 - In the `remove()` function
- ▶ This required a lot of failure handling code that was rarely tested
- ▶ To solve this problem, *device managed* allocations have been introduced.
- ▶ The idea is to associate resource allocation with the `struct device`, and automatically release those resources
 - When the device disappears
 - When the device is unbound from the driver
- ▶ Functions prefixed by `devm_`
- ▶ See [driver-api/driver-model/devres](#) for details



Device managed allocations: memory allocation example

- ▶ Normally done with `kmalloc(size_t, gfp_t)`, released with `kfree(void *)`
- ▶ Device managed with `devm_kmalloc(struct device *, size_t, gfp_t)`

Without devm functions

```
int foo_probe(struct platform_device *pdev)
{
    struct foo_t *foo = kmalloc(sizeof(struct foo_t),
                                GFP_KERNEL);
    /* Register to framework, store
     * reference to framework structure in foo */
    ...
    if (failure) {
        kfree(foo);
        return -EBUSY;
    }
    platform_set_drvdata(pdev, foo);
    return 0;
}

void foo_remove(struct platform_device *pdev)
{
    struct foo_t *foo = platform_get_drvdata(pdev);
    /* Retrieve framework structure from foo
     * and unregister it */
    ...
    kfree(foo);
}
```

With devm functions

```
int foo_probe(struct platform_device *pdev)
{
    struct foo_t *foo = devm_kmalloc(&pdev->dev,
                                     sizeof(struct foo_t),
                                     GFP_KERNEL);
    /* Register to framework, store
     * reference to framework structure in foo */
    ...
    if (failure)
        return -EBUSY;
    platform_set_drvdata(pdev, foo);
    return 0;
}

void foo_remove(struct platform_device *pdev)
{
    struct foo_t *foo = platform_get_drvdata(pdev);
    /* Retrieve framework structure from foo
     * and unregister it */
    ...
    /* foo automatically freed */
}
```



Device managed allocations caveats

- ▶ Cleanup is done when the `struct device` is cleaned up. There is no reference counting or anything like that.
- ▶ Don't use if the allocated memory is used outside of the device node. E.g. if the userspace device file is still open after remove.
- ▶ Be very careful when there are circular references.
- ▶ "Why is `devm_kzalloc()` harmful and what can we do about it", Laurent Pinchart, LPC 2022



Driver data structures and links



Driver-specific Data Structure

- ▶ Each *framework* defines a structure that a device driver must register to be recognized as a device in this framework
 - `struct uart_port` for serial ports, `struct net_device` for network devices, `struct fb_info` for framebuffers, etc.
- ▶ In addition to this structure, the driver usually needs to store additional information about each device
- ▶ This is typically done
 - By subclassing the appropriate framework structure
 - By storing a reference to the appropriate framework structure
 - Or by including your information in the framework structure



Driver-specific Data Structure Examples 1/2

- ▶ i.MX serial driver: `struct imx_port` is a subclass of `struct uart_port`

```
struct imx_port {  
    struct uart_port port;  
    struct timer_list timer;  
    unsigned int old_status;  
    int txirq, rxirq, rtsirq;  
    unsigned int have_rtscts:1;  
    [...]  
};
```




Driver-specific Data Structure Examples 2/2

- ▶ ds1305 RTC driver: `struct ds1305` has a reference to `struct rtc_device`

```
struct ds1305 {  
    struct spi_device    *spi;  
    struct rtc_device    *rtc;  
    [...]  
};  
  
static int ds1305_alarm_irq_enable(struct device *dev, unsigned int enabled)  
{  
    struct ds1305 *ds1305 = dev_get_drvdata(dev);  
    [ .. ]  
}
```

- ▶ rtl8150 network driver: `struct rtl8150` has a reference to `struct net_device` and is allocated within that framework structure.

```
struct rtl8150 {  
    unsigned long flags;  
    struct usb_device *udev;  
    struct tasklet_struct tl;  
    struct net_device *netdev;  
    [...]  
};
```



Links between structures 1/4


- ▶ The framework structure typically contains a `struct device *` pointer that the driver must point to the corresponding `struct device`
 - It's the relationship between the logical device (for example a network interface) and the physical device (for example the USB network adapter)
- ▶ The device structure also contains a `void *` pointer that the driver can freely use.
 - It's often used to link back the device to the higher-level structure from the framework.
 - It allows, for example, from the `struct platform_device` structure, to find the structure describing the logical device



Links between structures 2/4

```
static int serial_imx_probe(struct platform_device *pdev)
{
    struct imx_port *sport; /* per device structure */
    [...]
    sport = devm_kzalloc(&pdev->dev, sizeof(*sport), GFP_KERNEL);
    [...]
    /* setup the link between uart_port and the struct
     * device inside the platform_device */
    sport->port.dev = &pdev->dev;                // Arrow 1
    [...]
    /* setup the link between the struct device inside
     * the platform device to the imx_port structure */
    platform_set_drvdata(pdev, sport);            // Arrow 2
    [...]
    uart_add_one_port(&imx_reg, &sport->port);
}

static int serial_imx_remove(struct platform_device *pdev)
{
    /* retrieve the imx_port from the platform_device */
    struct imx_port *sport = platform_get_drvdata(pdev);
    [...]
    uart_remove_one_port(&imx_reg, &sport->port);
    [...]
}
```





Links between structures 3/4

```
static int ds1305_probe(struct spi_device *spi)
{
    struct ds1305          *ds1305;

    [...]

    /* set up driver data */
    ds1305 = devm_kzalloc(&spi->dev, sizeof(*ds1305), GFP_KERNEL);
    if (!ds1305)
        return -ENOMEM;
    ds1305->spi = spi;                // Arrow 1
    spi_set_drvdata(spi, ds1305);    // Arrow 2


    [...]

    ds1305->rtc = devm_rtc_allocate_device(&spi->dev);
                                   // Arrows 3 and 4

    [...]
}

static int ds1305_remove(struct spi_device *spi)
{
    struct ds1305 *ds1305 = spi_get_drvdata(spi);

    [...]
}
```





Links between structures 4/4

```
static int rtl8150_probe(struct usb_interface *intf,
                        const struct usb_device_id *id)
{
    struct usb_device *udev = interface_to_usbdev(intf);
    rtl8150_t *dev;
    struct net_device *netdev;

    netdev = alloc_etherdev(sizeof(rtl8150_t));
    dev = netdev_priv(netdev);

    [...]

    dev->udev = udev;      // Arrow 1
    dev->netdev = netdev;  // Arrow 2


    [...]

    usb_set_intfdata(intf, dev);      // Arrow 3
    SET_NETDEV_DEV(netdev, &intf->dev); // Arrow 4

    [...]
}

static void rtl8150_disconnect(struct usb_interface *intf)
{
    rtl8150_t *dev = usb_get_intfdata(intf);

    [...]
}
```





Practical lab - Expose the Nunchuk to user space




- ▶ Extend the Nunchuk driver to expose the Nunchuk features to user space applications, as an *input* device.
- ▶ Test the operation of the Nunchuk using `evtest`



Memory Management


© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!






Physical and virtual memory

Illustration on 32 bit systems






Virtual memory organization



- ▶ The top quarter reserved for kernel-space
 - Contains kernel code and core data structures
 - Allocations for loading modules
 - All kernel physical mappings
 - Identical in all address spaces
- ▶ The lower part is a per user process exclusive mapping
 - Process code and data (program, stack, ...)
 - Memory-mapped files
 - Each process has its own address space!
- ▶ The exact virtual mapping in-use is displayed in the kernel log early at boot time



Physical/virtual memory mapping on 32-bit systems






32-bit systems limitations

- ▶ Only less than 1GB memory addressable directly through kernel virtual addresses
- ▶ If more physical memory is present on the platform, part of the memory will not be accessible by kernel space, but can be used by user space
- ▶ To allow the kernel to access more physical memory:
 - Change the 3GB/1GB memory split to 2GB/2GB or 1GB/3GB ([CONFIG_VMSPLIT_2G](#) or [CONFIG_VMSPLIT_1G](#)) $\Rightarrow$ reduce total user memory available for each process
 - Activate *highmem* support if available for your architecture:
 - Allows kernel to map parts of its non-directly accessible memory
 - Mapping must be requested explicitly
 - Limited addresses ranges reserved for this usage
- ▶ See Arnd Bergmann's *4GB by 4GB split* presentation ([video and slides](#)) at Linaro Connect virtual 2020.




Physical/virtual memory mapping on 64-bit systems (4kiB-pages)





User space virtual address space

- ▶ When a process starts, the executable code is loaded in RAM and mapped into the process virtual address space.
- ▶ During execution, additional mappings can be created:
 - Memory allocations
 - Memory mapped files
 - mmap'ed areas
 - ...






Userspace memory allocations

- ▶ Userspace mappings can target the full memory
- ▶ When allocated, memory may not be physically allocated:
 - Kernel uses demand fault paging to allocate the physical page (the physical page is allocated when access to the virtual address generates a page fault)
 - ... or may have been swapped out, which also induces a page fault
 - See the `mlock/mlockall` system calls for workarounds
- ▶ User space memory allocation is allowed to over-commit memory (more than available physical memory) $\Rightarrow$ can lead to out of memory situations.
 - Can be prevented with the use of `/proc/sys/vm/overcommit_*`
- ▶ OOM killer kicks in and selects a process to kill to retrieve some memory. That's better than letting the system freeze.



Kernel memory allocators





Page allocator

- ▶ Appropriate for medium-size allocations
- ▶ A page is usually 4K, but can be made greater in some architectures (sh, mips: 4, 8, 16 or 64 KB, but not configurable in x86 or arm).
- ▶ Buddy allocator strategy, so only allocations of power of two number of pages are possible: 1 page, 2 pages, 4 pages, 8 pages, 16 pages, etc.
- ▶ Typical maximum size is 8192 KB, but it might depend on the kernel configuration.
- ▶ The allocated area is contiguous in the kernel virtual address space, but also maps to physically contiguous pages. It is allocated in the identity-mapped part of the kernel memory space.
 - This means that large areas may not be available or hard to retrieve due to physical memory fragmentation.
 - The *Contiguous Memory Allocator* (CMA) can be used to reserve a given amount of memory at boot (see <https://lwn.net/Articles/486301/>).



Page allocator API

- ▶ `unsigned long get_zeroed_page(gfp_t gfp_mask)`
 - Returns the virtual address of a free page, initialized to zero
 - `gfp_mask`: see the next pages for details.
- ▶ `unsigned long __get_free_page(gfp_t gfp_mask)`
 - Same, but doesn't initialize the contents
- ▶ `unsigned long __get_free_pages(gfp_t gfp_mask, unsigned int order)`
 - Returns the starting virtual address of an area of several contiguous pages in physical RAM, with `order` being $\log_2(\text{number_of_pages})$. Can be computed from the size with the `get_order()` function.
- ▶ `void free_page(unsigned long addr)`
 - Frees one page.
- ▶ `void free_pages(unsigned long addr, unsigned int order)`
 - Frees multiple pages. Need to use the same order as in allocation.



Page allocator flags

The most common ones are:

- ▶ [GFP_KERNEL](#)
 - Standard kernel memory allocation. The allocation may block in order to find enough available memory. Fine for most needs, except in interrupt handler context.
- ▶ [GFP_ATOMIC](#)
 - RAM allocated from code which is not allowed to block (interrupt handlers or critical sections). Never blocks, allows to access emergency pools, but can fail if no free memory is readily available.
- ▶ Others are defined in [include/linux/gfp_types.h](#).
See also the documentation in [core-api/memory-allocation](#)




SLAB allocator 1/2

- ▶ The SLAB allocator allows to create *caches*, which contain a set of objects of the same size. In English, *slab* means *tile*.
- ▶ The object size can be smaller or greater than the page size
- ▶ The SLAB allocator takes care of growing or reducing the size of the cache as needed, depending on the number of allocated objects. It uses the page allocator to allocate and free pages.
- ▶ SLAB caches are used for data structures that are present in many instances in the kernel: directory entries, file objects, network packet descriptors, process descriptors, etc.
 - See `/proc/slabinfo`
- ▶ They are rarely used for individual drivers.
- ▶ See `include/linux/slab.h` for the API



SLAB allocator 2/2





Different SLAB allocators

There are different, but API compatible, implementations of a SLAB allocator in the Linux kernel. A particular implementation is chosen at configuration time.

- ▶ **CONFIG_SLUB**: the default allocator, a good generic choice. It scales well and creates little fragmentation.
- ▶ **CONFIG_SLUB_TINY**: configure SLUB to achieve minimal memory footprint, sacrificing scalability, debugging and other features. Not recommended for systems with more than 16 MB of RAM.

.config - Linux/arm 6.7.0 Kernel Configuration

Choose SLAB allocator

```
SLAB (DEPRECATED)
<X> SLUB (Unqueued Allocator)
```



kmalloc allocator

- ▶ The kmalloc allocator is the general purpose memory allocator in the Linux kernel
- ▶ For small sizes, it relies on generic SLAB caches, named `kmalloc-XXX` in `/proc/slabinfo`
- ▶ For larger sizes, it relies on the page allocator
- ▶ The allocated area is guaranteed to be physically contiguous
- ▶ The allocated area size is rounded up to the size of the smallest SLAB cache in which it can fit (while using the SLAB allocator directly allows to have more flexibility)
- ▶ It uses the same flags as the page allocator (`GFP_KERNEL`, `GFP_ATOMIC`, etc.) with the same semantics.
- ▶ Maximum sizes, on x86 and arm (see <https://j.mp/YIGq6W>):
 - Per allocation: 4 MB
 - Total allocations: 128 MB
- ▶ Should be used as the primary allocator unless there is a strong reason to use another one.



kmalloc API 1/2

- ▶ `#include <linux/slab.h>`
- ▶ `void *kmalloc(size_t size, gfp_t flags);`
 - Allocate `size` bytes, and return a pointer to the area (virtual address)
 - `size`: number of bytes to allocate
 - `flags`: same flags as the page allocator
- ▶ `void kfree(const void *objp);`
 - Free an allocated area
- ▶ Example: (`drivers/infiniband/core/cache.c`)

```
struct ib_port_attr *tprops;  
tprops = kmalloc(sizeof *tprops, GFP_KERNEL);  
...  
kfree(tprops);
```



- ▶ `void *kzalloc(size_t size, gfp_t flags);`
 - Allocates a zero-initialized buffer
- ▶ `void *kcalloc(size_t n, size_t size, gfp_t flags);`
 - Allocates memory for an array of `n` elements of `size` size, and zeroes its contents.
- ▶ `void *krealloc(const void *p, size_t new_size, gfp_t flags);`
 - Changes the size of the buffer pointed by `p` to `new_size`, by reallocating a new buffer and copying the data, unless `new_size` fits within the alignment of the existing buffer.



devm_kmalloc functions

Allocations with automatic freeing when the corresponding device or module is unprobed.

- ▶ `void *devm_kmalloc(struct device *dev, size_t size, gfp_t gfp);`
- ▶ `void *devm_kzalloc(struct device *dev, size_t size, gfp_t gfp);`
- ▶ `void *devm_kcalloc(struct device *dev, size_t n, size_t size, gfp_t flags);`
- ▶ `void *devm_kfree(struct device *dev, void *p);`

Useful to immediately free an allocated buffer

For use in `probe()` functions, in which you have access to a `struct device` structure.



vmalloc allocator

- ▶ The `vmalloc()` allocator can be used to obtain memory zones that are contiguous in the virtual addressing space, but not made out of physically contiguous pages.
- ▶ The requested memory size is rounded up to the next page (not efficient for small allocations).
- ▶ The allocated area is in the kernel space part of the address space, but outside of the identically-mapped area
- ▶ Allocations of fairly large areas is possible (almost as big as total available memory, see <https://j.mp/YIGq6W> again), since physical memory fragmentation is not an issue.
- ▶ Not suitable for DMA purposes.
- ▶ API in `include/linux/vmalloc.h`
 - `void *vmalloc(unsigned long size);`
 - Returns a virtual address
 - `void vfree(void *addr);`



Kernel memory debugging

- ▶ KASAN (*Kernel Address Sanitizer*)
 - Dynamic memory error detector, to find use-after-free and out-of-bounds bugs.
 - Available on most architectures
 - See [dev-tools/kasan](#) for details.
- ▶ KFENCE (*Kernel Electric Fence*)
 - A low overhead alternative to KASAN, trading performance for precision. Meant to be used in production systems.
 - Available on most architectures.
 - See [dev-tools/kfence](#) for details.
- ▶ Kmemleak
 - Dynamic checker for memory leaks
 - This feature is available for all architectures.
 - See [dev-tools/kmemleak](#) for details.

KASAN and Kmemleak have a significant overhead. Only use them in development!



Kernel memory management: resources

Virtual memory and Linux, Alan Ott and Matt Porter, 2016

Great and much more complete presentation about this topic

<https://bit.ly/2Af1G2i> (video: <https://bit.ly/2Bwwv0C>)

Kernel Virtual Addresses (Small Mem)

Virtual Address Space

Physical Address Space

Kernel Virtual Addresses 0xFFFFFFFF (4GB)

Kernel Logical Addresses

PAGE_OFFSET

Userspace Addresses

0x00000000

Physical RAM

Embedded Linux Conference Europe


OpenIoT Summit Europe

22:29 / 51:18



I/O Memory


© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Memory-Mapped I/O

- ▶ Same address bus to address memory and I/O device registers
- ▶ Access to the I/O device registers using regular instructions
- ▶ Most widely used I/O method across the different architectures supported by Linux



Physical Memory
address space, accessed with
normal load/store instructions



Requesting I/O memory

- ▶ Tells the kernel which driver is using which I/O registers
- ▶ `struct resource *request_mem_region(unsigned long start, unsigned long len, char *name);`
- ▶ `void release_mem_region(unsigned long start, unsigned long len);`
- ▶ Allows to prevent other drivers from requesting the same I/O registers, but is purely voluntary.



/proc/iomem example - ARM 32 bit (BeagleBone Black, Linux 5.11)

```
40300000-4030ffff : 40300000.sram sram@0
44e00c00-44e00cff : 44e00c00.prm prm@c00
44e00d00-44e00dff : 44e00d00.prm prm@d00
44e00e00-44e00eff : 44e00e00.prm prm@e00
44e00f00-44e00fff : 44e00f00.prm prm@f00
44e01000-44e010ff : 44e01000.prm prm@1000
44e01100-44e011ff : 44e01100.prm prm@1100
44e01200-44e012ff : 44e01200.prm prm@1200
44e07000-44e07fff : 44e07000.gpio gpio@0
44e09000-44e0901f : serial
44e0b000-44e0bfff : 44e0b000.i2c i2c@0
44e10800-44e10a37 : pinctrl-single
44e10f90-44e10fcf : 44e10f90.dma-router dma-router@f90
48024000-48024fff : 48024000.serial serial@0
48042000-480423ff : 48042000.timer timer@0
48044000-480443ff : 48044000.timer timer@0
48046000-480463ff : 48046000.timer timer@0
48048000-480483ff : 48048000.timer timer@0
4804a000-4804a3ff : 4804a000.timer timer@0
4804c000-4804cfff : 4804c000.gpio gpio@0
48060000-48060fff : 48060000.mmc mmc@0
4819c000-4819cfff : 4819c000.i2c i2c@0
481a8000-481a8fff : 481a8000.serial serial@0
481ac000-481acfff : 481ac000.gpio gpio@0
481ae000-481aefff : 481ae000.gpio gpio@0
481d8000-481d8fff : 481d8000.mmc mmc@0
49000000-4900ffff : 49000000.dma edma3_cc
4a100000-4a1007ff : 4a100000.ethernet ethernet@0
4a101200-4a1012ff : 4a100000.ethernet ethernet@0
80000000-9fdfffff : System RAM
80008000-80cfffff : Kernel code
80e00000-80f3d807 : Kernel data
```




Mapping I/O memory in virtual memory

- ▶ Load/store instructions work with virtual addresses
- ▶ To access I/O memory, drivers need to have a virtual address that the processor can handle, because I/O memory is not mapped by default in virtual memory.
- ▶ The `ioremap` function satisfies this need:


```
#include <linux/io.h>
```

```
void __iomem *ioremap(phys_addr_t phys_addr, unsigned long size);  
void iounmap(void __iomem *addr);
```

- ▶ Caution: check that `ioremap()` doesn't return a NULL address!



ioremap()



`ioremap(0xAFFEBC00, 4096) = 0xCDEFA000`



Managed API

Using `request_mem_region()` and `ioremap()` in device drivers is now deprecated. You should use the below "managed" functions instead, which simplify driver coding and error handling:

- ▶ `devm_ioremap()`, `devm_iounmap()`
- ▶ `devm_ioremap_resource()`
 - Takes care of both the request and remapping operations!
- ▶ `devm_platform_ioremap_resource()`
 - Takes care of `platform_get_resource()`, `request_mem_region()` and `ioremap()`
 - Caution: unlike the other `devm_` functions, its first argument is of type `struct platform_device`, not a pointer to `struct device`:

```
base = devm_platform_ioremap_resource(pdev, 0);  
if (IS_ERR(base))  
    return PTR_ERR(base);
```



Accessing MMIO devices: using accessor functions

- ▶ Care must be taken when accessing MMIO registers
 - Memory mapped I/O can be weakly ordered
 - Developer's responsibility to enforcing proper ordering
 - Endianness of the device may be different than the CPU endianness
 - Endianness conversions might be required
 - Directly reading from or writing to addresses returned by `ioremap()` (*pointer dereferencing*) may not work on some architectures.
- ▶ A family of architecture-independent accessor functions are available covering most needs.
 - A few architecture-specific accessor functions also exists.



Ordering

- ▶ Reads/writes to MMIO-mapped registers of given device are done in program order
- ▶ However reads/writes to RAM can be re-ordered between themselves, and between MMIO-mapped read/writes
- ▶ Some of the accessor functions include memory barriers to help with this:
 - Write operation starts with a write memory barrier which prior writes cannot cross
 - Read operation ends with a read memory barrier which guarantees the ordering with regard to the subsequent reads
- ▶ Sometimes compiler/CPU reordering is not an issue, in this case the code may be optimized by dropping the memory barriers, using the raw or relaxed helpers
- ▶ See [Documentation/memory-barriers.txt](#)



MMIO access functions

- ▶ `read[b/w/l/q]` and `write[b/w/l/q]` for access to little-endian devices, includes memory barriers
- ▶ `ioread[8/16/32/64]` and `iowrite[8/16/32/64]` are very similar to `read/write` but also work with port I/O (not covered in the course), includes memory barriers
- ▶ `ioread[8/16/32/64]be` and `iowrite[8/16/32/64]be` for access to big-endian devices, includes memory barriers
- ▶ `__raw_read[b/w/l/q]` and `__raw_write[b/w/l/q]` for raw access: no endianness conversion, no memory barriers
- ▶ `read[b/w/l/q]_relaxed` and `write[b/q/l/w]_relaxed` for access to little-endian devices, without memory barriers
- ▶ All functions work on a `void __iomem *`



MMIO access functions summary

Name	Device endianness	Memory barriers
read/write	little	yes
ioread/iowrite	little	yes
ioreadbe/iowritebe	big	yes
__raw_read/__raw_write	native	no
read_relaxed/write_relaxed	little	no

More details at <https://docs.kernel.org/driver-api/device-io.html>



- ▶ Used to provide user space applications with direct access to physical addresses.
- ▶ Usage: open `/dev/mem` and read or write at given offset. What you read or write is the value at the corresponding physical address.
- ▶ Used by applications such as the X server to write directly to device memory.
- ▶ Easy to use from a shell with the `devmem2` program
- ▶ For security reasons, on `x86`, `arm`, `arm64`, `riscv`, `powerpc`, `parisc`, `s390`:
 - `CONFIG_STRICT_DEVMEM` restricts `/dev/mem` to non-RAM addresses (from v5.12)
 - `CONFIG_IO_STRICT_DEVMEM` goes beyond and only allows to access *idle* I/O ranges (not appearing in `/proc/iomem`).




- ▶ Add UART devices to the board device tree
- ▶ Access I/O registers to control the device and send first characters to it.



The misc subsystem

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!






Why a *misc* subsystem?

- ▶ The kernel offers a large number of **frameworks** covering a wide range of device types: input, network, video, audio, etc.
 - These frameworks allow to factorize common functionality between drivers and offer a consistent API to user space applications.
- ▶ However, there are some devices that **really do not fit in any of the existing frameworks**.
 - Highly customized devices implemented in a FPGA, or other weird devices for which implementing a complete framework is not useful.
- ▶ The drivers for such devices could be implemented directly as raw *character drivers* (with `cdev_init()` and `cdev_add()`).
- ▶ But there is a subsystem that makes this work a little bit easier: the **misc subsystem**.
 - It is really only a **thin layer** above the *character driver* API.
 - Another advantage is that devices are integrated in the Device Model (device files appearing in *devtmpfs*, which you don't have with raw character devices).



Misc subsystem diagram





Misc subsystem API (1/2)

- ▶ The misc subsystem API mainly provides two functions, to register and unregister a **single** *misc device*:
 - `int misc_register(struct miscdevice * misc);`
 - `void misc_deregister(struct miscdevice *misc);`
- ▶ A *misc device* is described by a `struct miscdevice` structure:

```
struct miscdevice {  
    int minor;  
    const char *name;  
    const struct file_operations *fops;  
    struct list_head list;  
    struct device *parent;  
    struct device *this_device;  
    const char *nodename;  
    umode_t mode;  
};
```



Misc subsystem API (2/2)

The main fields to be filled in `struct miscdevice` are:

- ▶ `minor`, the minor number for the device, or `MISC_DYNAMIC_MINOR` to get a minor number automatically assigned.
- ▶ `name`, name of the device, which will be used to create the device node if *devtmpfs* is used.
- ▶ `fops`, pointer to the same `struct file_operations` structure that is used for raw character drivers, describing which functions implement the *read*, *write*, *ioctl*, etc. operations.
- ▶ `parent`, pointer to the `struct device` of the underlying “physical” device (platform device, I2C device, etc.)



User space API for misc devices

- ▶ *misc devices* are regular character devices
- ▶ The operations they support in user space depends on the operations the kernel driver implements:
 - The `open()` and `close()` system calls to open/close the device.
 - The `read()` and `write()` system calls to read/write to/from the device.
 - The `ioctl()` system call to call some driver-specific operations.



Practical lab - Output-only serial port driver




- ▶ Extend the driver started in the previous lab by registering it into the *misc* subsystem.
- ▶ Implement serial output functionality through the *misc* subsystem.
- ▶ Test serial output using user space applications.



Processes, scheduling and interrupts

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Processes and scheduling




Process, thread?

- ▶ Confusion about the terms *process*, *thread* and *task*
- ▶ In UNIX, a process is created using `fork()` and is composed of
 - An address space, which contains the program code, data, stack, shared libraries, etc.
 - A single thread, which is the only entity known by the scheduler.
- ▶ Additional threads can be created inside an existing process, using `pthread_create()`
 - They run in the same address space as the initial thread of the process
 - They start executing a function passed as argument to `pthread_create()`




Process, thread: kernel point of view

- ▶ In kernel space, each thread running in the system is represented by a structure of type `struct task_struct`
- ▶ From a scheduling point of view, it makes no difference between the initial thread of a process and all additional threads created dynamically using `pthread_create()`



Process after `fork()`



Same process after
`pthread_create()`




Relation between execution mode, address space and context

- ▶ When speaking about *process* and *thread*, these concepts need to be clarified:
 - *Mode* is the level of privilege allowing to perform some operations:
 - *Kernel Mode*: in this level CPU can perform any operation allowed by its architecture; any instruction, any I/O operation, any area of memory accessed.
 - *User Mode*: in this level, certain instructions are not permitted (especially those that could alter the global state of the machine), some memory areas cannot be accessed.
 - Linux splits its *address space* in *kernel space* and *user space*
 - *Kernel space* is reserved for code running in *Kernel Mode*.
 - *User space* is the place where applications execute (accessible from *Kernel Mode*).
 - *Context* represents the current state of an execution flow.
 - The *process context* can be seen as the content of the registers associated to this process: execution register, stack register...
 - The *interrupt context* replaces the *process context* when the interrupt handler is executed.




A thread life





Execution of system calls




The execution of system calls takes place in the context of the thread requesting them.



Sleeping



Sleeping



Sleeping is needed when a userspace or kernelspace thread is waiting for data.



How to sleep with a wait queue 1/3

- ▶ Must declare a wait queue, which will be used to store the list of threads waiting for an event
- ▶ Dynamic queue declaration:
 - Typically one queue per device managed by the driver
 - It's convenient to embed the wait queue inside a per-device data structure.
 - Example from [drivers/net/ethernet/marvell/mvmdio.c](#):

```
struct orion_mdio_dev {  
    ...  
    wait_queue_head_t smi_busy_wait;  
};  
struct orion_mdio_dev *dev;  
...  
init_waitqueue_head(&dev->smi_busy_wait);
```
- ▶ Static queue declaration:
 - Using a global variable when a global resource is sufficient
 - `DECLARE_WAIT_QUEUE_HEAD(module_queue);`



How to sleep with a wait queue 2/3

Several ways to make a thread sleep

- ▶ `void wait_event(queue, condition);`
 - Sleeps until the wait queue is woken up **and** the given C expression is true. Caution: once entered, the thread sleep cannot be interrupted!
- ▶ `int wait_event_killable(queue, condition);`
 - Can be interrupted, but only by a *fatal* signal (`SIGKILL`). Returns `-ERESTARTSYS` if interrupted.
- ▶ `int wait_event_interruptible(queue, condition);`
 - The most common variant
 - Can be interrupted by any signal. Returns `-ERESTARTSYS` if interrupted.



How to sleep with a wait queue 3/3

- ▶ `int wait_event_timeout(queue, condition, timeout);`
 - Also stops sleeping when the thread is woken up **or** the timeout expired (a timer is used).
 - Returns 0 if the timeout elapsed, non-zero if the condition was met.
- ▶ `int wait_event_interruptible_timeout(queue, condition, timeout);`
 - Same as above, interruptible.
 - Returns 0 if the timeout elapsed, `-ERESTARTSYS` if interrupted, positive value if the condition was met.



How to sleep with a wait queue - Example

```
sig = wait_event_interruptible(ibmvtpm->wq,  
                               !ibmvtpm->tpm_processing_cmd);  
if (sig)  
    return -EINTR;
```

From [drivers/char/tpm/tpm_ibmvtpm.c](#)



Waking up!

Typically done by interrupt handlers when data sleeping threads are waiting for become available.

- ▶ `wake_up(&queue);`
 - Wakes up all threads in the wait queue
- ▶ `wake_up_interruptible(&queue);`
 - Wakes up all threads waiting in an interruptible sleep on the given queue




Exclusive vs. non-exclusive

- ▶ `wait_event_interruptible()` puts a thread in a non-exclusive wait.
 - All non-exclusive threads are woken up by `wake_up()` / `wake_up_interruptible()`
- ▶ `wait_event_interruptible_exclusive()` puts a thread in an exclusive wait.
 - `wake_up()` / `wake_up_interruptible()` wakes up all non-exclusive threads and only one exclusive thread
 - `wake_up_all()` / `wake_up_interruptible_all()` wakes up all non-exclusive and all exclusive threads
- ▶ Exclusive sleeps are useful to avoid waking up multiple threads when only one will be able to *consume* the event.
- ▶ Non-exclusive sleeps are useful when the event can *benefit* to multiple threads.



Sleeping and waking up - Implementation



The scheduler doesn't keep evaluating the sleeping condition!

- ▶ `wait_event(queue, cond);`

The thread is put in the `TASK_UNINTERRUPTIBLE` state.

- ▶ `wake_up(&queue);`

All threads waiting in `queue` are woken up, so they get scheduled later and have the opportunity to evaluate the condition again and go back to sleep if it is not met.

See [include/linux/wait.h](#) for implementation details.



How to sleep with completions 1/2

- ▶ Use [wait_for_completion\(\)](#) when no particular condition must be enforced at the time of the wake-up
 - Leverages the power of wait queues
 - Simplifies its use
 - Highly efficient using low level scheduler facilities
- ▶ Preparation of the completion structure:
 - Static declaration and initialization:
`DECLARE_COMPLETION(setup_done);`
 - Dynamic declaration:
`init_completion(&object->setup_done);`
 - The completion object should get a meaningful name (eg. not just “done”).
- ▶ Ready to be used by signal consumers and providers as soon as the completion object is initialized
- ▶ See [include/linux/completion.h](#) for the full API
- ▶ Internal documentation at [scheduler/completion](#)



How to sleep with completions 2/2

▶ Enter a wait state with

```
void wait_for_completion(struct completion *done)
```

- All `wait_event()` flavors are also supported, such as:
`wait_for_completion_timeout()`,
`wait_for_completion_interruptible()` / `_timeout()`,
`wait_for_completion_killable()` / `_timeout()`, etc.

▶ Wake up consumers with

```
void complete(struct completion *done)
```

- Several calls to `complete()` are valid, they will wake up the same number of threads waiting on this object (acts as a FIFO).
- A single `complete_all()` call would wake up all present and future threads waiting on this completion object

▶ Reset the counter with

```
void reinit_completion(struct completion *done)
```

- Resets the number of “done” completions still pending
- Mind not to call `init_completion()` twice, which could confuse the enqueued threads



- ▶ Use helpers which implement software loops or use hardware timers
 - `udelay()` waste CPU cycles in order to save a couple of context switches, suitable for $\leq 10\mu\text{s}$ or in atomic situations
 - `usleep()/usleep_range()/msleep()` put the thread in sleep for a given amount of micro/milliseconds (not suitable in atomic contexts)
 - If in doubt, use `fsleep()`, which will use the more suitable internal function depending on the period you've asked!



Waiting when hardware is involved

- ▶ When hardware is involved in the waiting process
 - but there is no interrupt available
 - or because a context switch would be too expensive
- ▶ Specific polling I/O accessors may be used:
 - Exhaustive list in `include/linux/iopoll.h`

```
int read[bwlq]_poll_timeout[_atomic](addr, val, cond,  
                                     delay_us, timeout_us)
```

 - `addr`: I/O memory location
 - `val`: Content of the register pointed with
 - `cond`: Boolean condition based on `val`
 - `delay_us`: Polling delay between reads
 - `timeout_us`: Timeout delay after which the operation fails and returns `-ETIMEDOUT`
 - `_atomic` variant uses `udelay()` instead of `usleep()`.
- ▶ Avoid implementing custom busy-wait loops if possible



Interrupt Management



Registering an interrupt handler 1/2

The *managed* API is recommended:

```
int devm_request_irq(struct device *dev, unsigned int irq, irq_handler_t handler,  
                    unsigned long irq_flags, const char *devname, void *dev_id);
```

- ▶ device for automatic freeing at device or module release time.
- ▶ irq is the requested IRQ channel. For platform devices, use `platform_get_irq()` to retrieve the interrupt number.
- ▶ handler is a pointer to the IRQ handler function
- ▶ irq_flags are option masks (see next slide)
- ▶ devname is the registered name (for `/proc/interrupts`). For platform drivers, good idea to use `pdev->name` which allows to distinguish devices managed by the same driver (example: `44e0b000.i2c`).
- ▶ dev_id is an opaque pointer. It can typically be used to pass a pointer to a per-device data structure. It cannot be NULL as it is used as an identifier for freeing interrupts on a shared line.



Registering an interrupt handler 2/2

Here are the most frequent `irq_flags` bit values in drivers (can be combined):

- ▶ **IRQF_SHARED**: interrupt channel can be shared by several devices.
 - When an interrupt is received, all the interrupt handlers registered on the same interrupt line are called.
 - This requires a hardware status register telling whether an IRQ was raised or not.
- ▶ **IRQF_ONESHOT**: for use by threaded interrupts (see next slides). Keeping the interrupt line disabled until the thread function has run.



Interrupt handler constraints

- ▶ No guarantee in which address space the system will be in when the interrupt occurs: can't transfer data to and from user space.
- ▶ Interrupt handler execution is managed by the CPU, not by the scheduler. Handlers can't run actions that may sleep, because there is nothing to resume their execution. In particular, need to allocate memory with `GFP_ATOMIC`.
- ▶ Interrupt handlers are run with all interrupts disabled on the local CPU (see <https://lwn.net/Articles/380931>). Therefore, they have to complete their job quickly enough, to avoiding blocking interrupts for too long.



/proc/interrupts on Raspberry Pi 2 (ARM, Linux 4.19)

```

      CPU0      CPU1      CPU2      CPU3
17:    1005317          0          0          0  ARMCTRL-level  1 Edge    3f00b880.mailbox
18:         36          0          0          0  ARMCTRL-level  2 Edge    VCHIQ doorbell
40:          0          0          0          0  ARMCTRL-level 48 Edge    bcm2708_fb DMA
42:    427715          0          0          0  ARMCTRL-level 50 Edge    DMA IRQ
56:   478426356          0          0          0  ARMCTRL-level 64 Edge    dwc_otg, dwc_otg_pcd, dwc_otg_hcd:usb1
80:    411468          0          0          0  ARMCTRL-level 88 Edge    mmc0
81:         502          0          0          0  ARMCTRL-level 89 Edge    uart-pl011
161:          0          0          0          0  bcm2836-timer  0 Edge    arch_timer
162:   10963772   6378711   16583353   6406625  bcm2836-timer  1 Edge    arch_timer
165:          0          0          0          0  bcm2836-pmu    9 Edge    arm-pmu
FIQ:
IPI0:          0          0          0          0  CPU wakeup interrupts
IPI1:          0          0          0          0  Timer broadcast interrupts
IPI2:   2625198   4404191   7634127   3993714  Rescheduling interrupts
IPI3:     3140     56405     49483     59648  Function call interrupts
IPI4:          0          0          0          0  CPU stop interrupts
IPI5:   2167923   477097    5350168   412699  IRQ work interrupts
IPI6:          0          0          0          0  completion interrupts
Err:          0
```

Note: interrupt numbers shown on the left-most column are virtual numbers when the Device Tree is used. The physical interrupt numbers can be found in `/sys/kernel/debug/irq/irqs/<nr>` files when `CONFIG_GENERIC_IRQ_DEBUGFS=y`.



Interrupt handler prototype

- ▶ `irqreturn_t foo_interrupt(int irq, void *dev_id)`
 - `irq`, the IRQ number
 - `dev_id`, the per-device pointer that was passed to `devm_request_irq()`
- ▶ Return value
 - `IRQ_HANDLED`: recognized and handled interrupt
 - `IRQ_NONE`: used by the kernel to detect spurious interrupts, and disable the interrupt line if none of the interrupt handlers has handled the interrupt.
 - `IRQ_WAKE_THREAD`: handler requests to wake the handler thread (see next slides)



Typical interrupt handler's job

- ▶ Acknowledge the interrupt to the device (otherwise no more interrupts will be generated, or the interrupt will keep firing over and over again)
- ▶ Read/write data from/to the device
- ▶ Wake up any process waiting for such data, typically on a per-device wait queue:
`wake_up_interruptible(&device_queue);`



Top half and bottom half processing


- ▶ Splitting the execution of interrupt handlers in 2 parts is sometimes needed
 - Hard IRQ handlers execute with all interrupts on the local CPU masked
 - No interrupt nesting mechanism, higher priority IRQs will be delayed
 - Need to block/sleep (eg. accessing I2C devices may involve sleeping)
- ▶ Top half
 - This is the real interrupt handler, which should complete as quickly as possible since all interrupts are disabled. It takes the data out of the device and if substantial post-processing is needed, schedules a bottom half to handle it.
- ▶ Bottom half
 - Is the general name for various mechanisms which allow to postpone the handling of interrupt-related work.
 - Implemented in Linux as softirqs, threaded handlers and workqueues
 - And yet, the abbreviation "bh" often means "softirqs" for historical reasons!



- ▶ Softirq handlers are callbacks executed once all interrupt handlers have completed, before the kernel resumes scheduling processes
 - They execute with all interrupts enabled
 - They run before the scheduler is in control, so sleeping is not allowed
 - A softirq handler can run simultaneously on multiple CPUs
- ▶ The number of softirqs is fixed, softirqs are not used by drivers, but by kernel subsystems (network, etc.). The list of softirqs is defined in [include/linux/interrupt.h](#)
- ▶ To avoid starving the system, there is typically a softirq budget (by default, softirqs can run 10 times in a row for a maximum time of 2ms) after which the callbacks are run in process context in a per CPU `ksoftirqd/N` kernel thread.
- ▶ A Tasklet is a deprecated mechanism which was dedicated to device drivers and was implemented on top of softirqs.



Softirq execution flow





Example usage of softirqs: NAPI

- ▶ Interface in the Linux kernel used for interrupt mitigation in network drivers
- ▶ Principle: when the network traffic exceeds a given threshold ("budget"), disable network interrupts and consume incoming packets through a polling function, instead of processing each new packet with an interrupt.
- ▶ This reduces overhead due to interrupts and yields better network throughput.
- ▶ The polling function is run by `napi_schedule()`, which uses `NET_RX_SOFTIRQ`.
- ▶ See https://en.wikipedia.org/wiki/New_API for details
- ▶ See also our commented network driver on <https://bootlin.com/pub/drivers/r6040-network-driver-with-comments.c>



Threaded interrupts

- ▶ It is possible to associate a threaded handler to a hard IRQ handler
 - The hard IRQ handler triggers it by returning `IRQ_WAKE_THREAD`
 - The threaded handler is executed inside a thread, in process context
- ▶ The interrupt line may be kept masked if the handler registration happened with `IRQF_ONESHOT`
- ▶ Sleeping/blocking is allowed in the threaded handler
 - The priority of the `irq/<nb>-<name>` thread can be tuned!
- ▶ Heavily used by `PREEMPT_RT`

```
int devm_request_threaded_irq(struct device *dev, unsigned int irq,
                             irq_handler_t handler, irq_handler_t thread_fn,
                             unsigned long flags, const char *name,
                             void *dev);
```

- ▶ handler: “hard IRQ” handler
- ▶ thread_fn: executed in a thread



Workqueues

- ▶ Workqueues are a general mechanism for deferring work. It is not limited in usage to handling interrupts.
 - Typically used for background jobs.
- ▶ Functions registered to run in workqueues are called works:
 - They can be created with the macro `INIT_WORK()`
 - When scheduled, they become threads (called workers) running in process context, which means:
 - All interrupts are enabled
 - Sleeping is allowed
 - Works can be queued on:
 - The default workqueue, with `schedule_work()`
 - A workqueue allocated by the subsystem or the drivers, with `alloc_workqueue()`
- ▶ The complete API is in `include/linux/workqueue.h`
- ▶ Example (`drivers/crypto/atmel-i2c.c`):

```
INIT_WORK(&work_data->work, atmel_i2c_work_handler);
schedule_work(&work_data->work);
```



Interrupt and deferred mechanisms execution constraints summary

Mechanism	Context	IRQs	Priority tuning	Can sleep	Typical Use
Hard IRQ	Interrupt	Disabled (local CPU)	No, FIFO	No	Fast handling, gen. purpose
Softirq	Interrupt	Enabled	No, fixed priorities	No	Fast handling, net, timers, RCU
Softirq	Process (ksoftirqd)	Enabled	One kthread per CPU	No	Softirq overflow
Threaded IRQ	Process (irq thread)	Enabled (IRQF_ONESHOT)	Yes	Yes	Interrupt handling needing to block
Workqueue	Process (worker)	Enabled	Yes, per queue	Yes	General purpose background worker



Interrupt management summary

- ▶ Device driver
 - In the `probe()` function, for each device, use `devm_request_irq()` to register an interrupt handler for the device's interrupt channel.
- ▶ Interrupt handler
 - Called when an interrupt is raised.
 - Acknowledge the interrupt
 - If needed, trigger a bottom half mechanism (typically taking care of handling data)
 - Wake up processes waiting for the data on a per-device queue
- ▶ Device driver
 - In the `remove()` function, for each device, the interrupt handler is automatically unregistered.



- ▶ Adding read capability to the character driver developed earlier.
- ▶ Register an interrupt handler for each device.
- ▶ Waiting for data to be available in the read file operation.
- ▶ Waking up the code when data are available from the devices.




Concurrent Access to Resources: Locking

© Copyright 2004-2026, Bootlin.

Creative Commons BY-SA 3.0 license.

Corrections, suggestions, contributions and translations are welcome!






Sources of concurrency issues

- ▶ In terms of concurrency, the kernel has the same constraint as a multi-threaded program: its state is global and visible in all executions contexts
- ▶ Concurrency arises because of
 - *Interrupts*, which interrupts the current thread to execute an interrupt handler. They may be using shared resources (memory addresses, hardware registers...)
 - *Kernel preemption*, if enabled, causes the kernel to switch from the execution of one thread to another. They may be using shared resources.
 - *Multiprocessing*, in which case code is really executed in parallel on different processors, and they may be using shared resources as well.
- ▶ The solution is to keep as much local state as possible and for the shared resources that can't be made local (such as hardware ones), use locking.



Concurrency protection with locks





mutex = **mutual exclusion**

- ▶ The kernel's main locking primitive. It's a *binary lock*. Note that *counting locks* (*semaphores*) are also available, but used 30x less frequently.
- ▶ The process requesting the lock blocks when the lock is already held. Mutexes can therefore only be used in contexts where sleeping is allowed.
- ▶ Mutex definition:
 - `#include <linux/mutex.h>`
- ▶ Initializing a mutex statically (unusual case):
 - `DEFINE_MUTEX(name);`
- ▶ Or initializing a mutex dynamically (the usual case, on a per-device basis):
 - `void mutex_init(struct mutex *lock);`




Locking and unlocking mutexes 1/2

- ▶ `void mutex_lock(struct mutex *lock);`
 - Tries to lock the mutex, sleeps otherwise.
 - Caution: cannot be interrupted, resulting in processes you cannot kill!
- ▶ `int mutex_lock_killable(struct mutex *lock);`
 - Same, but can be interrupted by a fatal (`SIGKILL`) signal. If interrupted, returns a non zero value and doesn't hold the lock. Test the return value!!!
- ▶ `int mutex_lock_interruptible(struct mutex *lock);`
 - Same, but can be interrupted by any signal.
- ▶ `void mutex_unlock(struct mutex *lock);`
 - Releases the lock. Do it as soon as you leave the critical section.



Spinlocks

- ▶ Locks to be used for code that is not allowed to sleep (interrupt handlers), or that doesn't want to sleep (critical sections). Be very careful not to call functions which can sleep!
- ▶ Originally intended for multiprocessor systems
- ▶ Spinlocks never sleep and keep spinning in a loop until the lock is available.
- ▶ The critical section protected by a spinlock is not allowed to sleep.






The spinlock API

- ▶ Spinlocks can be initialized:
 - Statically (unusual)
 - `DEFINE_SPINLOCK(my_lock);`
 - Dynamically (the usual case, on a per-device basis)
 - `void spin_lock_init(spinlock_t *lock);`
- ▶ They can be acquired and released with:
 - `void spin_lock(spinlock_t *lock);`
 - Used for locking in process context (critical sections in which you do not want to sleep) as well as atomic sections.
 - `void spin_unlock(spinlock_t *lock);`



Using spinlocks 1/2

- ▶ Manipulating spinlocks implies some care:




- ▶ So, kernel preemption on the local CPU is disabled. We need to avoid deadlocks (and unbounded latencies) because of preemption from processes that want to get the same lock.
- ▶ Disabling kernel preemption also disables migration to avoid the same kind of issue as pictured above from happening.



Using spinlocks 2/2

- ▶ We also need to avoid deadlocks because of interrupts that could want to get the same lock:



- ▶ `void spin_lock_irqsave(spinlock_t *lock, unsigned long flags);`
- ▶ `void spin_unlock_irqrestore(spinlock_t *lock, unsigned long flags);`
 - Disables/restores IRQs on the local CPU.
 - Typically used when the lock can be accessed in both process and interrupt context.



Using spinlocks 3/3

- ▶ `void spin_lock_bh(spinlock_t *lock);`
- ▶ `void spin_unlock_bh(spinlock_t *lock);`
 - Disables software interrupts, but not hardware ones.
 - Useful to protect shared data accessed in process context and in a soft interrupt (*bottom half*).
 - No need to disable hardware interrupts in this case.
- ▶ Note that reader/writer spinlocks also exist, allowing for multiple simultaneous readers.



Spinlock example

- ▶ From `drivers/tty/serial/uartlite.c`
- ▶ Spinlock structure embedded into `struct uart_port`

```
struct uart_port {  
    spinlock_t lock;  
    /* Other fields */  
};
```

- ▶ Spinlock taken/released with protection against interrupts


```
static unsigned int ulite_tx_empty(struct uart_port *port) {  
    unsigned long flags;  
  
    spin_lock_irqsave(&port->lock, flags);  
    /* Do something */  
    spin_unlock_irqrestore(&port->lock, flags);  
}
```




More deadlock situations

They can lock up your system. Make sure they never happen!

Rule 1: don't call a function that can try to get access to the same lock



Rule 2: if you need multiple locks, always acquire them in the same order!





Debugging locking

- ▶ Lock debugging: prove locking correctness
 - [CONFIG_PROVE_LOCKING](#)
 - Adds instrumentation to kernel locking code
 - Detect violations of locking rules during system life, such as:
 - Locks acquired in different order (keeps track of locking sequences and compares them).
 - Spinlocks acquired in interrupt handlers and also in process context when interrupts are enabled.
 - Not suitable for production systems but acceptable overhead in development.
 - See [locking/lockdep-design](#) for details
- ▶ [CONFIG_DEBUG_ATOMIC_SLEEP](#) allows to detect code that incorrectly sleeps in atomic section (while holding lock typically).
 - Warning displayed in `dmesg` in case of such violation.



Concurrency issues

- ▶ Kernel Concurrency SANitizer framework
- ▶ `CONFIG_KCSAN`, introduced in Linux 5.8.
- ▶ Dynamic race detector relying on compile time instrumentation.
- ▶ Can find concurrency issues (mainly data races) in your system.
- ▶ See dev-tools/kcsan and <https://lwn.net/Articles/816850/> for details.



Alternatives to locking

As we have just seen, locking can have a strong negative impact on system performance. In some situations, you could do without it.

- ▶ By using lock-free algorithms like *Read Copy Update* (RCU).
 - RCU API available in the kernel
 - See <https://en.wikipedia.org/wiki/Read-copy-update> for a coverage of how it works.
- ▶ When relevant, use atomic operations.



- ▶ Conditions where RCU is useful:
 - Frequent reads but infrequent writes
 - Focus on getting consistent data rather than getting the latest data
- ▶ Kind of enforces ownership by enforcing space/time synchronization
- ▶ RCU API ([Documentation/RCU/whatisRCU.rst](https://www.kernel.org/doc/html/latest/Documentation/RCU/whatisRCU.rst)):
 - `rcu_read_lock()` and `rcu_read_unlock()`: reclaim/release read access
 - `synchronize_rcu()`, `call_rcu()` or `kfree_rcu()`: wait for pre-existing readers
 - `rcu_assign_pointer()`: update RCU-protected pointer
 - `rcu_dereference()`: load RCU-protected pointer
- ▶ RCU mentorship session by Paul E. McKenney: <https://youtu.be/K-4TI5gFsig>



RCU example: ensuring consistent accesses (1/2)

Unsafe read/write

```
struct myconf { int a, b; } *shared_conf; /* initialized */

unsafe_get(int *cur_a, int *cur_b)
{
    *cur_a = shared_conf->a;
    /* What if *shared_conf gets updated now? The assignment is inconsistent! */
    *cur_b = shared_conf->b;
};

unsafe_set(int new_a, int new_b)
{
    shared_conf->a = new_a;
    shared_conf->b = new_b;
};
```



RCU example: ensuring consistent accesses (2/2)

Safe read/write with RCU

```
struct myconf { int a, b; } *shared_conf; /* initialized */

safe_get(int *cur_a, int *cur_b)
{
    struct myconf *temp;

    rcu_read_lock();
    temp = rcu_dereference(shared_conf);
    *cur_a = temp->a;
    /* If *shared_conf is updated, temp->a and temp->b will remain consistent! */
    *cur_b = temp->b;
    rcu_read_unlock();
};

safe_set(int new_a, int new_b)
{
    struct myconf *newconf = kmalloc(...);
    struct myconf *oldconf;

    oldconf = rcu_dereference(shared_conf);
    newconf->a = new_a;
    newconf->b = new_b;
    rcu_assign_pointer(shared_conf, newconf);
    /* Readers might still have a reference over the old struct here... */
    synchronize_rcu();
    /* ...but not here! No more readers of the old struct, kfree() is safe! */
    kfree(oldconf);
};
```



Atomic variables 1/2

```
#include <linux/atomic.h>
```

- ▶ Useful when the shared resource is an integer value
- ▶ Even an instruction like `n++` is not guaranteed to be atomic on all processors!
- ▶ Ideal for RMW (Read-Modify-Write) operations
- ▶ Main atomic operations on `atomic_t` (signed integer, at least 24 bits):
 - Set or read the counter:
 - `void atomic_set(atomic_t *v, int i);`
 - `int atomic_read(atomic_t *v);`
 - Operations without return value:
 - `void atomic_inc(atomic_t *v);`
 - `void atomic_dec(atomic_t *v);`
 - `void atomic_add(int i, atomic_t *v);`
 - `void atomic_sub(int i, atomic_t *v);`



Atomic variables 2/2

- ▶ Similar functions testing the result:
 - `int atomic_inc_and_test(...);`
 - `int atomic_dec_and_test(...);`
 - `int atomic_sub_and_test(...);`
- ▶ Functions returning the new value:
 - `int atomic_inc_return(...);`
 - `int atomic_dec_return(...);`
 - `int atomic_add_return(...);`
 - `int atomic_sub_return(...);`



Atomic bit operations

- ▶ Supply very fast, atomic operations
- ▶ On most platforms, apply to an `unsigned long * type`.
- ▶ Apply to a `void * type` on a few others.
- ▶ Ideal for bitmaps
- ▶ Set, clear, toggle a given bit:
 - `void set_bit(int nr, unsigned long *addr);`
 - `void clear_bit(int nr, unsigned long *addr);`
 - `void change_bit(int nr, unsigned long *addr);`
- ▶ Test bit value:
 - `int test_bit(int nr, unsigned long *addr);`
- ▶ Test and modify (return the previous value):
 - `int test_and_set_bit(...);`
 - `int test_and_clear_bit(...);`
 - `int test_and_change_bit(...);`



Kernel locking: summary and references

- ▶ Use mutexes in code that is allowed to sleep
- ▶ Use spinlocks in code that is not allowed to sleep (interrupts) or for which sleeping would be too costly (critical sections)
- ▶ Use atomic operations to protect integers or addresses

See [kernel-hacking/locking](#) in kernel documentation for many details about kernel locking mechanisms.

Further reading: see the classical [dining philosophers problem](#) for a nice illustration of synchronization and concurrency issues.




Image source: https://en.wikipedia.org/wiki/Dining_philosophers_problem




- ▶ Add locking to the driver to prevent concurrent accesses to shared resources



Direct Memory Access

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!






DMA main principles



DMA integration


DMA (*Direct Memory Access*) is used to copy data directly between devices and RAM, without going through the CPU.





Peripheral DMA


Some device controllers embed their own DMA controller and therefore can do DMA on their own.





DMA controllers


Other device controllers rely on an external DMA controller (on the SoC). Their drivers need to submit DMA descriptors to this controller.





DMA descriptors


DMA descriptors describe the various attributes of a DMA transfer, and are chained.





Cache constraints


- ▶ The CPU can access memory through a data cache
 - Using the cache can be more efficient (faster accesses to the cache than the bus)
- ▶ But the DMA does not access the CPU cache, so one needs to take care of cache coherency (cache content vs. memory content):
 - When the CPU reads from memory accessed by DMA, the relevant cache lines must be invalidated to force reading from memory again
 - When the CPU writes to memory before starting DMA transfers, the cache lines must be flushed/cleaned in order to force the data to reach the memory





DMA addressing constraints

- ▶ Memory and devices have physical addresses: `phys_addr_t`
- ▶ CPUs usually access memory through an MMU, using virtual pointers: `void *`
- ▶ DMA controllers do not access memory through the MMU and thus cannot manipulate virtual addresses, instead they access a `dma_addr_t` through either:
 - physical addresses directly
 - an IOMMU, in which case a specific mapping must be created





DMA memory allocation constraints

The APIs must remain generic and handle all cases transparently, hence:

- ▶ Each memory chunk accessed by the DMA shall be physically contiguous, which means one can use:
 - any memory allocated by `kmalloc()` (up to 128 KB)
 - any memory allocated by `__get_free_pages()` (up to 8MB)
 - block I/O and networking buffers, designed to support DMA
- ▶ Unless the buffer is smaller than one page, one cannot use:
 - kernel memory allocated with `vmalloc()`
 - user memory allocated with `malloc()`
 - Almost all the time userspace relies on the kernel to allocate the buffers and `mmap()` them to be usable from userspace (requires a dedicated user API)



Kernel APIs for DMA



dma-mapping vs. dmaengine vs. dma-buf

The dma-mapping API:

- ▶ Allocates and manages DMA buffers
- ▶ Offers generic interfaces to handle coherency
- ▶ Manages IO-MMU DMA mappings when relevant
- ▶ See [core-api/dma-api](#) and [core-api/dma-api-howto](#)

The dmaengine API:

- ▶ Abstracts the DMA controller
- ▶ Offers generic functions to configure, queue, trigger, stop transfers
- ▶ Unused when dealing with peripheral DMA
- ▶ See [driver-api/dmaengine/client](#) and

The dma-buf API:

- ▶ Enables sharing DMA buffers between devices within the kernel
- ▶ Not covered in this training



▶ Coherent mappings

- The kernel allocates a suitable buffer and sets the mapping for the driver
- Can simultaneously be accessed by the CPU and device
- So, has to be in a cache coherent memory area
- Usually allocated for the whole time the module is loaded
 - Can be expensive to setup and use on some platforms
 - Typically implemented by disabling cache on ARM

▶ Streaming mappings

- Use an already allocated buffer
- The driver provides a buffer, the kernel just sets the mapping
- Mapping set up for each transfer (keeps DMA registers free on the hardware)



dma-mapping: memory addressing constraints

- ▶ Any device capable of being master on a bus shall assert its addressing capabilities.
- ▶ The default addressing capability is assumed to be 32-bit.
- ▶ If the platform supports it, the DMA addressing capability can be:
 - increased (eg. need to access highmem)
 - decreased (eg. ISA devices, where `kmalloc()` buffers can also be allocated in the first part of the RAM with `GFP_DMA`)
 - Linux stores this capability in a per-device mask
- ▶ DMA mappings can fail because a buffer is out of reach
 - `swiotlb` may solve some situations by leveraging internal bounce buffers
- ▶ In all cases, the DMA mask shall be consistent before allocating/mapping buffers

```
int dma_set_mask_and_coherent(struct device *dev, u64 mask)
```




dma-mapping: Allocating coherent memory mappings

The kernel takes care of both buffer allocation and mapping:

```
#include <linux/dma-mapping.h>

void *                               /* Output: buffer address */
dma_alloc_coherent(
    struct device *dev, /* device structure */
    size_t size,        /* Needed buffer size in bytes */
    dma_addr_t *handle, /* Output: DMA bus address */
    gfp_t gfp          /* Standard GFP flags */
);

void dma_free_coherent(struct device *dev,
    size_t size, void *cpu_addr, dma_addr_t handle);
```



dma-mapping: Setting up streaming memory mappings (single)

Works on already allocated buffers:

```
#include <linux/dma-mapping.h>
```

```
dma_addr_t dma_map_single(  
    struct device *,           /* device structure */  
    void *,                   /* input: buffer to use */  
    size_t,                   /* buffer size */  
    enum dma_data_direction /* Either DMA_BIDIRECTIONAL,  
                             * DMA_TO_DEVICE or  
                             * DMA_FROM_DEVICE */  
);  
  
void dma_unmap_single(struct device *dev, dma_addr_t handle,  
    size_t size, enum dma_data_direction dir);
```



dma-mapping: Setting up streaming memory mappings (multiples)

A scatterlist using the scatter-gather library can be used to map several buffers and link them together

```
#include <linux/dma-mapping.h>
#include <linux/scatterlist.h>

struct scatterlist sglist[NENTS], *sg;
int i, count;

sg_init_table(sglist, NENTS);
sg_set_buf(&sglist[0], buf0, len0);
sg_set_buf(&sglist[1], buf1, len1);

count = dma_map_sg(dev, sglist, NENTS, DMA_TO_DEVICE);
for_each_sg(sglist, sg, count, i) {
    dma_address[i] = sg_dma_address(sg);
    dma_len[i] = sg_dma_len(sg);
}
...
dma_unmap_sg(dev sglist, count, DMA_TO_DEVICE);
```



dma-mapping: Setting up streaming I/O mappings

Physical addresses with MMIO registers might need to be remapped in order to be accessed through an IO-MMU:

```
#include <linux/dma-mapping.h>
```

```
dma_addr_t dma_map_resource(  
    struct device *,           /* device structure */  
    phys_addr_t,              /* input: resource to use */  
    size_t,                   /* buffer size */  
    enum dma_data_direction, /* Either DMA_BIDIRECTIONAL,  
                             * DMA_TO_DEVICE or  
                             * DMA_FROM_DEVICE */  
    unsigned long attrs,      /* optional attributes */  
);  
  
void dma_unmap_resource(struct device *dev, dma_addr_t handle,  
    size_t size, enum dma_data_direction dir, unsigned long attrs);
```



dma-mapping: Verifying DMA memory mappings

- ▶ All mapping helpers can fail and return errors
- ▶ The right way to check the validity of the returned `dma_addr_t` is to call:
`int dma_mapping_error(struct device *dev, dma_addr_t dma_addr)`
 - May give additional clues if `CONFIG_DMA_API_DEBUG` is enabled.



dma-mapping: Syncing streaming DMA mappings

- ▶ In general streaming mappings are:
 - mapped right before use with DMA
 - `MEM_TO_DEV`: caches are flushed
 - unmapped right after
 - `DEV_TO_MEM`: cache lines are invalidated
- ▶ The CPU shall only access the buffer after unmapping!
- ▶ If however the same memory region has to be used for several DMA transfers, the same mapping can be kept in place. In this case the data must be synchronized before access:
 - The CPU needs to access the data:

```
dma_sync_single_for_cpu(dev, dma_handle, size, direction);  
dma_sync_sg_for_cpu(dev, sglist, nents, direction);
```
 - The device needs to access the data:

```
dma_sync_single_for_device(dev, dma_handle, size, direction);  
dma_sync_sg_for_device(dev, sglist, nents, direction);
```




Starting DMA transfers

- ▶ If the device you're writing a driver for is doing peripheral DMA, no external API is involved.
- ▶ If it relies on an external DMA controller, you'll need to
 1. Ask the hardware to use DMA, so that it will drive its request line
 2. Use Linux `dmaengine` framework, especially its slave API



The dmaengine framework





dmaengine: Slave API: Initial configuration

Steps to start a DMA transfer with dmaengine:

1. Request a channel for exclusive use with `dma_request_chan()`, or one of its variants
 - This channel pointer will be used all along
 - Returns a pointer over a `struct dma_chan` which can also be an error pointer
2. Configure the engine by filling a `struct dma_slave_config` structure and passing it to `dmaengine_slave_config()`:

```
struct dma_slave_config txconf = {};
```

```
/* Tell the engine what configuration we want on a given channel:
 * direction, access size, burst length, source and destination).
 * Source being memory, there is no buswidth or maxburst limitation
 * and each buffer will be different. */
txconf.direction = DMA_MEM_TO_DEV;
txconf.dst_addr_width = DMA_SLAVE_BUSWIDTH_1_BYTE;
txconf.dst_maxburst = TX_TRIGGER;
txconf.dst_addr = fifo_dma_addr;
ret = dmaengine_slave_config(dma->txchan, &txconf);
```



dmaengine: Slave API: Per-transfer configuration (1/2)

1. Create a descriptor with all the required configuration for the next transfer with:

```
struct dma_async_tx_descriptor *  
dmaengine_prep_slave_single(struct dma_chan *chan, dma_addr_t buf,  
                           size_t len, enum dma_transfer_direction dir,  
                           unsigned long flags);  
  
struct dma_async_tx_descriptor *  
dmaengine_prep_slave_sg(struct dma_chan *chan, struct scatterlist *sgl,  
                        unsigned int sg_len, enum dma_transfer_direction dir,  
                        unsigned long flags);  
  
struct dma_async_tx_descriptor *  
dmaengine_prep_dma_cyclic(struct dma_chan *chan, dma_addr_t buf, size_t buf_len,  
                          size_t period_len, enum dma_data_direction dir);
```

- ▶ A common flag is:
 - **DMA_PREP_INTERRUPT**: Generates an interrupt once done
- ▶ The descriptor returned can be used to fill-in a callback:

```
desc->callback = foo_dma_complete;  
desc->callback_param = foo_dev;
```



2. Queue the next operation:

```
dma_cookie_t cookie;
```

```
cookie = dmaengine_submit(desc);  
ret = dma_submit_error(cookie);  
if (ret)  
    ...
```

3. Trigger the queued transfers

```
dma_async_issue_pending(chan);
```

3bis. In case anything went wrong or the device should stop being used, it is possible to terminate all ongoing transactions with:

```
dmaengine_terminate_sync(chan);
```



Examples

- ▶ Commented network driver, with both streaming and coherent mappings:
<https://bootlin.com/pub/drivers/r6040-network-driver-with-comments.c>
- ▶ Example of usage of the slave API: look at the code for `stm32_i2c_prep_dma_xfer()`.



- ▶ Setup streaming mappings with the `dma-mapping` API
- ▶ Configure a DMA controller with the `dmaengine` API
- ▶ Configure the hardware to trigger DMA transfers
- ▶ Wait for DMA completion



Kernel debugging

© Copyright 2004-2026, Bootlin.

Creative Commons BY-SA 3.0 license.

Corrections, suggestions, contributions and translations are welcome!





Debugging/tracing using logs 1/4

▶ Good old `printk()`!

- Works in all contexts
- Can specify a log level ranging from 0 (emergency) to 7 (debug)
- Be careful of the delays introduced when logs are spitted out on a serial console at 115200 bauds
 - A `*_ratelimited()` version exists to limit the amount of print if called too often
- Not recommended for upstream contributions

Example:

```
printk("in probe\n");
```

Here's what you get in the kernel log:

```
[ 1.878382] in probe
```

All other logging facilities are based on it.



Debugging/tracing using logs 2/4

▶ The `pr_*`() family of functions

- They include the log level in the name:
`pr_emerg()`, `pr_alert()`, `pr_crit()`, `pr_err()`, `pr_warn()`, `pr_notice()`,
`pr_info()`, `pr_cont()` and the special `pr_debug()` (see next pages)
- They allow setting a manual prefix (eg. eases grepping):
`#define pr_fmt(fmt) "foo: " fmt`

▶ Also defined in `include/linux/printk.h`

Example:

```
pr_info("in probe\n");
```

Here's what you get in the kernel log:

```
[    1.878382] in probe
```

or similarly with a manual format:

```
[    1.878382] foo: in probe
```




Debugging/tracing using logs 3/4

► The `dev_*()` family of functions

- They include a formatted standard prefix with the device name:
`dev_emerg()`, `dev_alert()`, `dev_crit()`, `dev_err()`, `dev_warn()`, `dev_notice()`,
`dev_info()` and the special `dev_dbg()` (see next pages)
- They additionally take a pointer to `struct device` as first argument
- Defined in `include/linux/dev_printk.h`
- To be used in device drivers

Example:

```
dev_info(&pdev->dev, "in probe\n");
```

Here's what you get in the kernel log:

```
[    1.878382] serial 48024000.serial: in probe
[    1.884873] serial 481a8000.serial: in probe
```



Debugging/tracing using logs 4/4

- ▶ The kernel defines many more format specifiers than the standard `printf()` existing ones.
 - `%p`: Display the hashed value of pointer by default.
 - `%px`: Always display the address of a pointer (use carefully on non-sensitive addresses).
 - `%pK`: Display hashed pointer value, zeros or the pointer address depending on `kptr_restrict` sysctl value.
 - `%pOF`: Device-tree node format specifier.
 - `%pr`: Resource structure format specifier.
 - `%pa`: Physical address display (work on all architectures 32/64 bits)
 - `%pe`: Error pointer (displays the string corresponding to the error number)
- ▶ See [core-api/printk-formats](#) for an exhaustive list of format specifiers
- ▶ Also features a helper to dump entire buffers with a `hexdump` like display: [print_hex_dump\(\)](#)



pr_debug() and dev_dbg()

- ▶ When the driver is compiled with `DEBUG` defined, all these messages are compiled and printed at the debug level. `DEBUG` can be defined by `#define DEBUG` at the beginning of the driver, or using `ccflags-$(CONFIG_DRIVER) += -DDEBUG` in the `Makefile`
- ▶ When the kernel is compiled with `CONFIG_DYNAMIC_DEBUG`, then these messages can dynamically be enabled on a per-file, per-module or per-message basis, by writing commands to `/proc/dynamic_debug/control`. Note that messages are not enabled by default.
 - Details in [admin-guide/dynamic-debug-howto](#)
 - Very powerful feature to only get the debug messages you're interested in.
- ▶ When neither `DEBUG` nor `CONFIG_DYNAMIC_DEBUG` are used, these messages are not compiled in.



Configuring the priority

- ▶ Each message is associated to a priority, as specified in [include/linux/kern_levels.h](#).
- ▶ All the messages, regardless of their priority, are stored in the kernel log ring buffer
 - Typically accessed using the `dmesg` command
- ▶ Messages with a priority lower than the `loglevel` also appear on the console
- ▶ The `loglevel` can be changed:
 - in the kernel configuration using `CONFIG_CONSOLE_LOGLEVEL_DEFAULT`
 - on the cmdline with `loglevel=` (see [admin-guide/kernel-parameters](#))
 - at runtime through `/proc/sys/kernel/printk` ([admin-guide/sysctl/kernel](#))
- ▶ Examples:
 - `loglevel=0`: no message on the console (see also: `quiet`)
 - `loglevel=8`: all messages on the console (see also: `ignore_loglevel`)



A virtual filesystem to export debugging information to user space.

- ▶ Kernel configuration: [CONFIG_DEBUG_FS](#)
 - Kernel hacking -> Debug Filesystem
- ▶ The debugging interface disappears when Debugfs is configured out.
- ▶ You can mount it as follows:
 - `sudo mount -t debugfs none /sys/kernel/debug`
- ▶ First described on <https://lwn.net/Articles/115405/>
- ▶ API documented in the Linux Kernel Filesystem API: [filesystems/debugfs](#) The debugfs filesystem



DebugFS API

- ▶ Create a sub-directory for your driver:

```
struct dentry *debugfs_create_dir(const char *name,  
                                struct dentry *parent);
```

- ▶ Expose an integer as a file in DebugFS. Example:

```
struct dentry *debugfs_create_u8(const char *name, mode_t mode,  
                                struct dentry *parent, u8 *value);
```

- u8, u16, u32, u64 for decimal representation
- x8, x16, x32, x64 for hexadecimal representation

- ▶ Expose a binary blob as a file in DebugFS:

```
struct dentry *debugfs_create_blob(const char *name,  
                                  mode_t mode, struct dentry *parent,  
                                  struct debugfs_blob_wrapper *blob);
```

- ▶ Also possible to support writable DebugFS files or customize the output using the more generic `debugfs_create_file()` function.



Using Magic SysRq

Functionnality provided by serial drivers

- ▶ Allows to run multiple debug / rescue commands even when the kernel seems to be in deep trouble
 - On PC: press [Alt] + [Prnt Scrn] + <character> simultaneously ([SysRq] = [Alt] + [Prnt Scrn])
 - On embedded: in the console, send a break character (Picocom: press [Ctrl] + a followed by [Ctrl] + \), then press <character>
- ▶ Example commands:
 - h: show available commands
 - s: sync all mounted filesystems
 - b: reboot the system
 - n: makes RT processes nice-able.
 - w: shows the kernel stack of all sleeping processes
 - t: shows the kernel stack of all running processes
 - You can even register your own!
- ▶ Detailed in [admin-guide/sysrq](#)



kgdb - A kernel debugger

- ▶ `CONFIG_KGDB` in *Kernel hacking*.
- ▶ The execution of the kernel is fully controlled by `gdb` from another machine, connected through a serial line.
- ▶ Can do almost everything, including inserting breakpoints in interrupt handlers.
- ▶ Feature supported for the most popular CPU architectures
- ▶ `CONFIG_GDB_SCRIPTS` allows to build GDB python scripts that are provided by the kernel.
 - See [process/debugging/kgdb](https://bootlin.com/process/debugging/kgdb) for more information



Using kgdb (1/2)

- ▶ Details available in the kernel documentation: [process/debugging/kgdb](#)
- ▶ You must include a kgdb I/O driver. One of them is kgdb over serial console (kgdboc: kgdb over console, enabled by [CONFIG_KGDB_SERIAL_CONSOLE](#))
- ▶ Configure kgdboc at boot time by passing to the kernel:
 - `kgdboc=<tty-device>,<bauds>.`
 - For example: `kgdboc=ttys0,115200`
- ▶ Or at runtime using sysfs:
 - `echo ttys0 > /sys/module/kgdboc/parameters/kgdboc`
 - If the console does not have polling support, this command will yield an error.



Using kgdb (2/2)


- ▶ Then also pass `kgdbwait` to the kernel: it makes `kgdb` wait for a debugger connection.
- ▶ Boot your kernel, and when the console is initialized, interrupt the kernel with a break character and then `g` in the serial console (see our *Magic SysRq* explanations).
- ▶ On your workstation, start `gdb` as follows:
 - `arm-linux-gdb ./vmlinux`
 - `(gdb) set remotebaud 115200`
 - `(gdb) target remote /dev/ttyS0`
- ▶ Once connected, you can debug a kernel the way you would debug an application program.
- ▶ On GDB side, the first threads represent the CPU context (`ShadowCPU<x>`), then all the other threads represents a task.



Debugging with a JTAG interface

Two types of JTAG dongles

- ▶ The ones offering a `gdb` compatible interface, over a serial port or an Ethernet connection. `gdb` can directly connect to them.
- ▶ The ones not offering a `gdb` compatible interface are generally supported by OpenOCD (Open On Chip Debugger): <https://openocd.sourceforge.net/>
 - OpenOCD is the bridge between the `gdb` debugging language and the JTAG interface of the target CPU.
 - See the very complete documentation: <https://openocd.org/pages/documentation.html>
 - For each board, you'll need an OpenOCD configuration file (ask your supplier)





Early traces

- ▶ If something breaks before the `tty` layer, serial driver and serial console are properly registered, you might just have nothing else after "Starting kernel..."
- ▶ On ARM, if your platform implements it, you can activate (`CONFIG_DEBUG_LL` and `CONFIG_EARLY_PRINTK`), and add `earlyprintk` to the kernel command line
 - Assembly routines to just push a character and wait for it to be sent
 - Extremely basic, but is part of the uncompressed section, so available even if the kernel does not uncompress correctly!
- ▶ On other platforms, hoping that your serial driver implements `OF_EARLYCON_DECLARE()`, you can enable `CONFIG_SERIAL_EARLYCON`
 - The kernel will try to hook an appropriate `earlycon` UART driver using the `stdout-path` of the device-tree.



More kernel debugging tips 1/2

- ▶ Make sure `CONFIG_KALLSYMS_ALL` is enabled
 - To get oops messages with symbol names instead of raw addresses
 - Turned on by default
- ▶ Make sure `CONFIG_DEBUG_INFO` is also enabled
 - This way, the kernel is compiled with `$(CROSSCOMPILE)gcc -g`, which keeps the source code inside the binaries.
- ▶ If your device is not probed, try enabling `CONFIG_DEBUG_DRIVER`
 - Extremely verbose!
 - Will enable all the debug logs in the device-driver core section



More kernel debugging tips 2/2

Device Tree output can be better understood with the in-tree `dtc/dtx_diff` script

- ▶ Show the origin of each property/node: `scripts/dtc/dtx_diff -T <dts>`

```
main_i2c3: i2c@20030000 { /* arch/arm64/boot/dts/ti/k3-am62-main.dtsi:450:26-460:4,  
                        arch/arm64/boot/dts/ti/k3-am625-beagleplay.dts:820:12-825:3,  
                        arch/arm64/boot/dts/ti/k3-am625-beagleplay-custom.dts:20:12-25:3 */  
  
    ...  
    reg = <0x00 0x20030000 0x00 0x100>; /* arch/arm64/boot/dts/ti/k3-am62-main.dtsi:452:3-452:38 */  
    status = "okay"; /* arch/arm64/boot/dts/ti/k3-am625-beagleplay.dts:824:2-824:18 */  
  
    joystick@52 { /* arch/arm64/boot/dts/ti/k3-am625-beagleplay-custom.dts:21:14-24:4 */  
        compatible = "nintendo,nunchuk"; /* arch/arm64/boot/dts/ti/k3-am625-beagleplay-custom.dts:23  
        reg = <0x52>; /* arch/arm64/boot/dts/ti/k3-am625-beagleplay-custom.dts:22:3-22:16 */  
    };  
};
```

- ▶ Can also be used to diff DTS



Getting help and reporting bugs


- ▶ If you are using a custom kernel from a hardware vendor, contact that company. The community will have less interest supporting a custom kernel.
- ▶ Otherwise, or if this doesn't work, try to reproduce the issue on the latest version of the kernel.
- ▶ Make sure you investigate the issue as much as you can: see [admin-guide/bug-bisect](#)
- ▶ Check for previous bugs reports. Use web search engines, accessing public mailing list archives.
- ▶ If you're the first to face the issue, it's very useful for others to report it, even if you cannot investigate it further.
- ▶ If the subsystem you report a bug on has a mailing list, use it. Otherwise, contact the official maintainer (see the [MAINTAINERS](#) file). Always give as many useful details as possible.



Debugging resources

Checkout Bootlin's debugging training!

- ▶ Linux debugging, profiling, tracing and performance analysis training
- ▶ <https://bootlin.com/doc/training/debugging/>
- ▶ Slides and labs are available for free





Practical lab - Kernel debugging




- ▶ Use the dynamic debug feature.
- ▶ Add debugfs entries
- ▶ Load a broken driver and see it crash
- ▶ Analyze the error information dumped by the kernel.
- ▶ Disassemble the code and locate the exact C instruction which caused the failure.



Power Management

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





- ▶ Several power management *building blocks*
 - Clock framework
 - CPUidle (system wide)
 - Frequency and voltage scaling (system wide)
 - Suspend and resume (system wide)
 - Runtime power management (peripheral specific)
 - Power domains (SoC specific)
- ▶ Independent *building blocks* that can be improved gradually during development



Clock framework (1)

- ▶ Generic framework to manage clocks used by devices in the system
- ▶ Allows to reference count clock users and to shutdown the unused clocks to save power
- ▶ Simple API described in [include/linux/clock.h](#).
 - [clk_get\(\)](#) to lookup and obtain a reference to a clock producer
 - [clk_put\(\)](#) to free the clock source
 - [clk_prepare_enable\(\)](#) to inform the system when the clock source should be running
 - [clk_disable_unprepare\(\)](#) to inform the system when the clock source is no longer required.
 - [clk_get_rate\(\)](#) to obtain the current clock rate (in Hz) for a clock source
 - [clk_set_rate\(\)](#) to set the current clock rate (in Hz) of a clock source




Clock framework (2)

The common clock framework

- ▶ Allows to declare the available clocks and their association to devices in the Device Tree
- ▶ Provides a *debugfs* representation of the clock tree
- ▶ Is implemented in `drivers/clock/`



Diagram overview of the common clock framework





Clock framework (3)

The interface of the CCF divided into two halves:

- ▶ Common Clock Framework core
 - Common definition of `struct clk`
 - Common implementation of the `clk.h` API (defined in `drivers/clk/clk.c`)
 - `struct clk_ops`: operations invoked by the `clk` API implementation
 - Not supposed to be modified when adding a new driver
- ▶ Hardware-specific
 - Callbacks registered with `struct clk_ops` and the corresponding hardware-specific structures
 - Has to be written for each new hardware clock
 - Example: `drivers/clk/mvebu/clk-cpu.c`



Clock framework (4)

Hardware clock operations: device tree

- ▶ The **device tree** is the **mandatory way** to declare a clock and to get its resources, as for any other driver using DT we have to:
 - **Parse** the device tree to **setup** the clock: the resources but also the properties are retrieved.
 - Declare the **compatible** clocks and associate each to an **initialization** function using `CLK_OF_DECLARE()`
 - Example: `arch/arm/boot/dts/marvell/armada-xp.dtsi` and `drivers/clk/mvebu/armada-xp.c`

See our presentation about the Clock Framework for more details:

<https://bootlin.com/pub/conferences/2013/elce/common-clock-framework-how-to-use-it/>



Saving power in the idle loop

- ▶ The idle loop is what you run when there's nothing left to run in the system.
- ▶ `arch_cpu_idle()` implemented in all architectures in `arch/<arch>/kernel/process.c`
- ▶ Example: `arch/arm/kernel/process.c`
- ▶ The CPU can run power saving HLT instructions, enter NAP mode, and even disable the timers (tickless systems).
- ▶ See also https://en.wikipedia.org/wiki/Idle_loop



Adding support for multiple idle levels

- ▶ Modern CPUs have several sleep states offering different power savings with associated wake up latencies
- ▶ The *dynamic tick* feature allows to remove the periodic timer tick to save power, and to know when the next event is scheduled, for smarter sleeps.
- ▶ CPUidle infrastructure to change sleep states
 - Platform-specific driver defining sleep states and transition operations
 - Platform-independent governors
 - Available in particular for x86/ACPI and most ARM SoCs
 - See [admin-guide/pm/cpuidle](#) in kernel documentation.



<https://en.wikipedia.org/wiki/PowerTOP>

- ▶ With dynamic ticks, allows to fix parts of kernel code and applications that wake up the system too often.
- ▶ PowerTOP allows to track the worst offenders
- ▶ Now available on ARM cpus implementing CPUidle
- ▶ Also gives you useful hints for reducing power.
- ▶ Try it on your x86 laptop:
`sudo powertop`



Frequency and voltage scaling (1)

Frequency and voltage scaling possible through the `cpufreq` kernel infrastructure.

- ▶ Generic infrastructure: `drivers/cpufreq/cpufreq.c` and `include/linux/cpufreq.h`
- ▶ Generic governors, responsible for deciding frequency and voltage transitions
 - performance: maximum frequency
 - powersave: minimum frequency
 - ondemand: measures CPU consumption to adjust frequency
 - conservative: often better than ondemand. Only increases frequency gradually when the CPU gets loaded.
 - schedutil: Tightly integrated with the scheduler, making per-policy decisions, RT tasks running at full speed.
 - userspace: leaves the decision to a user space daemon.
- ▶ This infrastructure can be controlled from `/sys/devices/system/cpu/cpu<n>/cpufreq/`



Frequency and voltage scaling (2)

- ▶ CPU frequency drivers are in `drivers/cpufreq/`. Example: `drivers/cpufreq/omap-cpufreq.c`
- ▶ Must implement the operations of the `cpufreq_driver` structure and register them using `cpufreq_register_driver()`
 - `init()` for initialization
 - `exit()` for cleanup
 - `verify()` to verify the user-chosen policy
 - `setpolicy()` or `target()` to actually perform the frequency change
- ▶ See documentation in `cpu-freq/` for useful explanations



Suspend and resume (to / from RAM)

- ▶ Infrastructure in the kernel to support suspend and resume
- ▶ System on Chip hooks
 - Define operations (at least `valid()` and `enter()`) `struct platform_suspend_ops` structure. See the documentation for this structure for details about possible operations and the way they are used.
 - Registered using the `suspend_set_ops()` function
 - See [arch/arm/mach-at91/pm.c](#)
- ▶ Device driver hooks
 - pm operations (`suspend()` and `resume()` hooks) in the `struct device_driver` as a `struct dev_pm_ops` structure in (`struct platform_driver`, `struct usb_driver`, etc.)
 - See [drivers/net/ethernet/cadence/macb_main.c](#)
- ▶ *Hibernate to disk* is based on suspend to RAM, copying the RAM contents (after a simulated suspend) to a swap partition.



Triggering suspend / hibernate

- ▶ `struct suspend_ops` functions are called by the `enter_state()` function. `enter_state()` also takes care of executing the suspend and resume functions for your devices.
- ▶ Read `kernel/power/suspend.c`
- ▶ The execution of this function can be triggered from user space:
 - `echo mem > /sys/power/state` (suspend to RAM)
 - `echo disk > /sys/power/state` (hibernate to disk)
- ▶ Systemd can also manage suspend and hibernate for you, and offers customizations
 - `systemctl suspend` or `systemctl hibernate`.
 - See <https://www.man7.org/linux/man-pages/man8/systemd-suspend.service.8.html>



Runtime power management

- ▶ Managing per-device idle, each device being managed by its device driver independently from others.
- ▶ According to the kernel configuration interface: *Enable functionality allowing I/O devices to be put into energy-saving (low power) states at run time (or autosuspended) after a specified period of inactivity and woken up in response to a hardware-generated wake-up event or a driver's request.*
- ▶ New hooks must be added to the drivers: `runtime_suspend()`, `runtime_resume()`, `runtime_idle()` in the `struct dev_pm_ops` structure in `struct device_driver`.
- ▶ API and details on [power/runtime_pm](#)
- ▶ See [drivers/net/ethernet/cadence/macb_main.c](#) again.



Generic PM Domains (pmdomain, ex-genpd)

- ▶ Generic infrastructure to implement power domains based on Device Tree descriptions, allowing to group devices by the physical power domain they belong to. This sits at the same level as bus type for calling PM hooks.
- ▶ All the devices in the same PD get the same state at the same time.
- ▶ Specifications and examples available at Documentation/devicetree/bindings/power/power_domain.txt
- ▶ Driver example: <drivers/pmdomain/rockchip/pm-domains.c>
See [_pd_power_on\(\)](#), [_pd_power_off\(\)](#), [_pm_add_one_domain\(\)](#)
- ▶ DT example: look for [rockchip,px30-power-controller](#) ([arch/arm64/boot/dts/rockchip/px30.dtsi](#)) and find PD definitions and corresponding devices.
- ▶ See Kevin Hilman's talk at Kernel Recipes 2017:
<https://youtu.be/SctfvoskABM>




- ▶ [power/](#) in kernel documentation.
 - Will give you many useful details.
- ▶ Introduction to kernel power management, Kevin Hilman (Kernel Recipes 2015)
 - <https://www.youtube.com/watch?v=juJJZ0RgVwI>
- ▶ Linux Power Management Features, Their Relationships and Interactions — Théo Lebrun (Embedded Linux Conference Europe 2024)
 - https://www.youtube.com/watch?v=_jb6U40ZCZk



Kernel Resources

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Linux Weekly News

- ▶ <https://lwn.net/>
- ▶ The weekly digest off all Linux and free software information sources
- ▶ In depth technical discussions about the kernel
- ▶ Coverage of the features accepted in each merge window
- ▶ Subscribe to finance the editors (\$7 / month)
- ▶ Articles available for non subscribers after 1 week.



LWN.net



Useful Online Resources

- ▶ Kernel documentation
 - <https://kernel.org/doc/>
- ▶ Linux kernel mailing list FAQ
 - <https://subspace.kernel.org/etiquette.html>
 - Complete Linux kernel FAQ
 - Read this before asking a question to the mailing list
- ▶ Linux kernel mailing lists
 - <https://lore.kernel.org/>
 - Easy browsing and referencing of all e-mail threads
 - Easy access to an mbox in order to answer to e-mails you were not Cc'ed to
- ▶ Kernel Newbies
 - <https://kernelnewbies.org/>
 - Articles, presentations, HOWTOs, recommended reading, useful tools for people getting familiar with Linux kernel or driver development.
 - Glossary: <https://kernelnewbies.org/KernelGlossary>
 - In depth coverage of the new features in each kernel release:
<https://kernelnewbies.org/LinuxChanges>
- ▶ The <https://elinux.org> wiki




International Conferences (1)

▶ Embedded Linux Conference:

- <https://embeddedlinuxconference.com/>
- Organized by the Linux Foundation
- Once per year, alternating North America/Europe
- Very interesting kernel and user space topics for embedded systems developers. Many kernel and embedded project maintainers are present.
- Presentation slides and videos freely available on https://elinux.org/ELC_Presentations

▶ Linux Plumbers

- <https://lpc.events/>
- About the low-level plumbing of Linux: kernel, audio, power management, device management, multimedia, etc.
- Not really a conventional conference with formal presentations, but rather a place where contributors on each topic meet, share their progress and make plans for work ahead.





International Conferences (2)

- ▶ Kernel Recipes: <https://kernel-recipes.org/>
 - Well attended conference in Europe (Paris), only one track at a time, with a format that really allows for discussions.
- ▶ linux.conf.au:
<https://linux.org.au/linux-conf-au/>
 - In Australia / New Zealand
 - Features a few presentations by key kernel hackers.
- ▶ Currently, most conferences are available on-line. They are much more affordable and often free.



LINUXCONFAU



After the course

Continue to learn:

- ▶ Run your labs again on your own hardware. The Nunchuk lab should be rather straightforward, but the serial lab will be quite different if you use a different processor.
- ▶ Learn by reading the kernel code by yourself, ask questions and propose improvements.
- ▶ Implement and share drivers for your own hardware, of course!

Hobbyists can make their first contributions by:

- ▶ Helping with tasks keeping the kernel code clean and up-to-date:
<https://kernelnewbies.org/KernelJanitors/ToDo>
- ▶ Proposing fixes for issues reported by the *Coccinelle* tool:
`make coccicheck`
- ▶ Participating to improving drivers in [drivers/staging/](#)
- ▶ Investigating and do the triage of issues reported by Coverity Scan:
<https://scan.coverity.com/projects/linux>



Contribute your changes

Recommended resources

- ▶ See [process/submitting-patches](https://kernelnewbies.org/UpstreamMerge) for guidelines and <https://kernelnewbies.org/UpstreamMerge> for very helpful advice to have your changes merged upstream (by Rik van Riel).
- ▶ Watch the *Write and Submit your first Linux kernel Patch* talk by Greg. K.H: <https://www.youtube.com/watch?v=LLBrBBImJt4>
- ▶ How to Participate in the Linux Community (by Jonathan Corbet).
A guide to the kernel development process.
https://www.static.linuxfound.org/sites/lfcorp/files/How-Participate-Linux-Community_0.pdf




Last slides

© Copyright 2004-2026, Bootlin.

Creative Commons BY-SA 3.0 license.

Corrections, suggestions, contributions and translations are welcome!





Thank you!
And may the Source be with you



Rights to copy

© Copyright 2004-2026, Bootlin

License: Creative Commons Attribution - Share Alike 3.0

<https://creativecommons.org/licenses/by-sa/3.0/legalcode>

You are free:

- ▶ to copy, distribute, display, and perform the work
- ▶ to make derivative works
- ▶ to make commercial use of the work

Under the following conditions:

- ▶ **Attribution.** You must give the original author credit.
- ▶ **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.
- ▶ For any reuse or distribution, you must make clear to others the license terms of this work.
- ▶ Any of these conditions can be waived if you get permission from the copyright holder.


Your fair use and other rights are in no way affected by the above.

Document sources: <https://github.com/bootlin/training-materials/>



Backup slides

© Copyright 2004-2026, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





mmap



- ▶ Possibility to have parts of the virtual address space of a program mapped to the contents of a file
- ▶ Particularly useful when the file is a device file
- ▶ Allows to access device I/O memory and ports without having to go through (expensive) read, write or ioctl calls
- ▶ One can access to current mapped files by two means:
 - `/proc/<pid>/maps`
 - `pmap <pid>`




/proc/<pid>/maps

start-end	perm	offset	major:minor	inode	mapped file name
...					
7f4516d04000-7f4516d06000	rw-s	1152a2000	00:05	8406	/dev/dri/card0
7f4516d07000-7f4516d0b000	rw-s	120f9e000	00:05	8406	/dev/dri/card0
...					
7f4518728000-7f451874f000	r-xp	00000000	08:01	268909	/lib/x86_64-linux-gnu/libexpat.so.1.5.2
7f451874f000-7f451894f000	---p	00027000	08:01	268909	/lib/x86_64-linux-gnu/libexpat.so.1.5.2
7f451894f000-7f4518951000	r--p	00027000	08:01	268909	/lib/x86_64-linux-gnu/libexpat.so.1.5.2
7f4518951000-7f4518952000	rw-p	00029000	08:01	268909	/lib/x86_64-linux-gnu/libexpat.so.1.5.2
...					
7f451da4f000-7f451dc3f000	r-xp	00000000	08:01	1549	/usr/bin/Xorg
7f451de3e000-7f451de41000	r--p	001ef000	08:01	1549	/usr/bin/Xorg
7f451de41000-7f451de4c000	rw-p	001f2000	08:01	1549	/usr/bin/Xorg
...					



mmap overview





How to Implement mmap - User space

- ▶ Open the device file
- ▶ Call the `mmap` system call (see `man mmap` for details):

```
void * mmap(  
    void *start,    /* Often 0, preferred starting address */  
    size_t length,  /* Length of the mapped area */  
    int prot,       /* Permissions: read, write, execute */  
    int flags,      /* Options: shared mapping, private copy... */  
    int fd,         /* Open file descriptor */  
    off_t offset    /* Offset in the file */  
);
```

- ▶ You get a virtual address you can write to or read from.



How to Implement mmap - Kernel space

- ▶ Character driver: implement an `mmap` file operation and add it to the driver file operations:

```
int (*mmap) (  
    struct file *,           /* Open file structure */  
    struct vm_area_struct * /* Kernel VMA structure */  
);
```

- ▶ Initialize the mapping.
 - Can be done in most cases with the `remap_pfn_range()` function, which takes care of most of the job.



remap_pfn_range()

- ▶ *pfn*: page frame number
- ▶ The most significant bits of the page address (without the bits corresponding to the page size).

```
#include <linux/mm.h>
```

```
int remap_pfn_range(  
    struct vm_area_struct *, /* VMA struct */  
    unsigned long virt_addr, /* Starting user  
                             * virtual address */  
    unsigned long pfn,       /* pfn of the starting  
                             * physical address */  
    unsigned long size,      /* Mapping size */  
    pgprot_t prot           /* Page permissions */  
);
```



Simple mmap implementation

```
static int acme_mmap
(struct file * file, struct vm_area_struct *vma)
{
    size = vma->vm_end - vma->vm_start;

    if (size > ACME_SIZE)
        return -EINVAL;

    if (remap_pfn_range(vma,
                        vma->vm_start,
                        ACME_PHYS >> PAGE_SHIFT,
                        size,
                        vma->vm_page_prot))
        return -EAGAIN;

    return 0;
}
```



- ▶ <https://bootlin.com/pub/mirror/devmem2.c>, by Jan-Derk Bakker
- ▶ Very useful tool to directly peek (read) or poke (write) I/O addresses mapped in physical address space from a shell command line!
 - Very useful for early interaction experiments with a device, without having to code and compile a driver.
 - Uses `mmap` to `/dev/mem`.
 - Examples (b: byte, h: half, w: word)
 - `devmem2 0x000c0004 h` (reading)
 - `devmem2 0x000c0008 w 0xffffffff` (writing)
 - `devmem` is now available in BusyBox, making it even easier to use.



mmap summary

- ▶ The device driver is loaded. It defines an `mmap` file operation.
- ▶ A user space process calls the `mmap` system call.
- ▶ The `mmap` file operation is called.
- ▶ It initializes the mapping using the device physical address.
- ▶ The process gets a starting address to read from and write to (depending on permissions).
- ▶ The MMU automatically takes care of converting the process virtual addresses into physical ones.
- ▶ Direct access to the hardware without any expensive `read` or `write` system calls



Useful general-purpose kernel APIs



- ▶ In `include/linux/string.h`
 - Memory-related: `memset()`, `memcpy()`, `memmove()`, `memscan()`, `memcmp()`, `memchr()`
 - String-related: `strcpy()`, `strcat()`, `strcmp()`, `strchr()`, `strrchr()`, `strlen()` and variants
 - Allocate and copy a string: `kstrdup()`, `kstrndup()`
 - Allocate and copy a memory area: `kmemdup()`
- ▶ In `include/linux/kernel.h`
 - String to int conversion: `simple_strtoul()`, `simple_strtol()`, `simple_strtoull()`, `simple_strtoll()`
 - Other string functions: `sprintf()`, `sscanf()`



Linked lists

- ▶ Convenient linked-list facility in `include/linux/list.h`
 - Used in thousands of places in the kernel
- ▶ Add a `struct list_head` member to the structure whose instances will be part of the linked list. It is usually named `node` when each instance needs to only be part of a single list.
- ▶ Define the list with the `LIST_HEAD()` macro for a global list, or define a `struct list_head` element and initialize it with `INIT_LIST_HEAD()` for lists embedded in a structure.
- ▶ Then use the `list_*`() API to manipulate the list
 - Add elements: `list_add()`, `list_add_tail()`
 - Remove, move or replace elements: `list_del()`, `list_move()`, `list_move_tail()`, `list_replace()`
 - Test the list: `list_empty()`
 - Iterate over the list: `list_for_each_*`() family of macros



Linked lists examples 1/2

drivers/i2c/busses/i2c-stm32f7.c

```
/**
 * struct stm32f7_i2c_timings - private I2C output parameters
 * @node: List entry
 * @presc: Prescaler value
 * @scldel: Data setup time
 * @sdadel: Data hold time
 * @sclh: SCL high period (master mode)
 * @scll: SCL low period (master mode)
 */
struct stm32f7_i2c_timings {
    struct list_head node;
    u8 presc;
    u8 scldel;
    u8 sdadel;
    u8 sclh;
    u8 scll;
};
```



Linked lists examples 2/2

drivers/i2c/busses/i2c-stm32f7.c

```
static int stm32f7_i2c_compute_timing(/* ... */)
{
    struct stm32f7_i2c_timings *v;
    struct list_head solutions;
    INIT_LIST_HEAD(&solutions);
    /* ... */

    for (p = 0; p < STM32F7_PRESC_MAX; p++) {
        for (l = 0; l < STM32F7_SCLDEL_MAX; l++) {
            v = kmalloc(sizeof(*v), GFP_KERNEL);
            v->presc = p;
            v->scldel = l;
            list_add_tail(&v->node, &solutions);
        }
    }

    list_for_each_entry(v, &solutions, node) {
        /* ... */
    }
}
```